



# **Intro to Neo4j and Graph Databases**

Neo4j Webinar March 2016

# ABOUT ME

- William Lyon
- Developer Relations Engineer @neo4j
- <http://neo4j.com/developer>



[will@neo4j.com](mailto:will@neo4j.com)

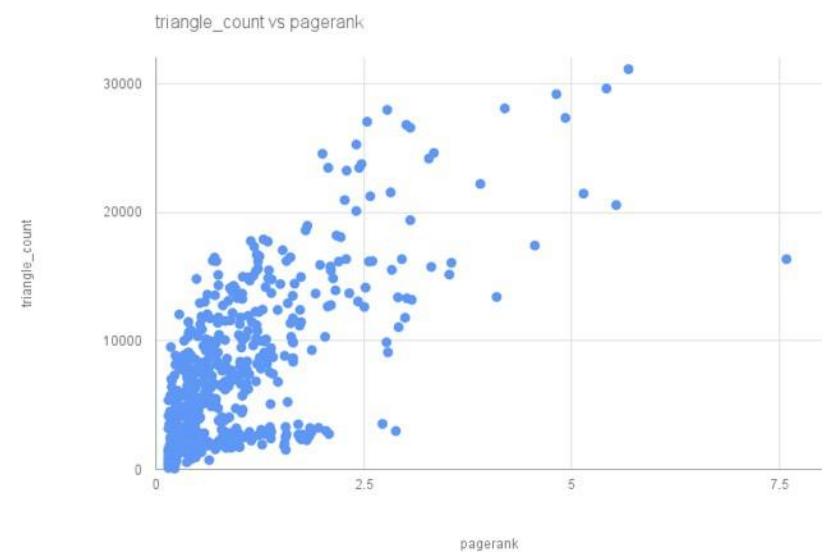
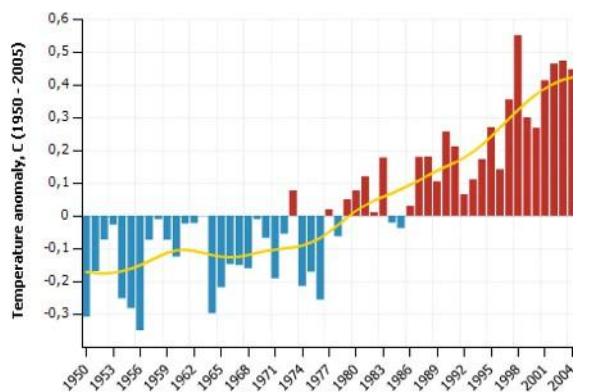
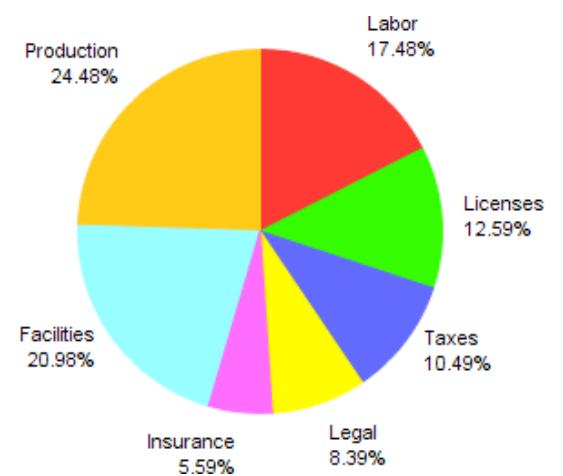
@lyonwj



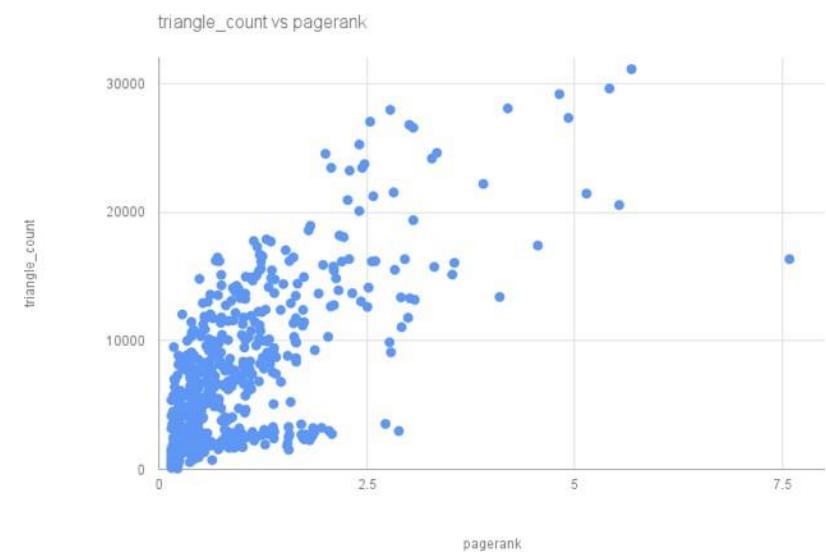
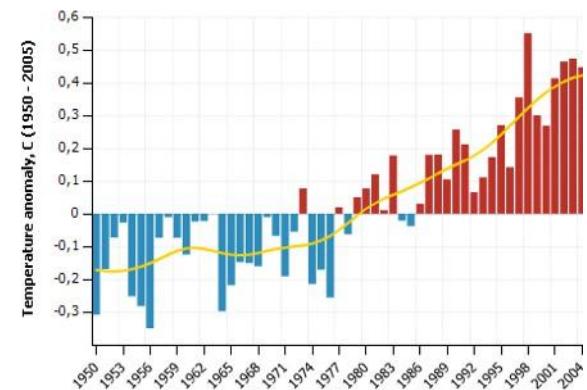
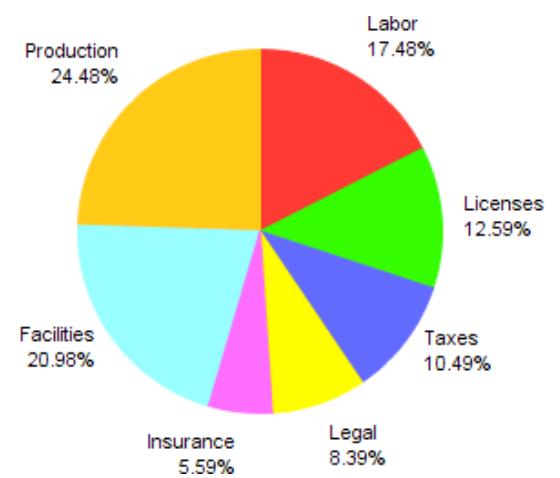
# Agenda

- What is a graph [database]?
- Use cases - why graphs?
- Neo4j product overview
  - Labeled property graph data model
  - Cypher query language
- RDBMS to graph
- Resources
- Questions?

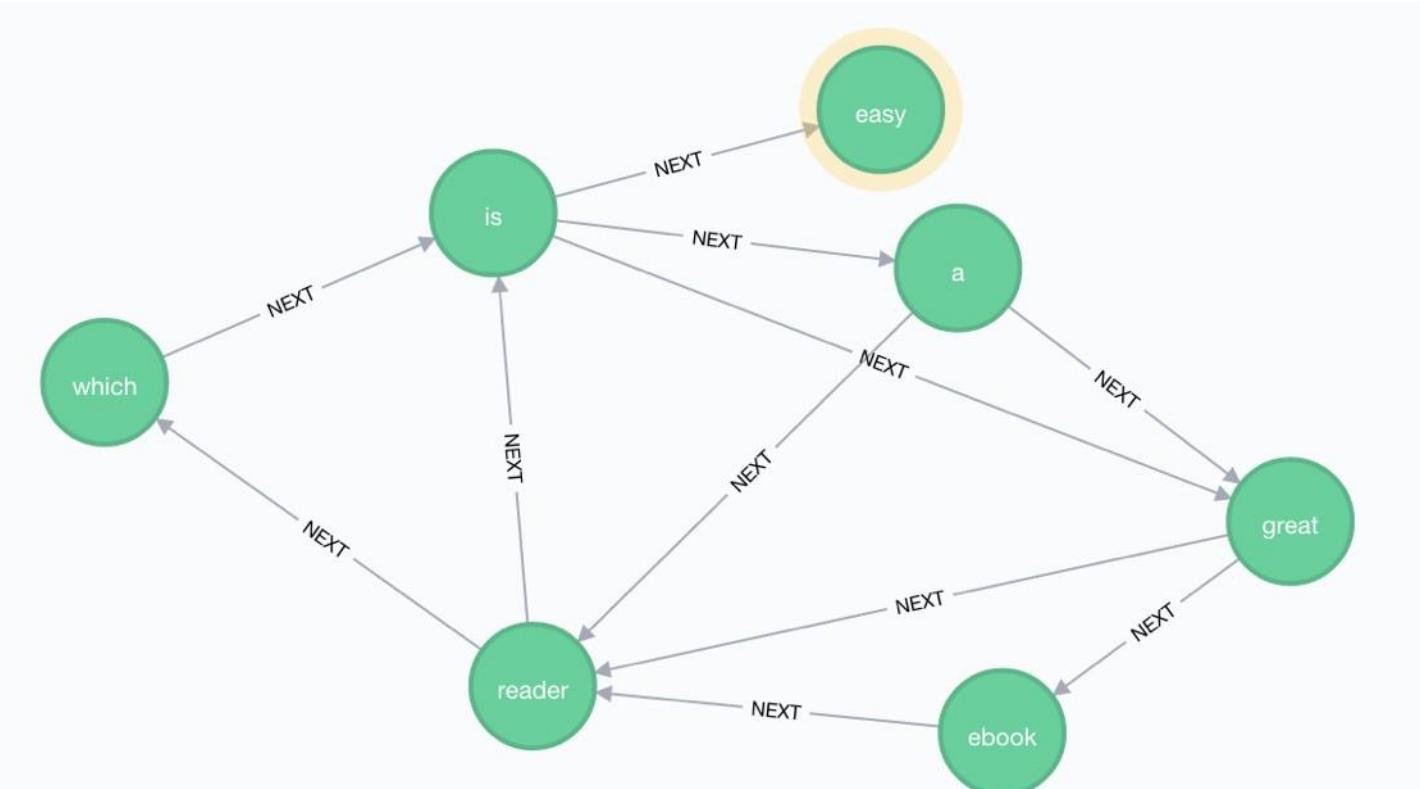
# Chart

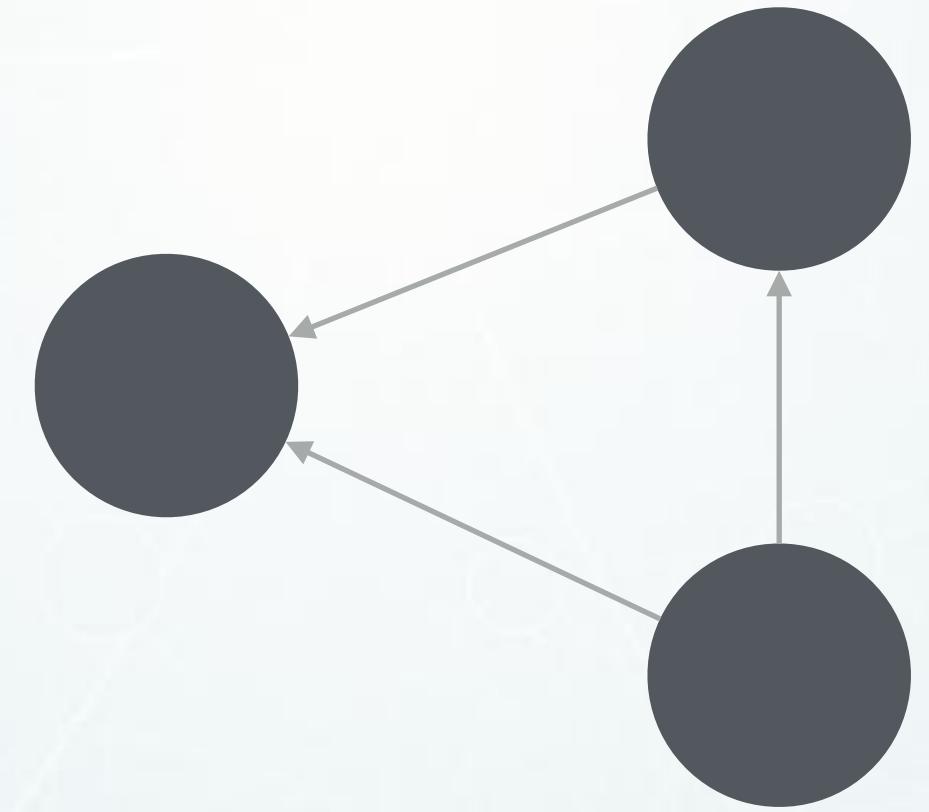


# Chart

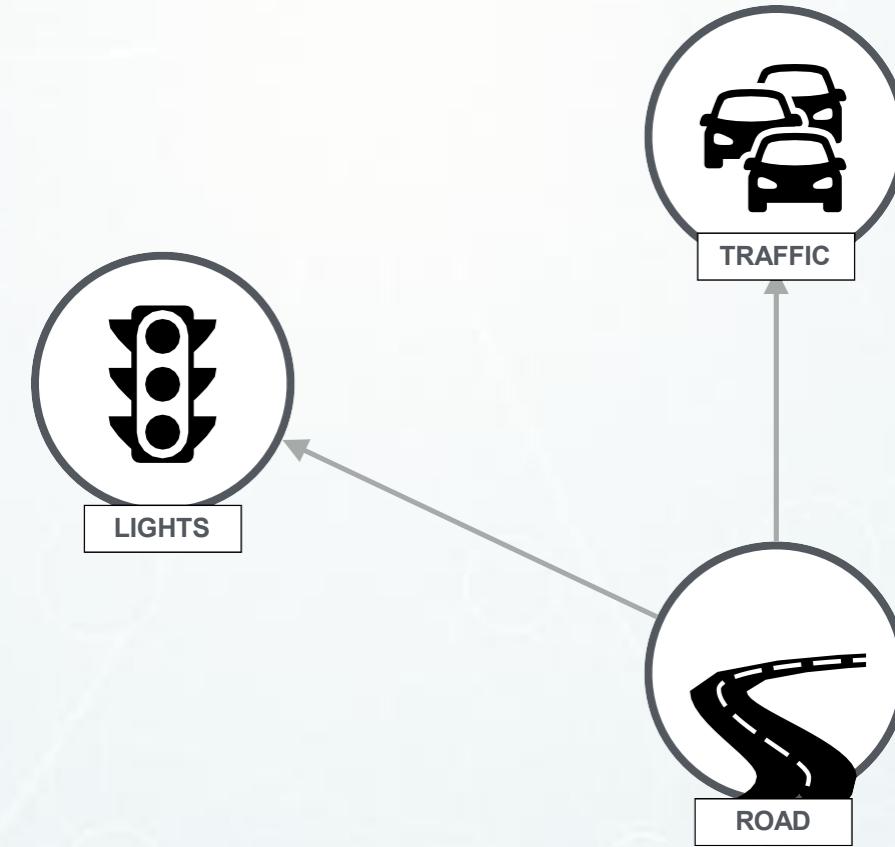


# Graph

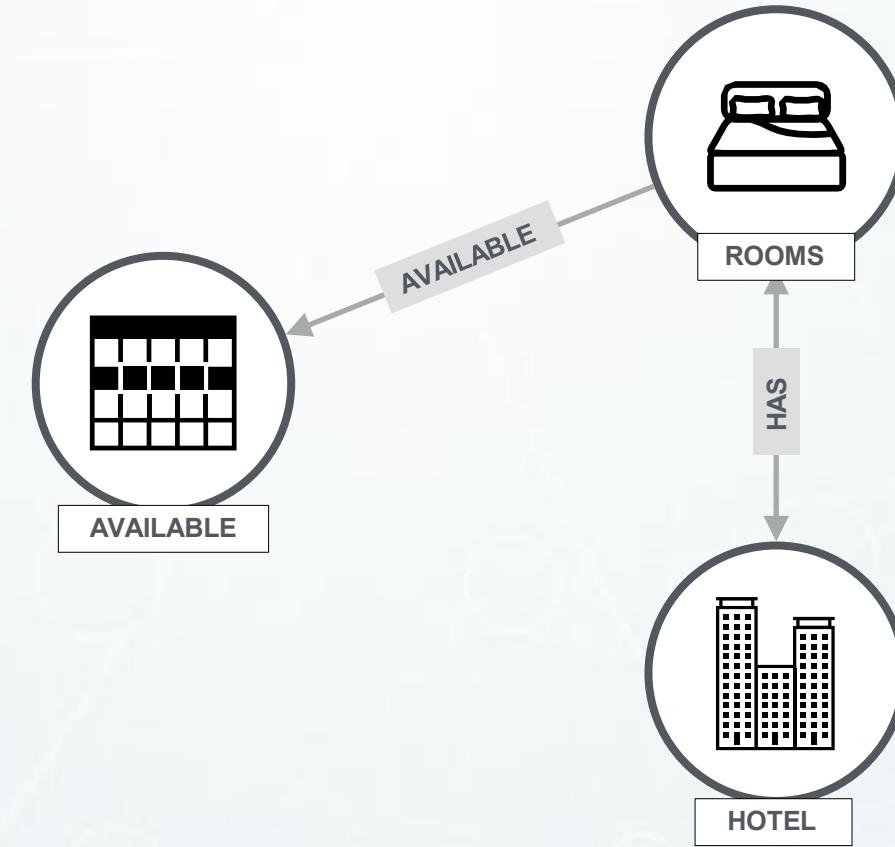




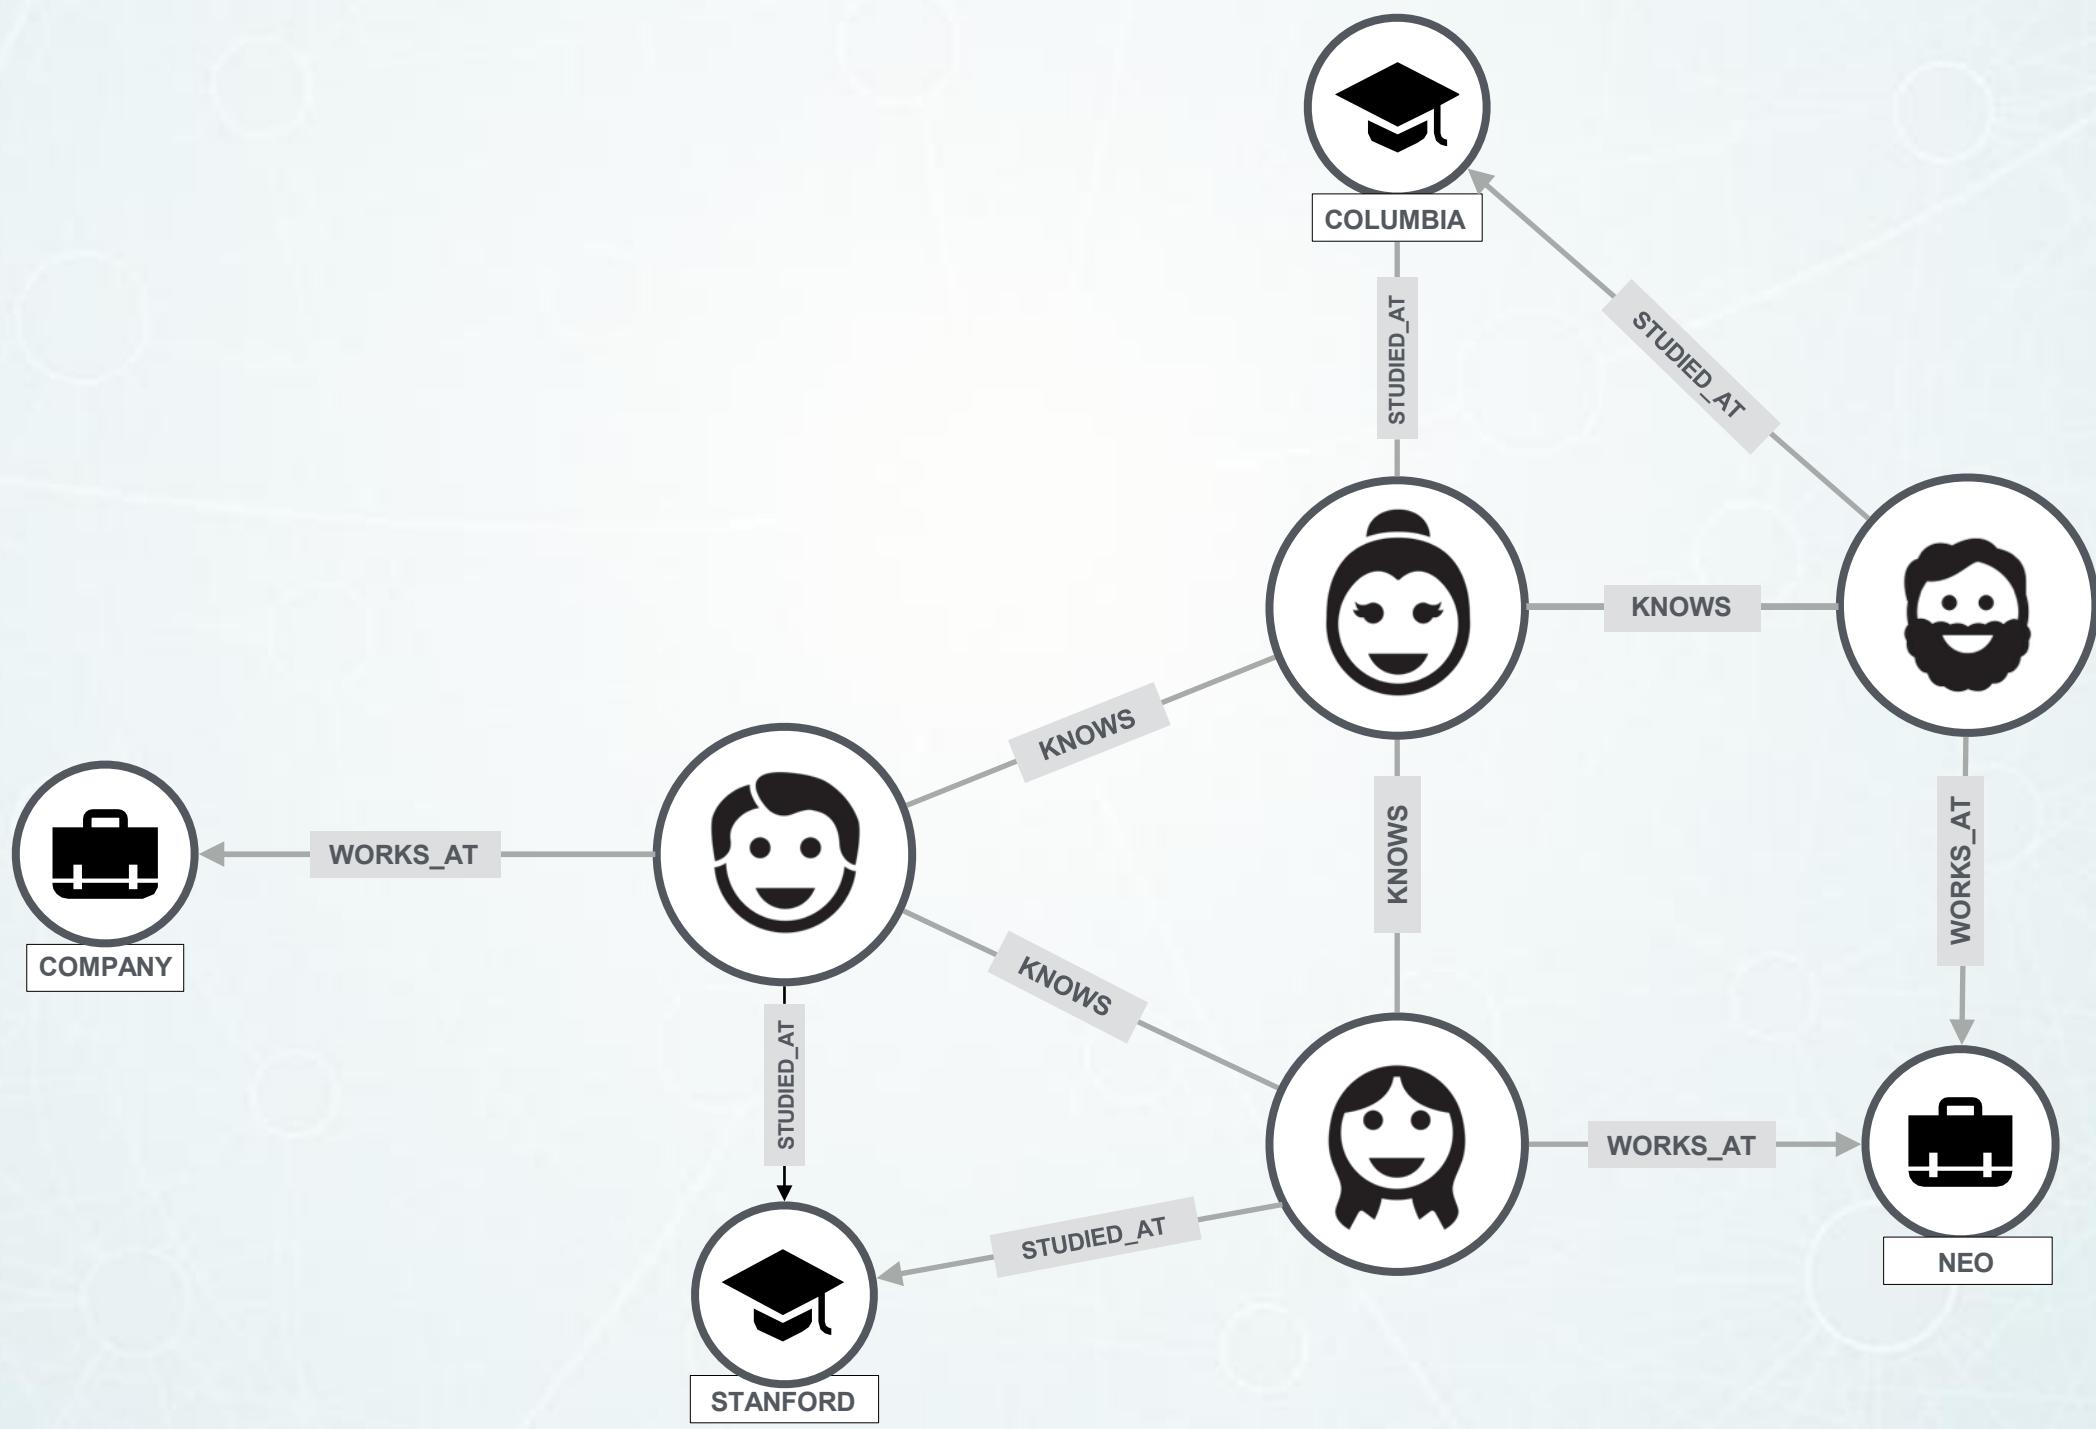
**A Graph Is Connected Data**



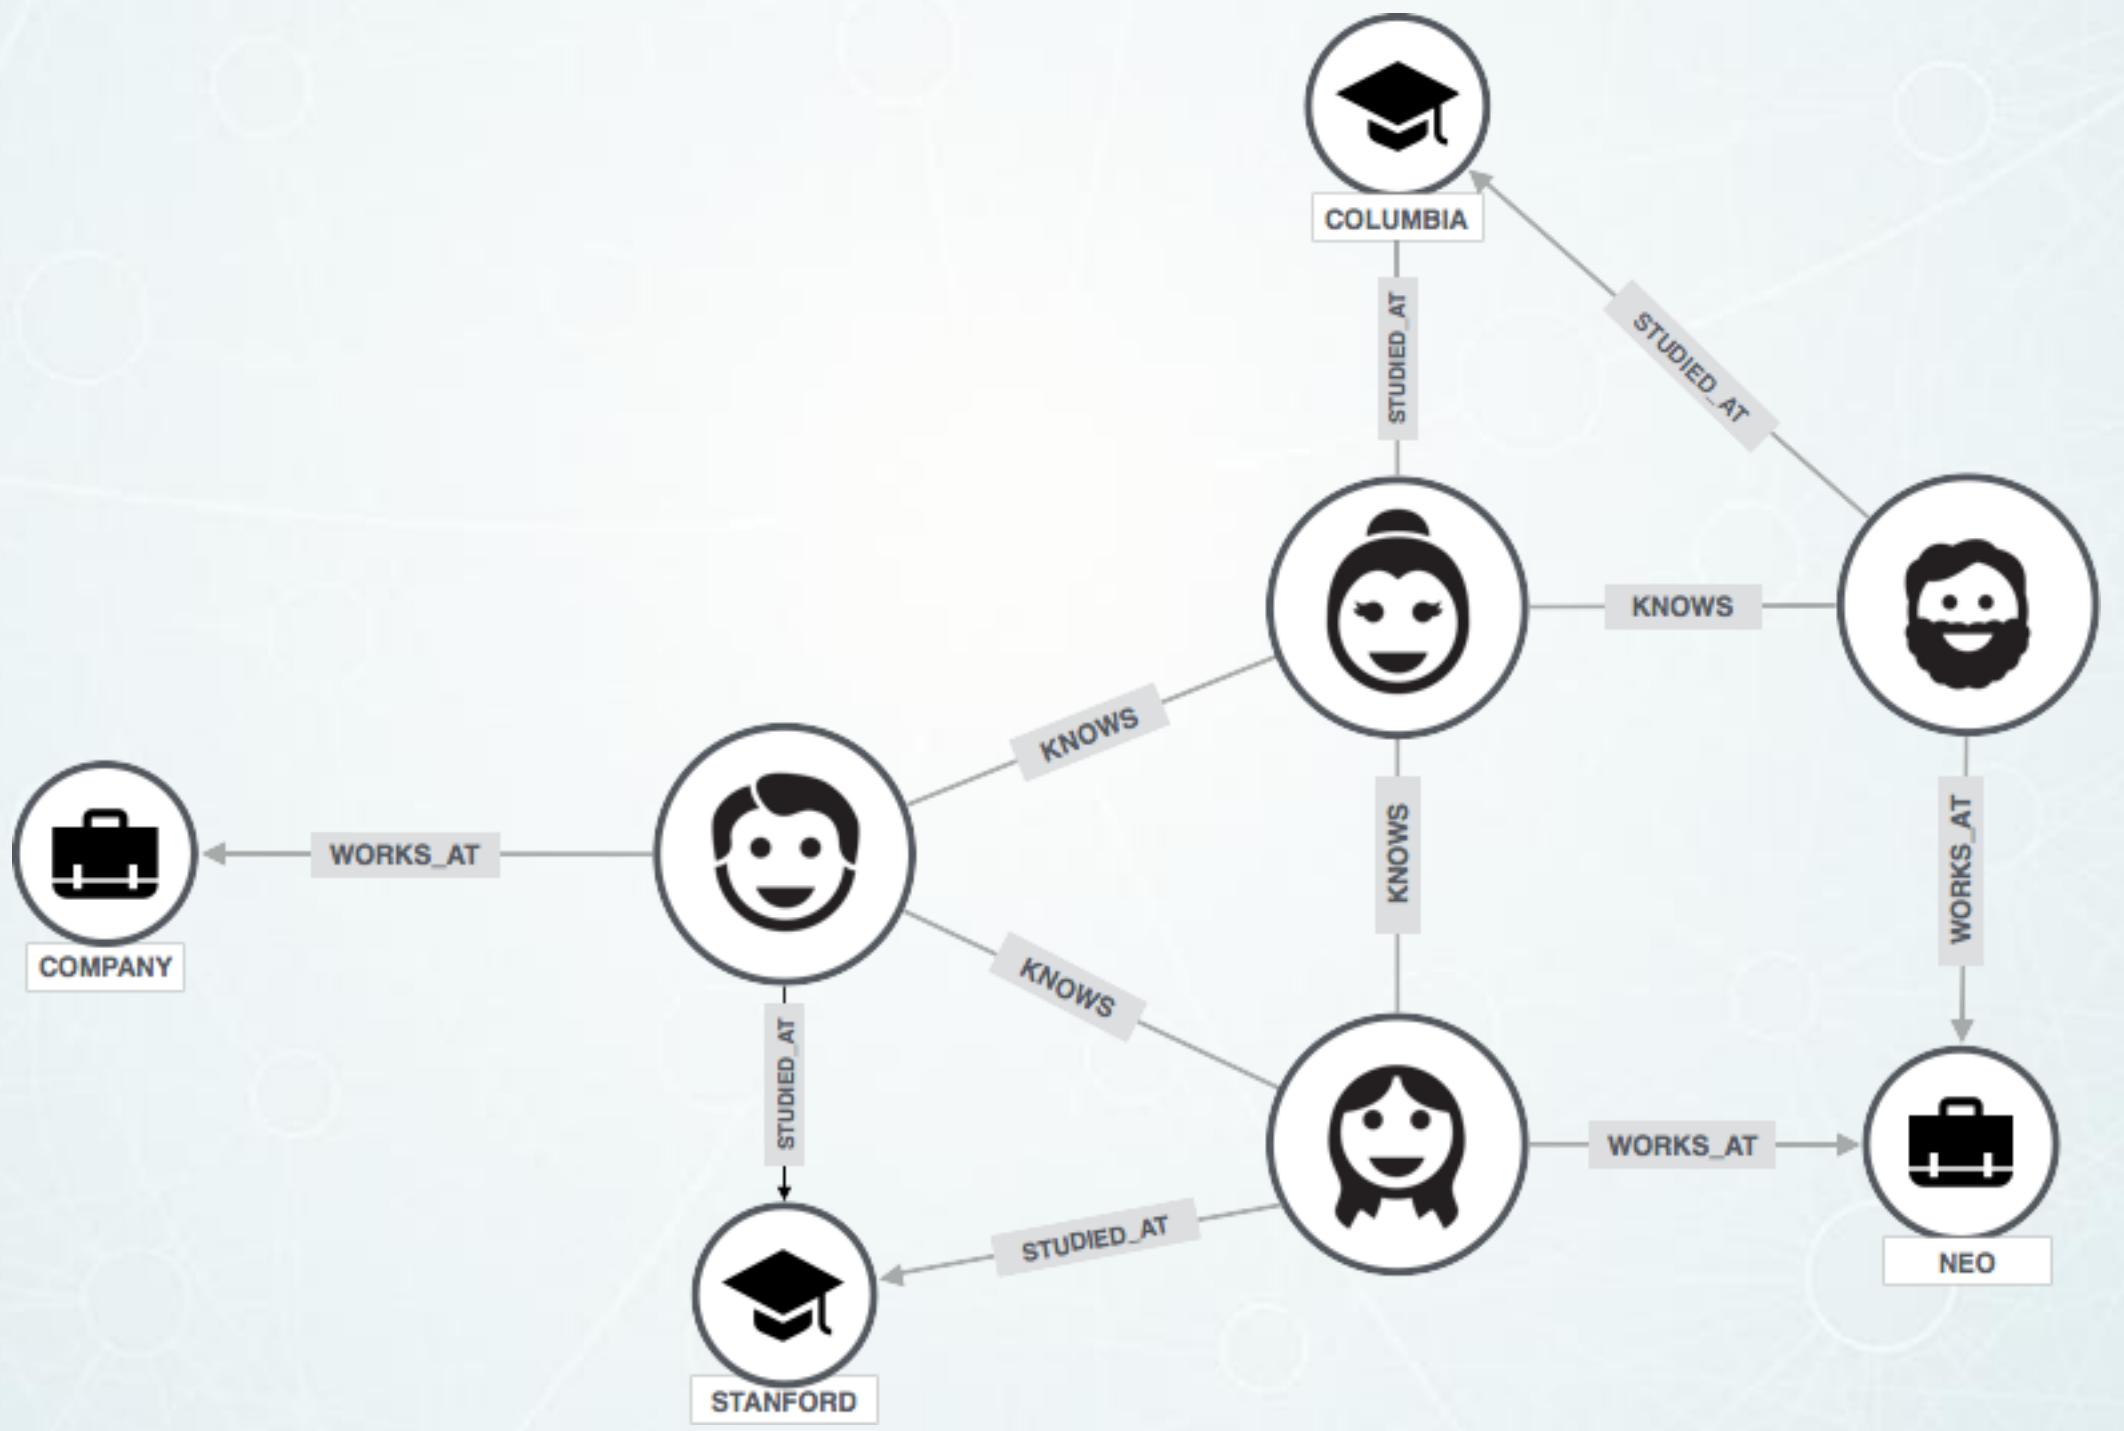
# A Graph Is Connected Data



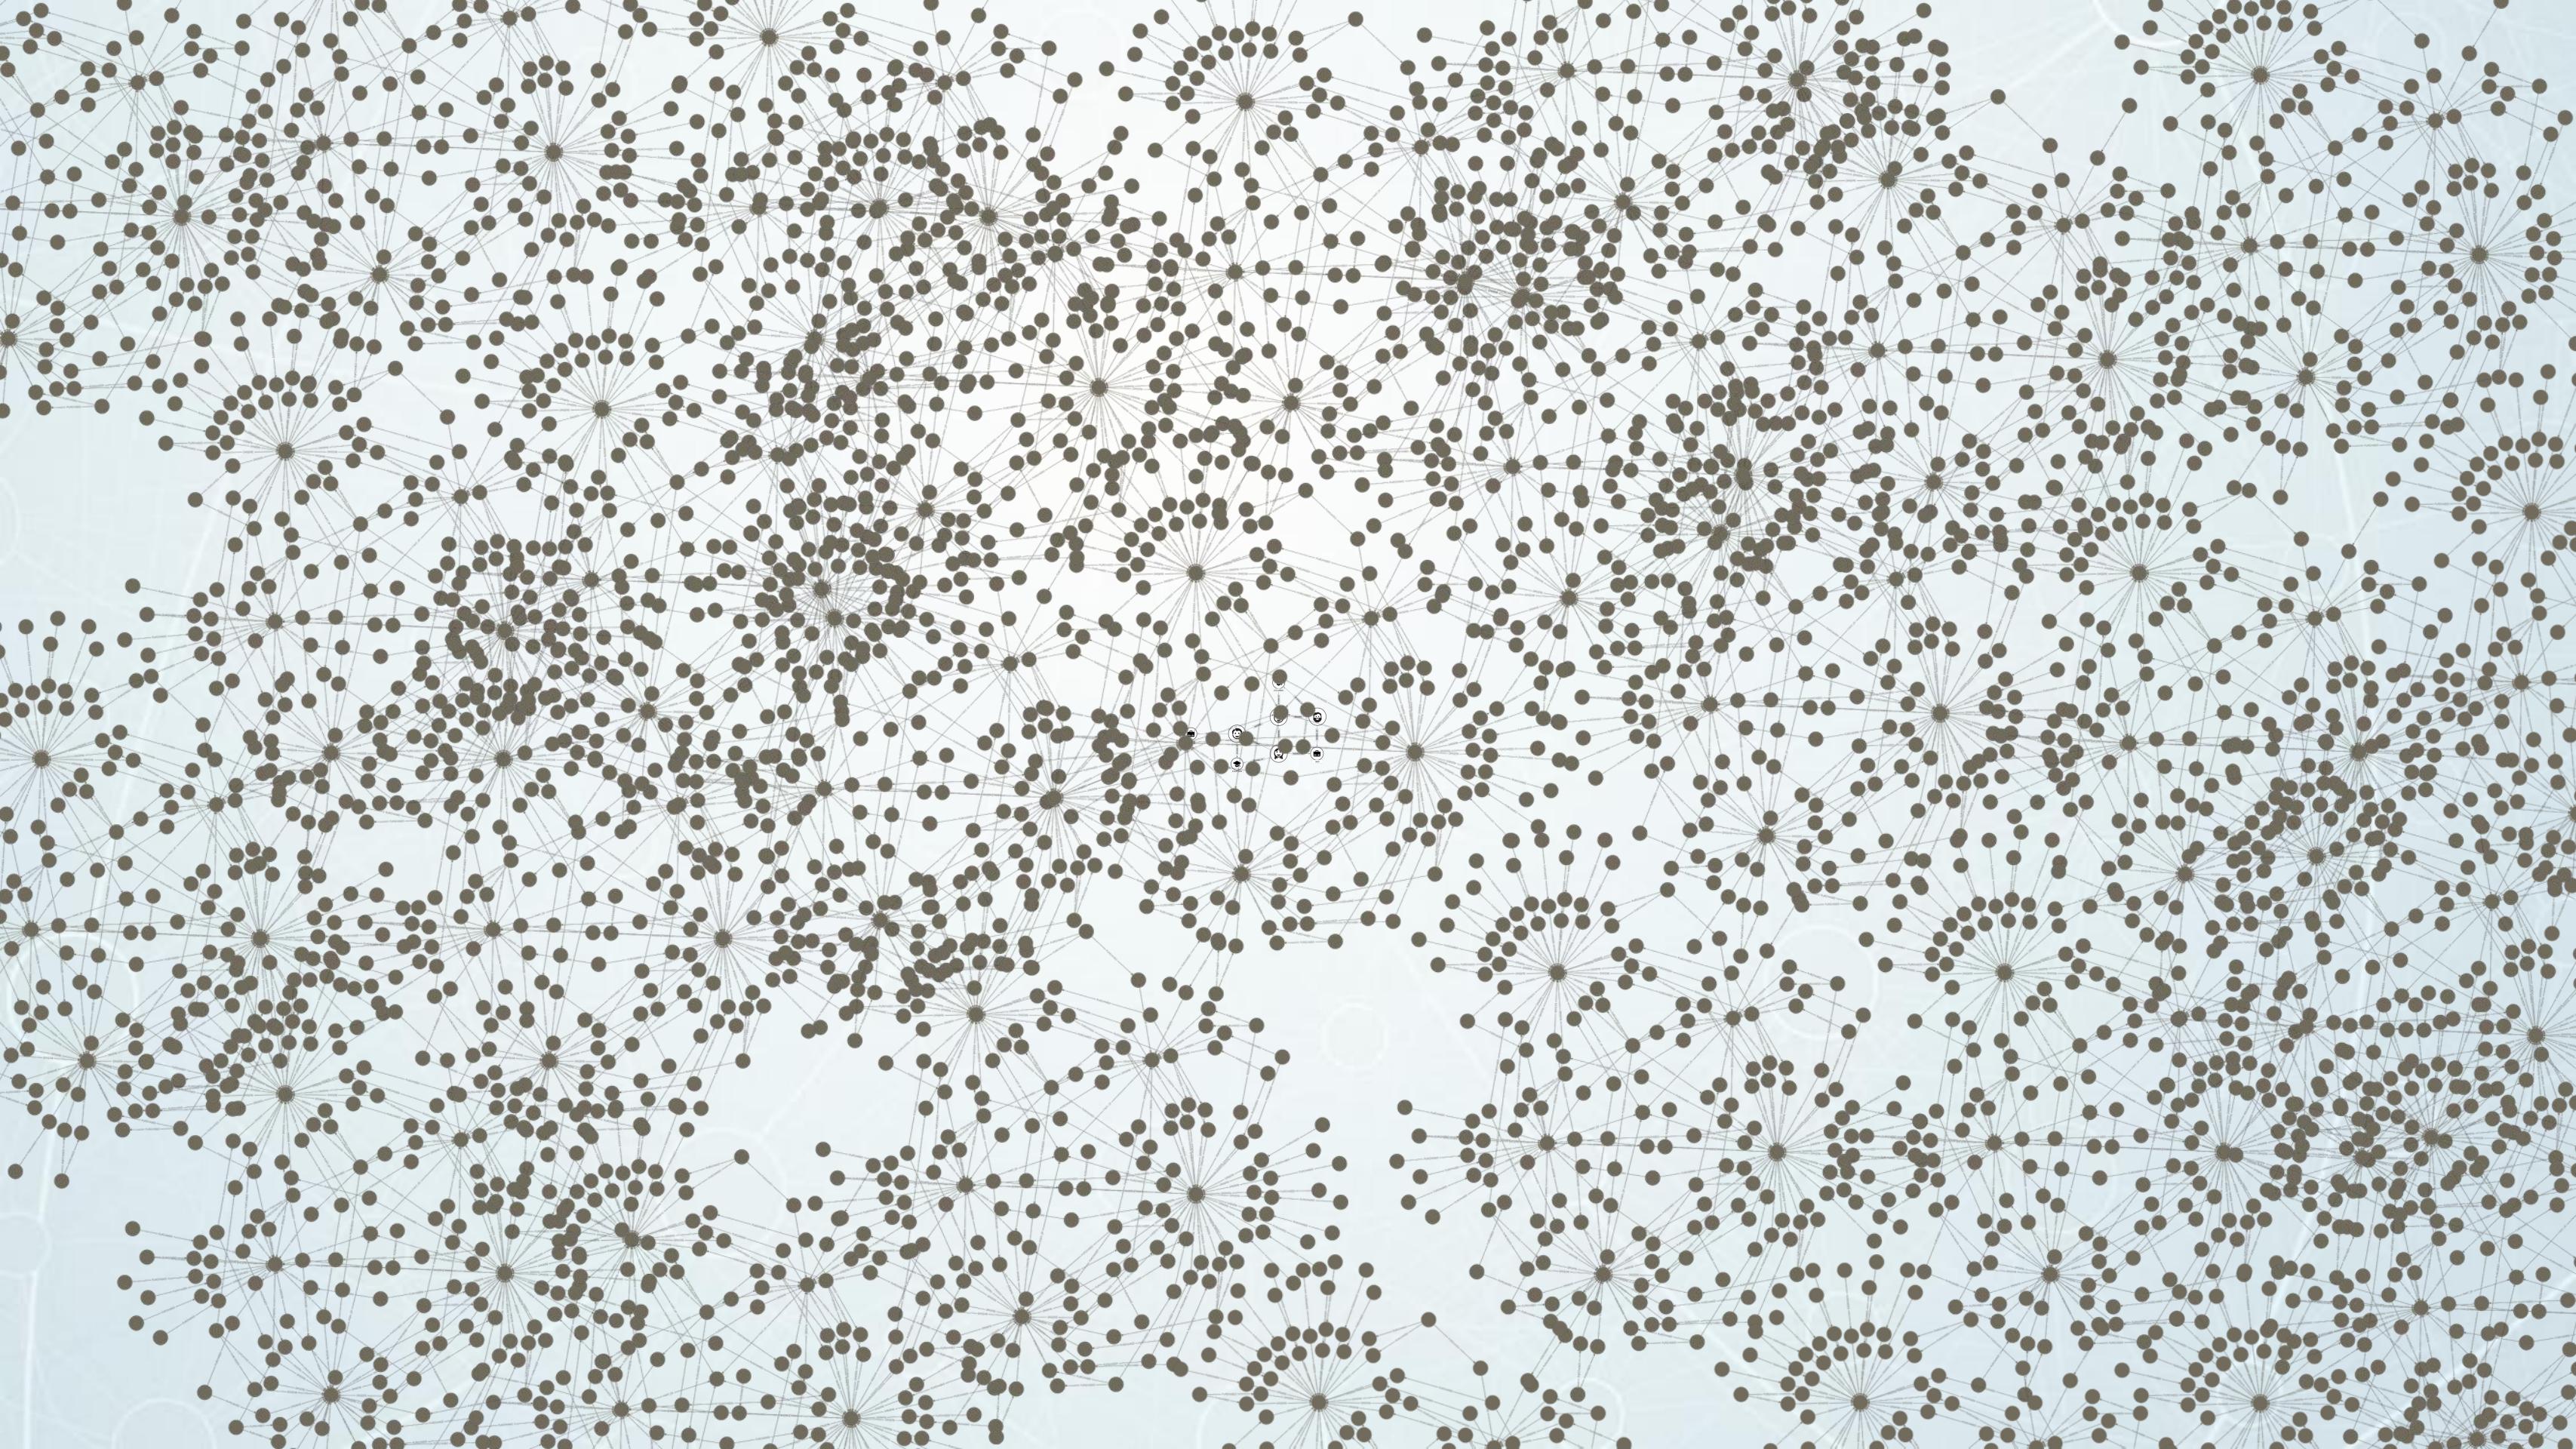
# A Graph Is Connected Data



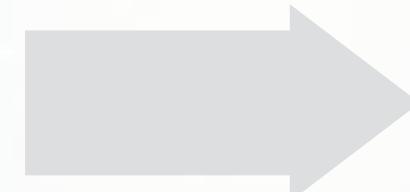
# A Graph Is Connected Data



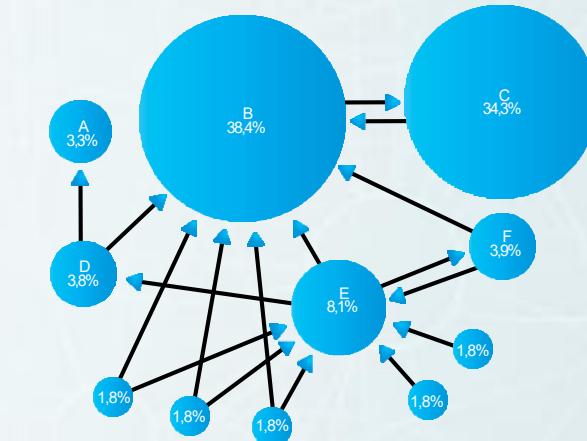
# A Graph Is Connected Data



# Use of Graphs has created some of the most successful companies in the world

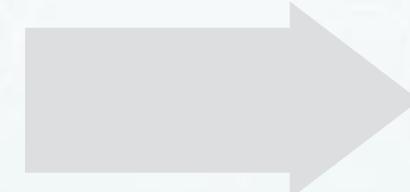


# Google

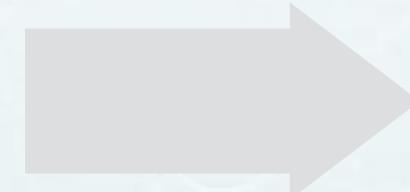
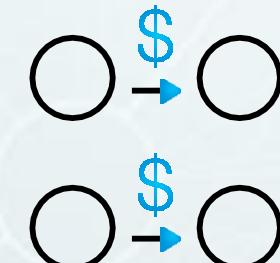
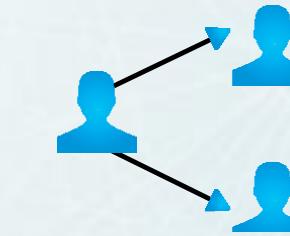


# monster

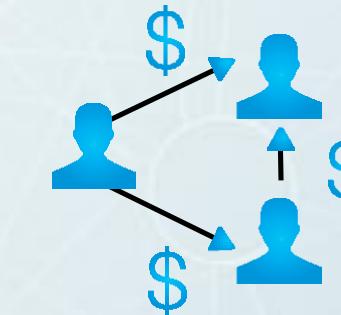
Find better.™



# LinkedIn



# PayPal





# Today we see graph-projects in virtually every industry



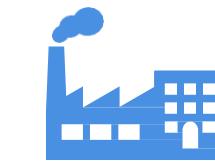
Finance



Social networks



HR &  
Recruiting



Manufacturing  
& Logistics



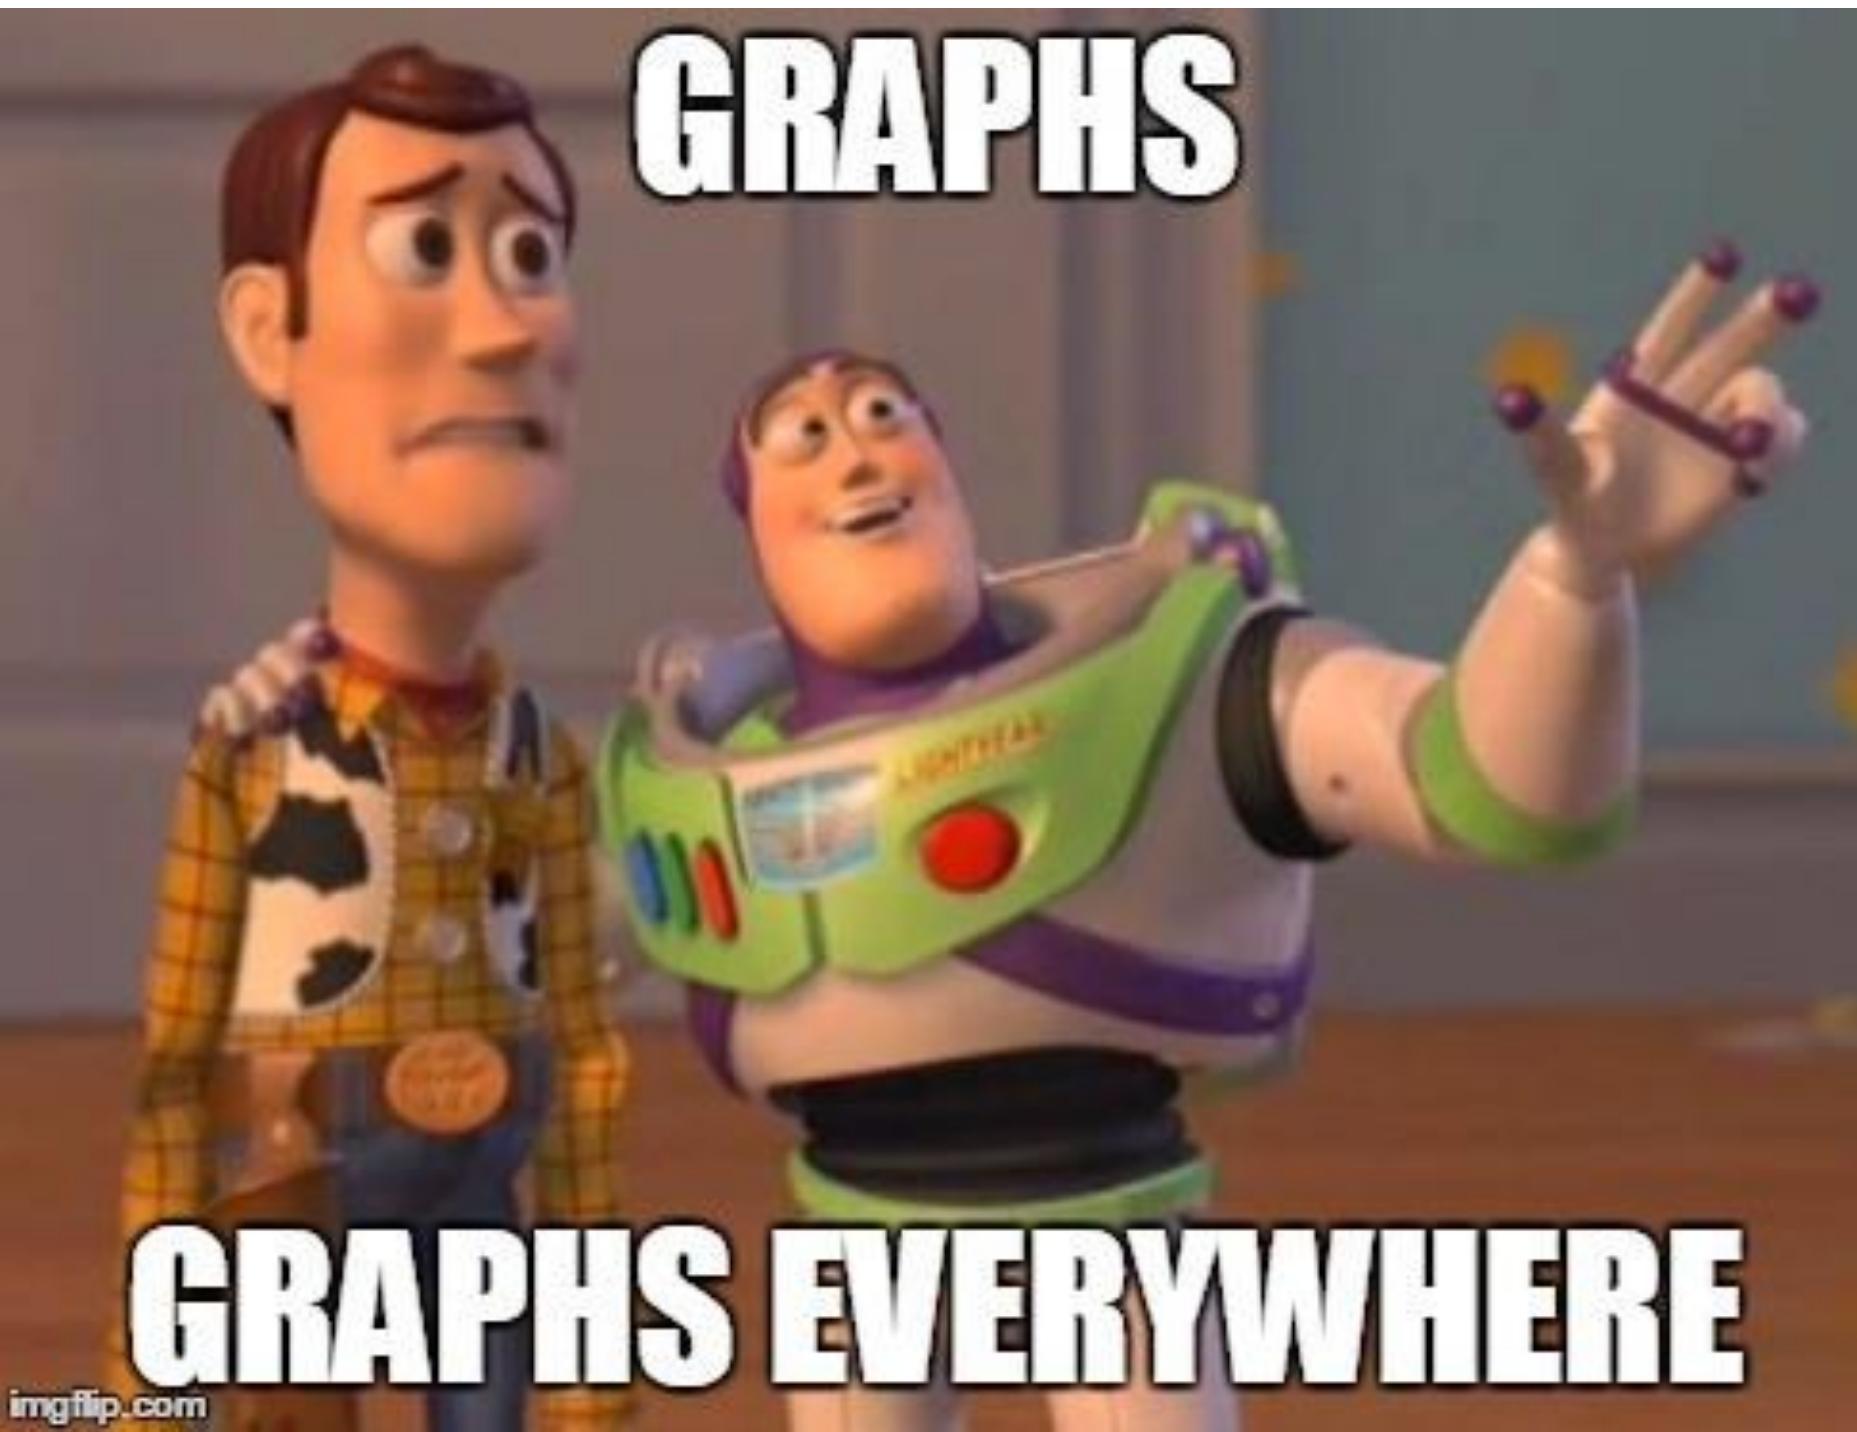
Health Care



Telco



Retail



# NEO4j USE CASES

Real Time Recommendations

Master Data Management

Fraud Detection

Graph Based Search

Network & IT-Operations

Identity & Access Management

# NEO4j USE CASES

Real Time Recommendations

Master Data Management

Fraud Detection

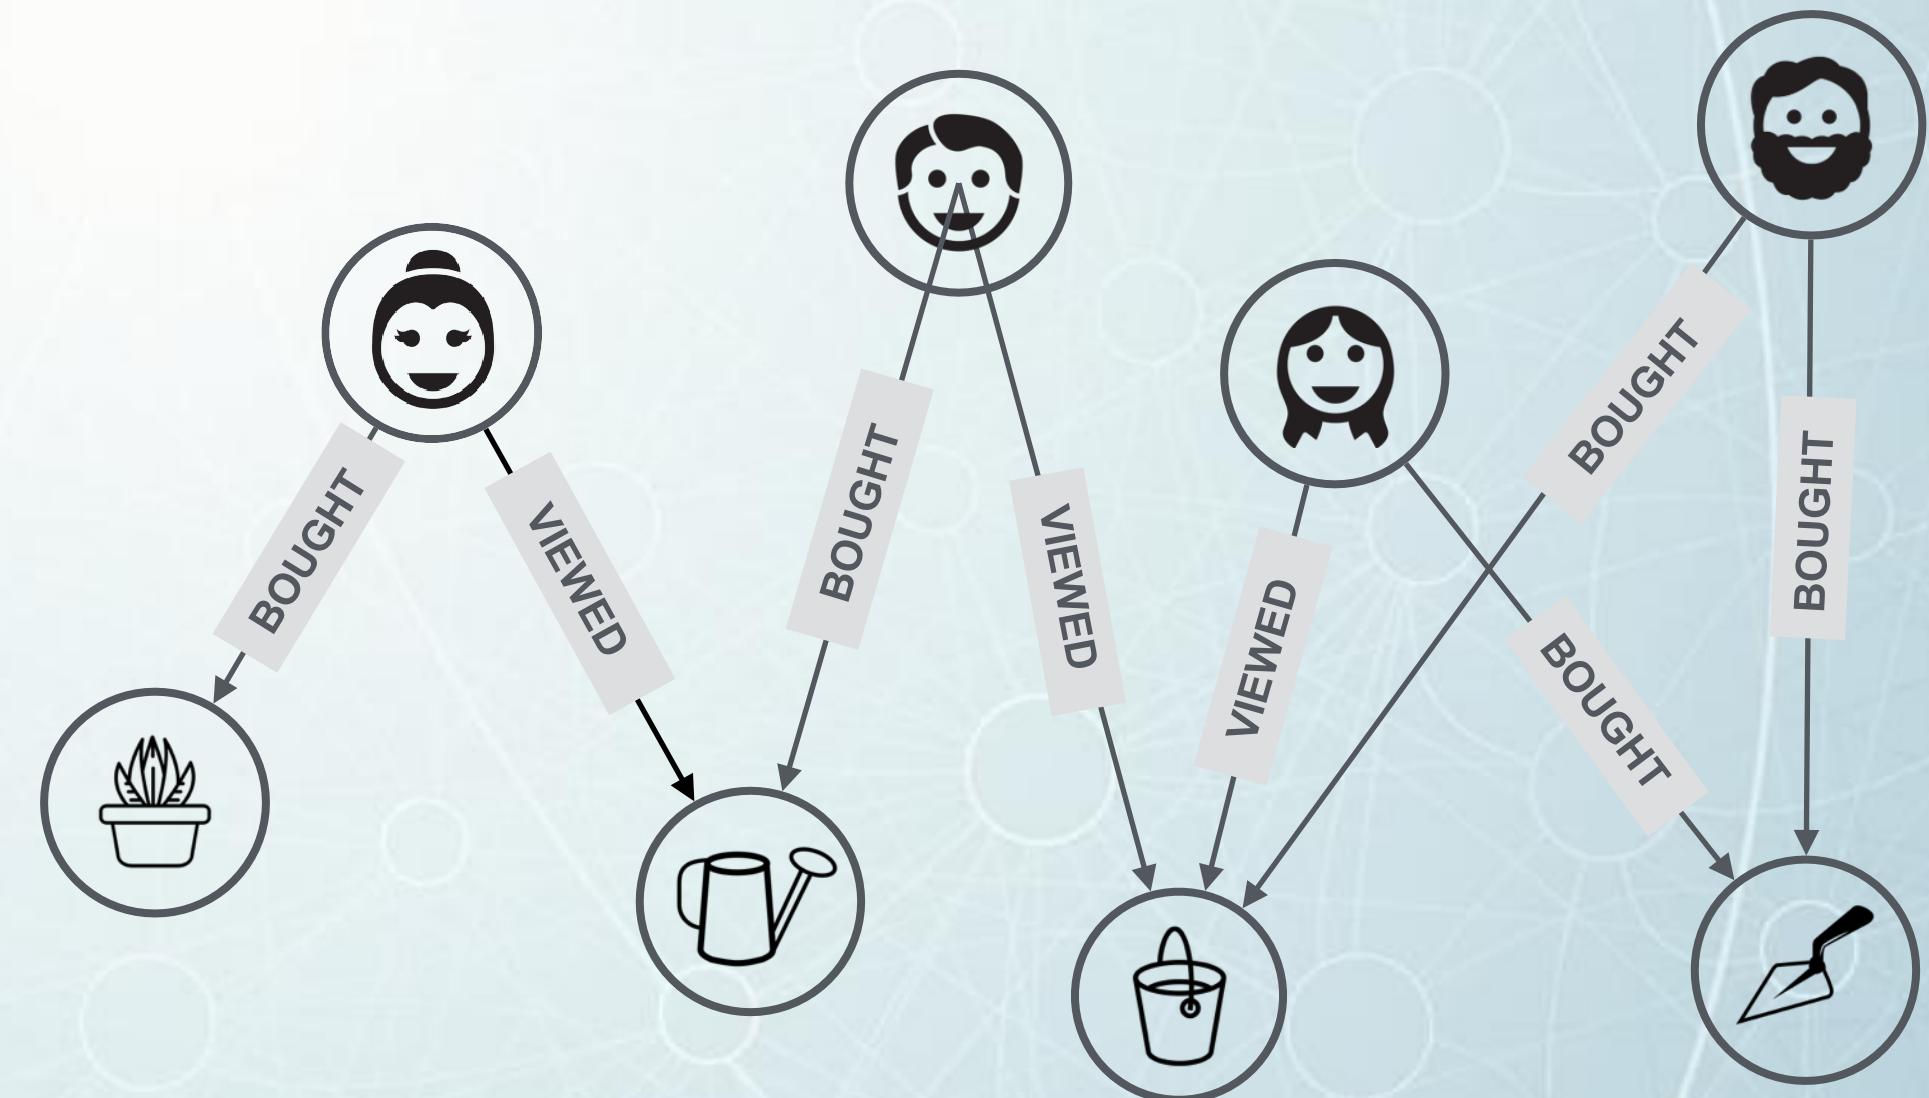
Graph Based Search

Network & IT-Operations

Identity & Access Management



## GRAPH THINKING: Real Time Recommendations





# NEO4j USE CASES

Real Time Recommendations

Master Data Management

Fraud Detection

Graph Based Search

Network & IT-Operations

Identity & Access Management



All

All Departments

Daily Savings Center

My Local Store

Tips & Ideas

Savings Catcher



FREE st



“As the current market leader in graph databases, and with enterprise features for scalability and availability, Neo4j is the right choice to meet our demands.”

Marcos Wada  
Software Developer, Walmart

Rollbacks

# NEO4j USE CASES

Real Time Recommendations

Master Data Management

Fraud Detection

Graph Based Search

Network & IT-Operations

Identity & Access Management



## GRAPH THINKING: Master Data Management



# NEO4j USE CASES

Real Time Recommendations

Master Data Management

Fraud Detection

Graph Based Search

Network & IT-Operations

Identity & Access Management



Neo4j is the heart of Cisco HMP: used for governance and single source of truth and a one-stop shop for all of Cisco's hierarchies.

# NEO4j USE CASES

Real Time Recommendations

Master Data Management

Fraud Detection

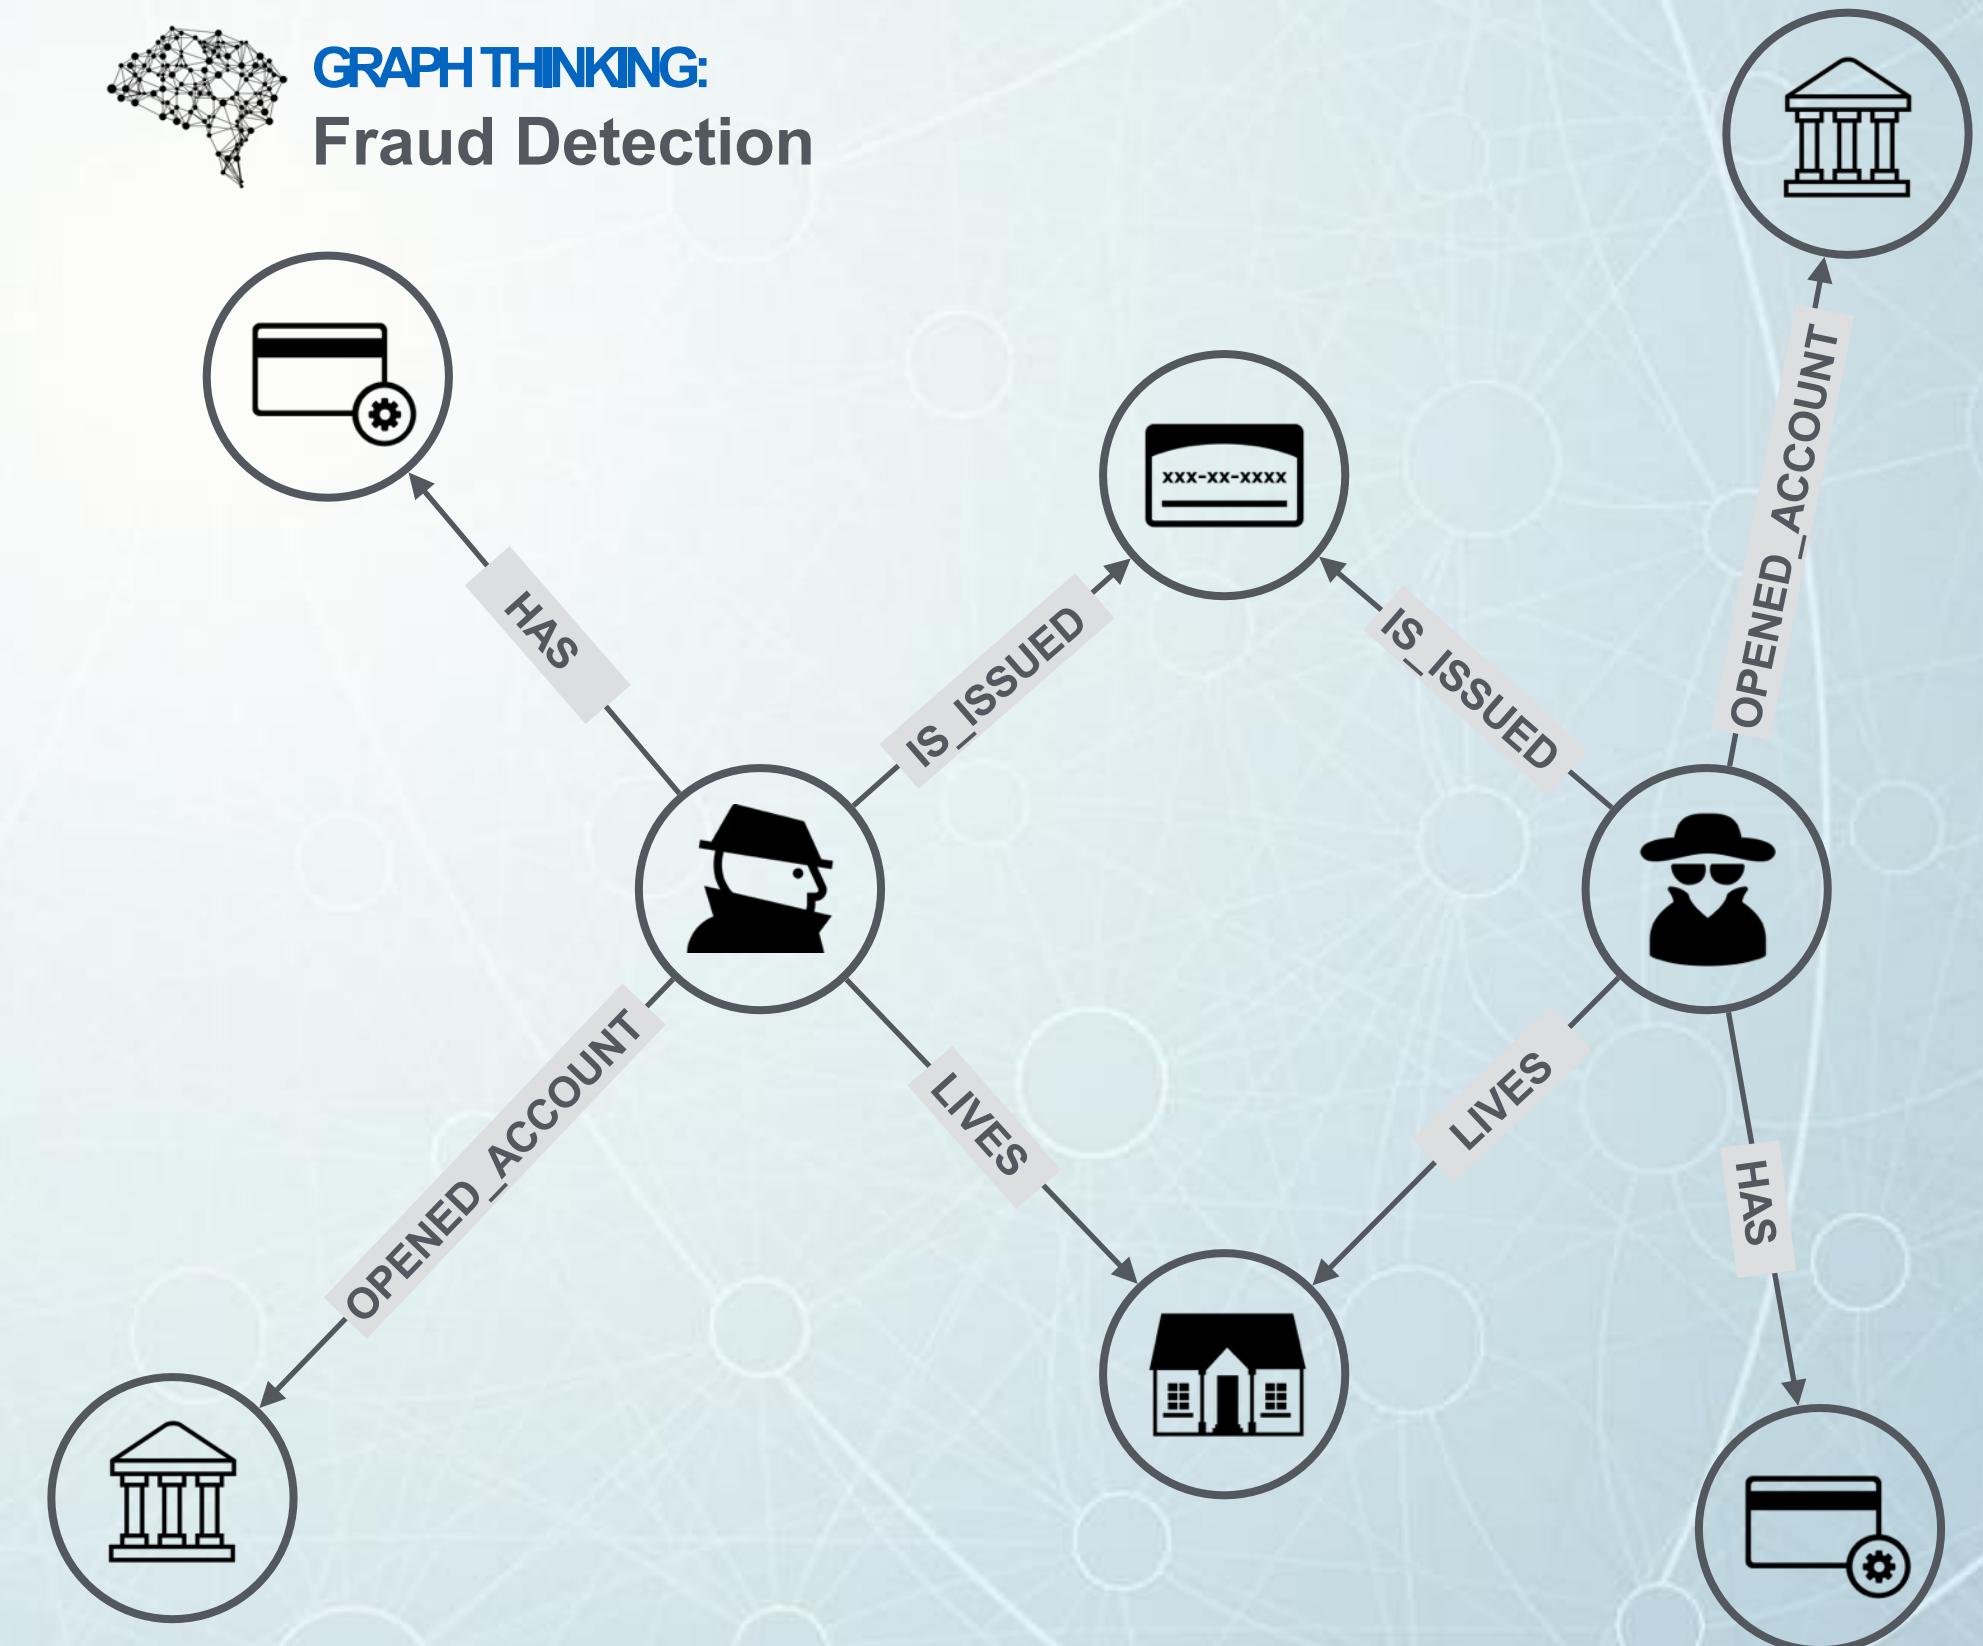
Graph Based Search

Network & IT-Operations

Identity & Access Management



## GRAPH THINKING: Fraud Detection



# NEO4j USE CASES

Real Time Recommendations

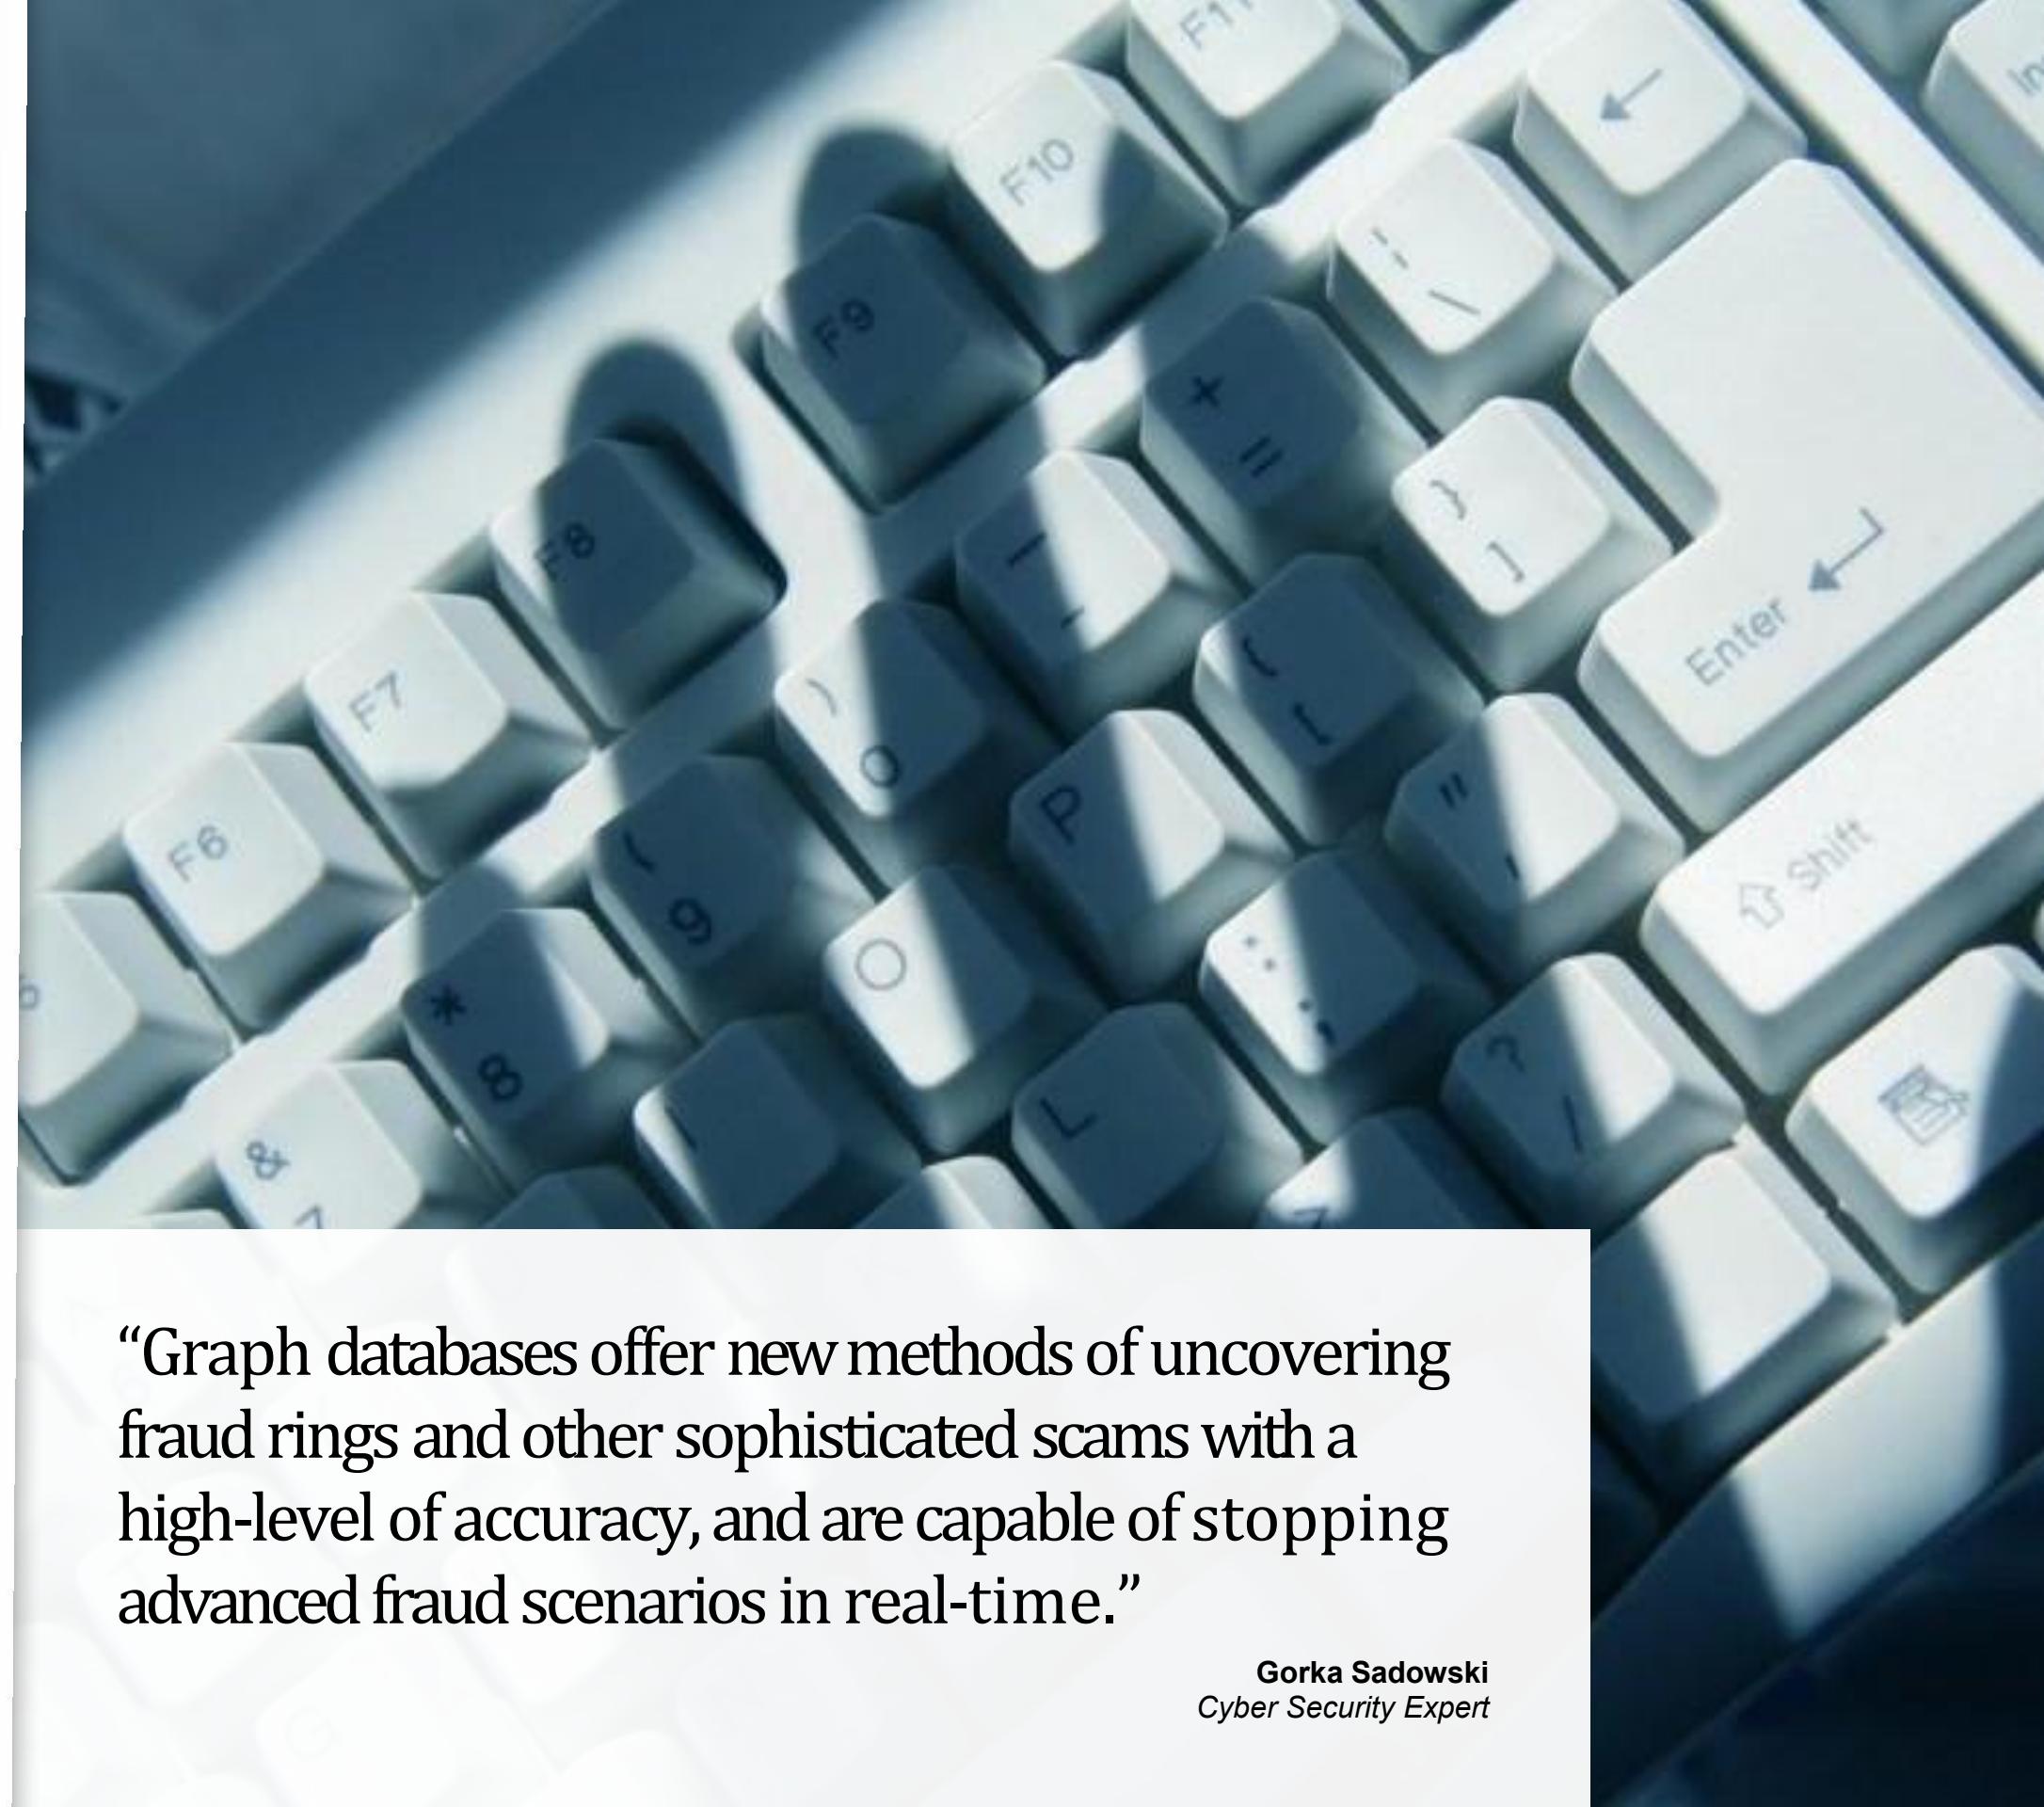
Master Data Management

Fraud Detection

Graph Based Search

Network & IT-Operations

Identity & Access Management



“Graph databases offer new methods of uncovering fraud rings and other sophisticated scams with a high-level of accuracy, and are capable of stopping advanced fraud scenarios in real-time.”

Gorka Sadowski  
*Cyber Security Expert*

# NEO4j USE CASES

Real Time Recommendations

Master Data Management

Fraud Detection

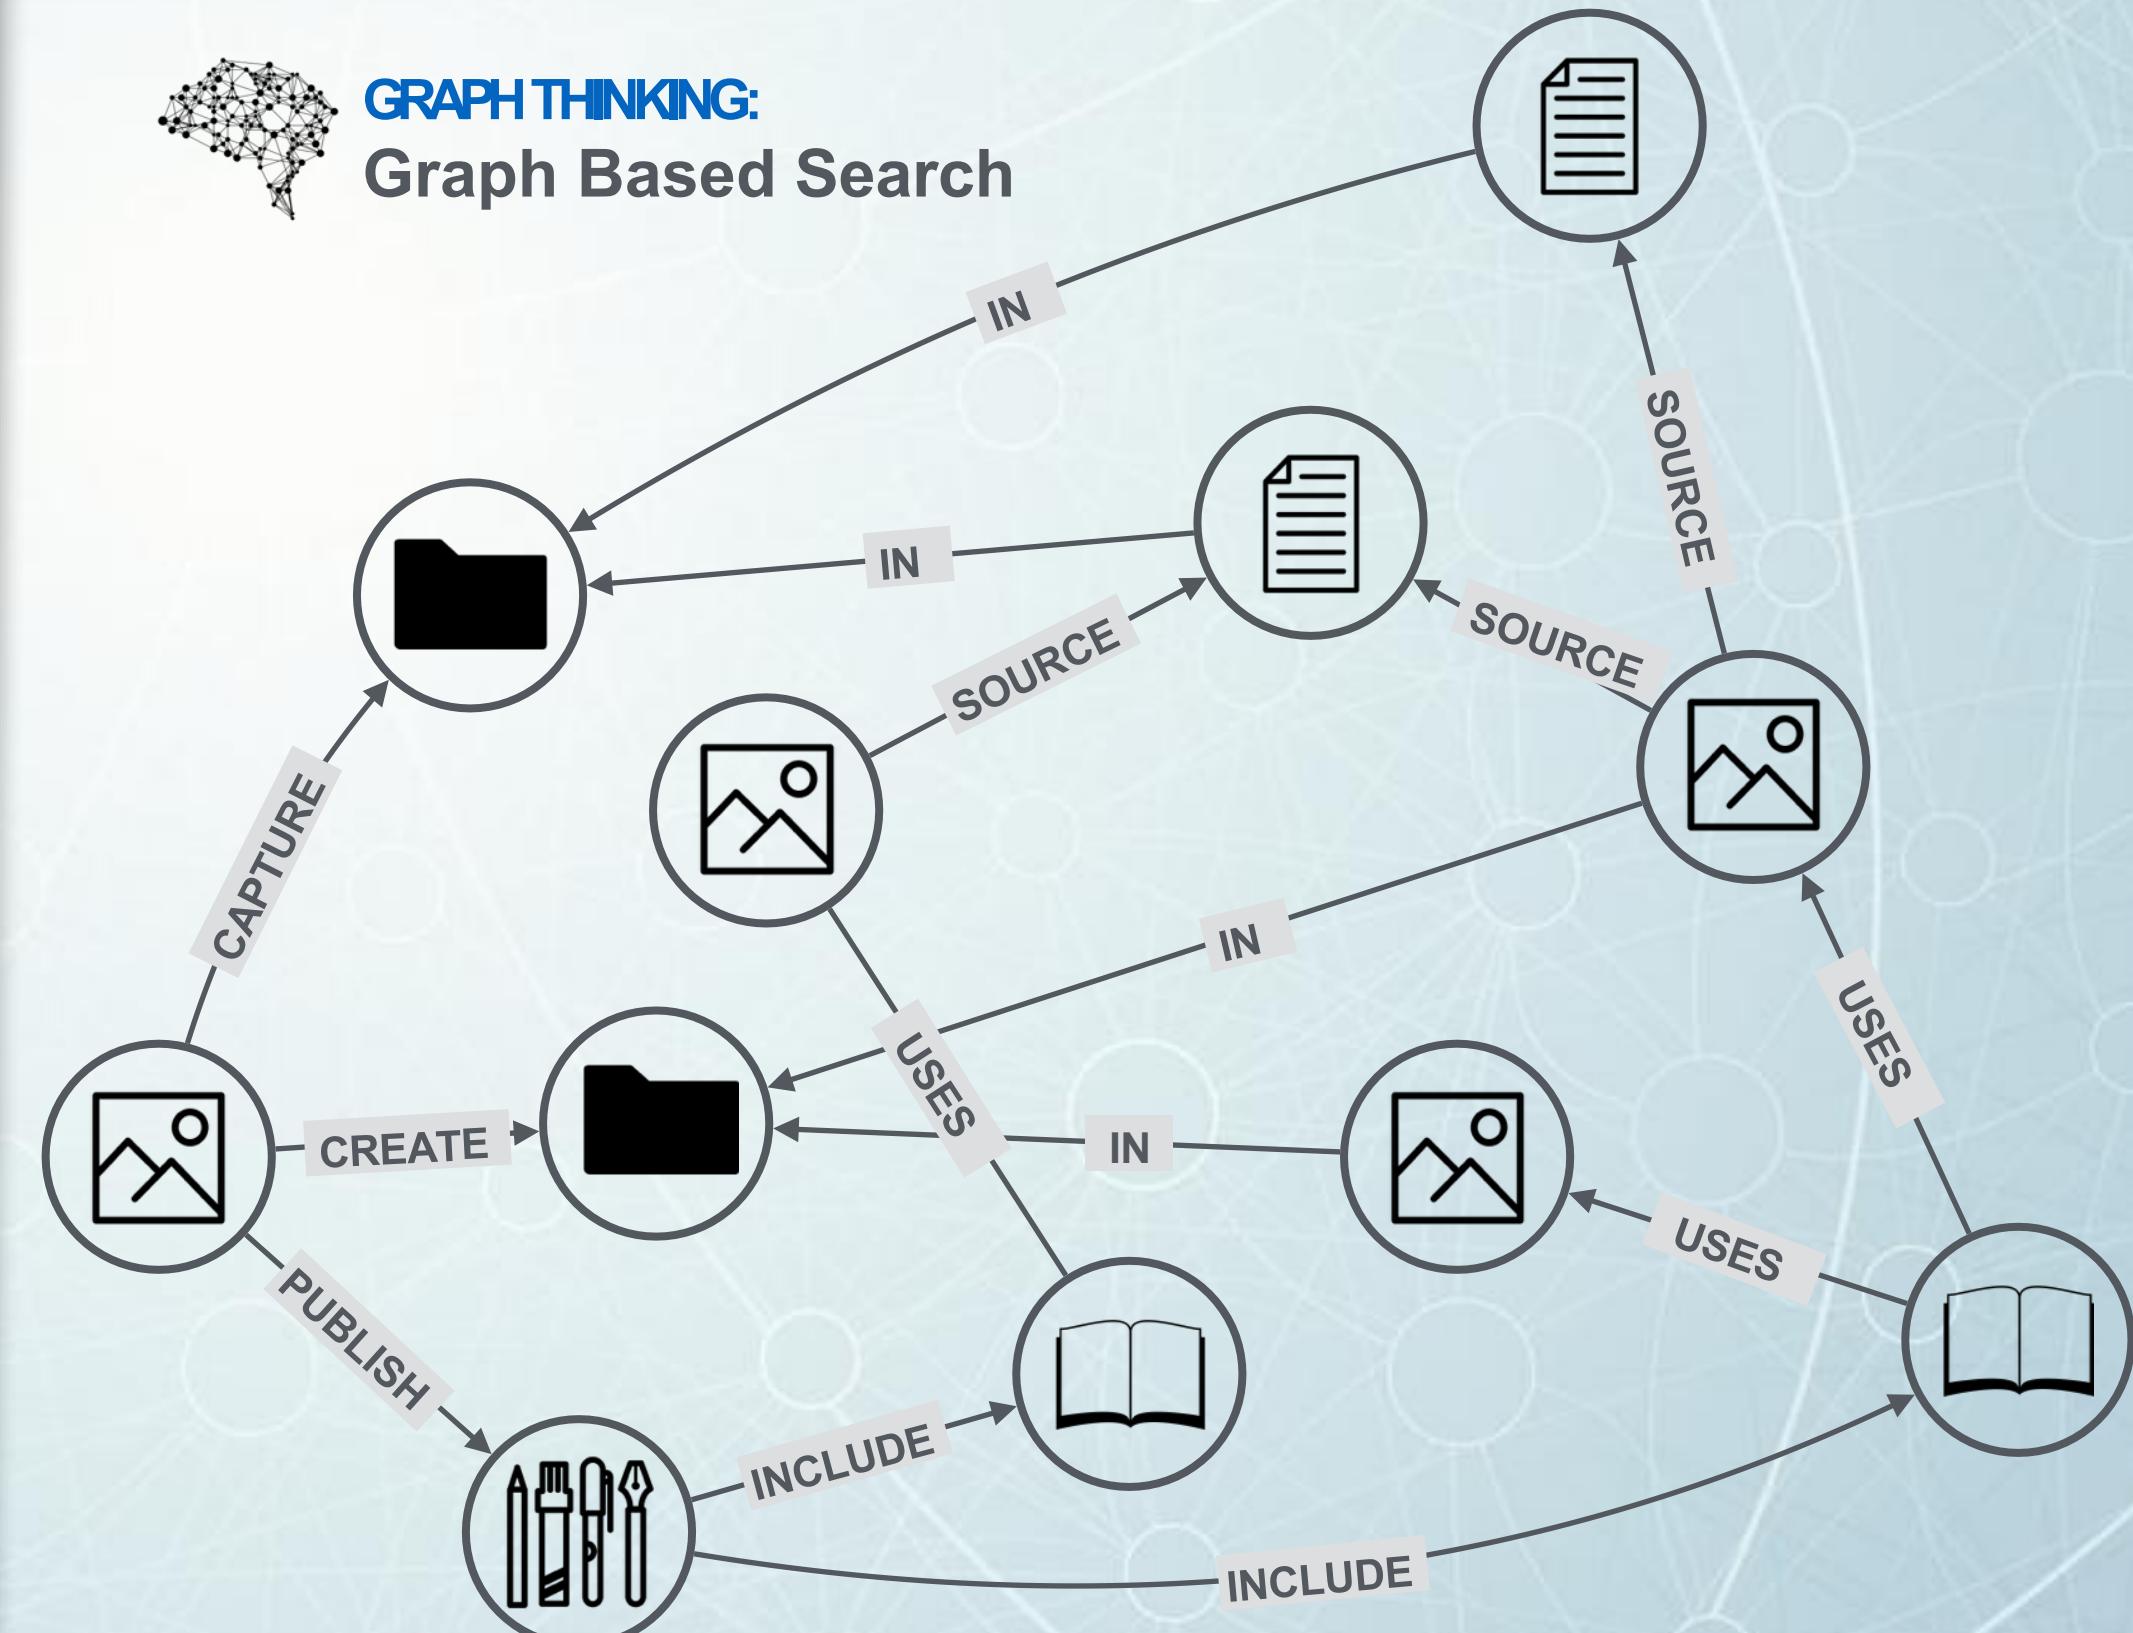
Graph Based Search

Network & IT-Operations

Identity & Access Management



## GRAPH THINKING: Graph Based Search



# NEO4j USE CASES

Real Time Recommendations

Master Data Management

Fraud Detection

Graph Based Search

Network & IT-Operations

Identity & Access Management



**Lufthansa**

Uses Neo4j to manage the digital assets inside of its next generation in-flight entertainment system.

# NEO4j USE CASES

Real Time Recommendations

Master Data Management

Fraud Detection

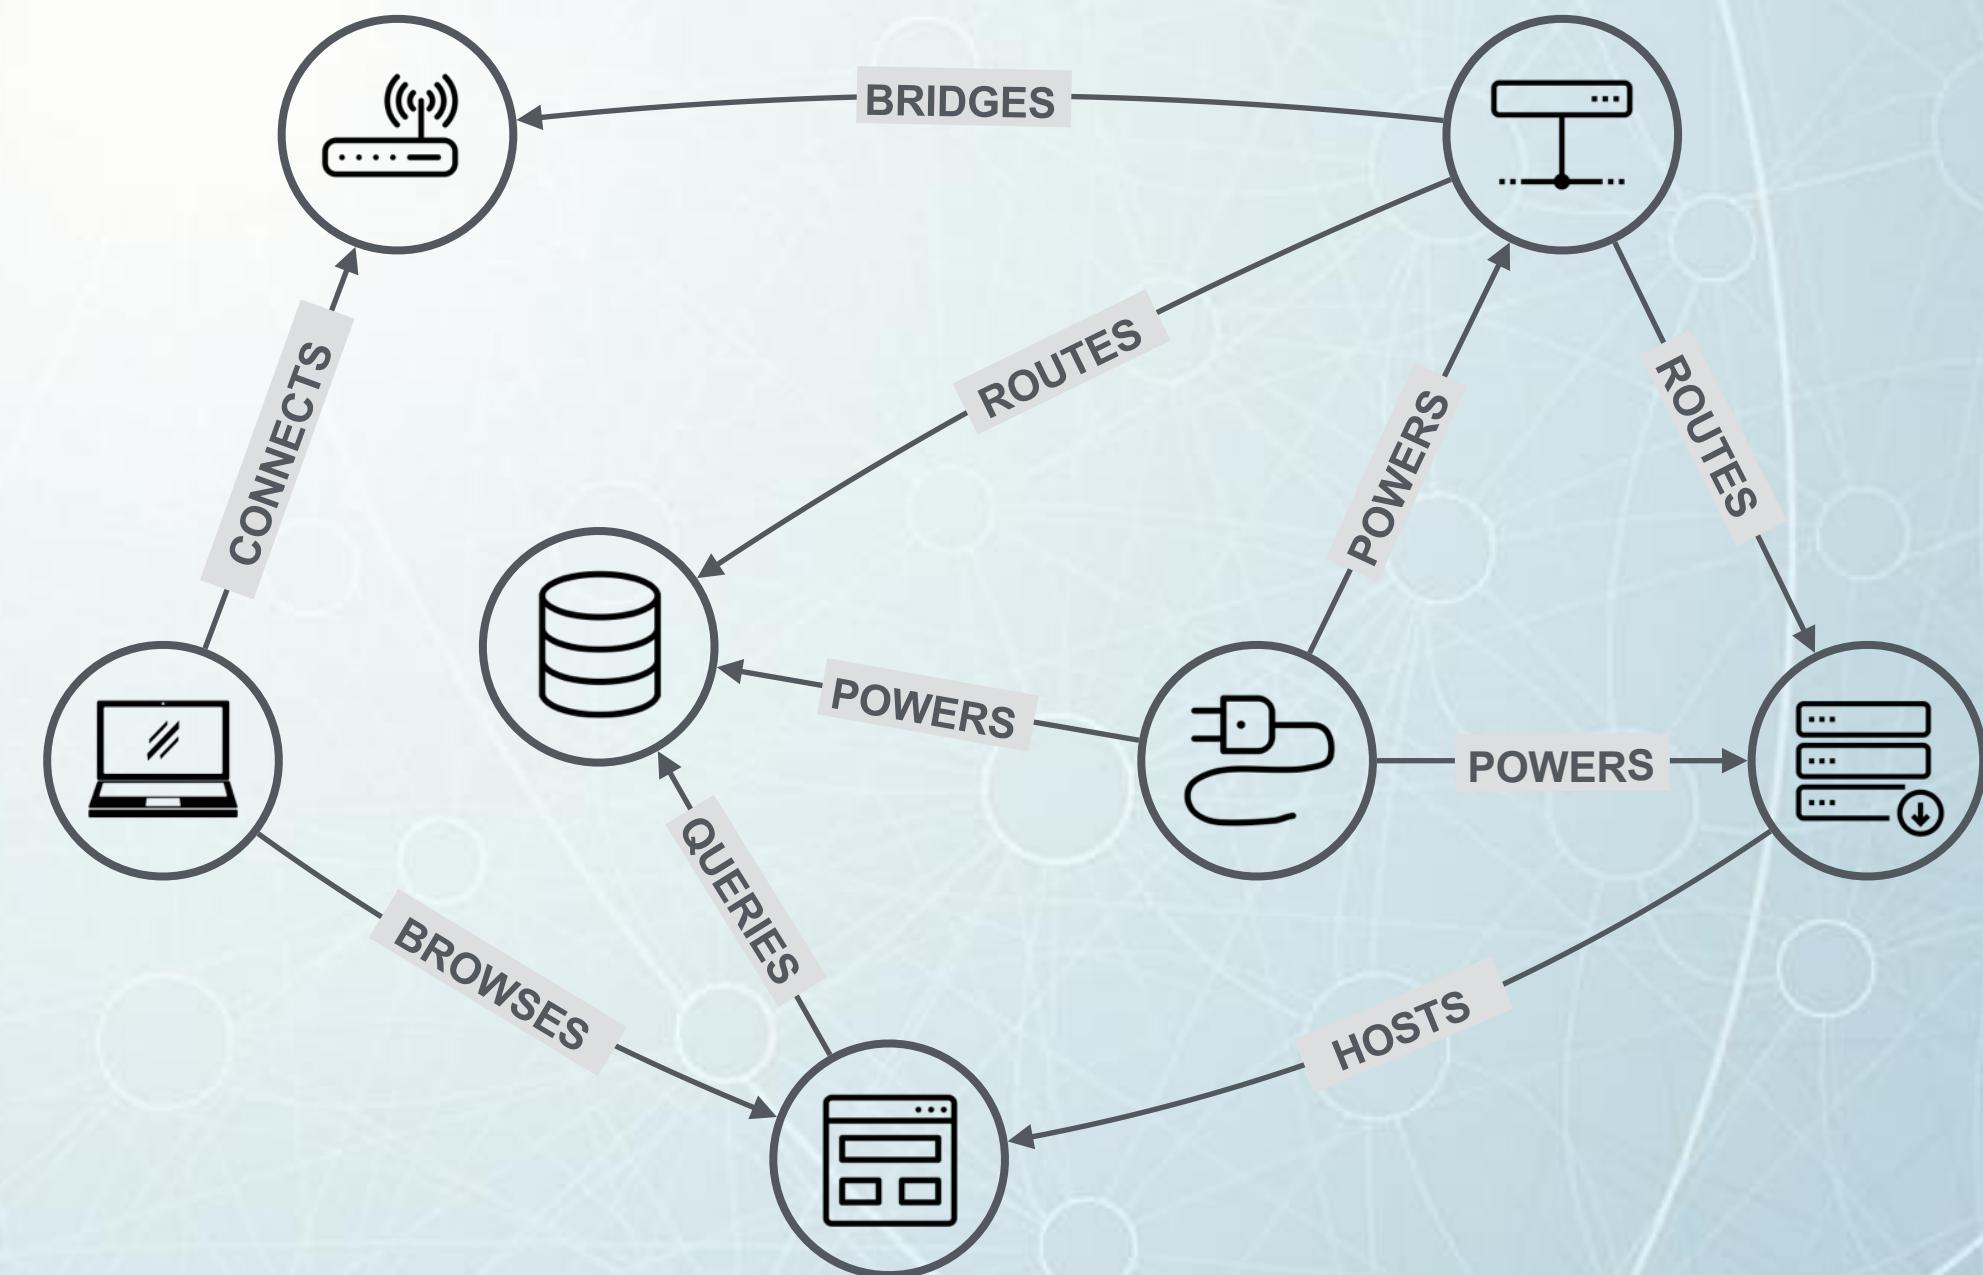
Graph Based Search

Network & IT-Operations

Identity & Access Management



## GRAPH THINKING: Network & IT-Operations



# NEO4j USE CASES

Real Time Recommendations

Master Data Management

Fraud Detection

Graph Based Search

Network & IT-Operations

Identity & Access Management



Uses Neo4j for network topology analysis  
for big telco service providers

# NEO4j USE CASES

Real Time Recommendations

Master Data Management

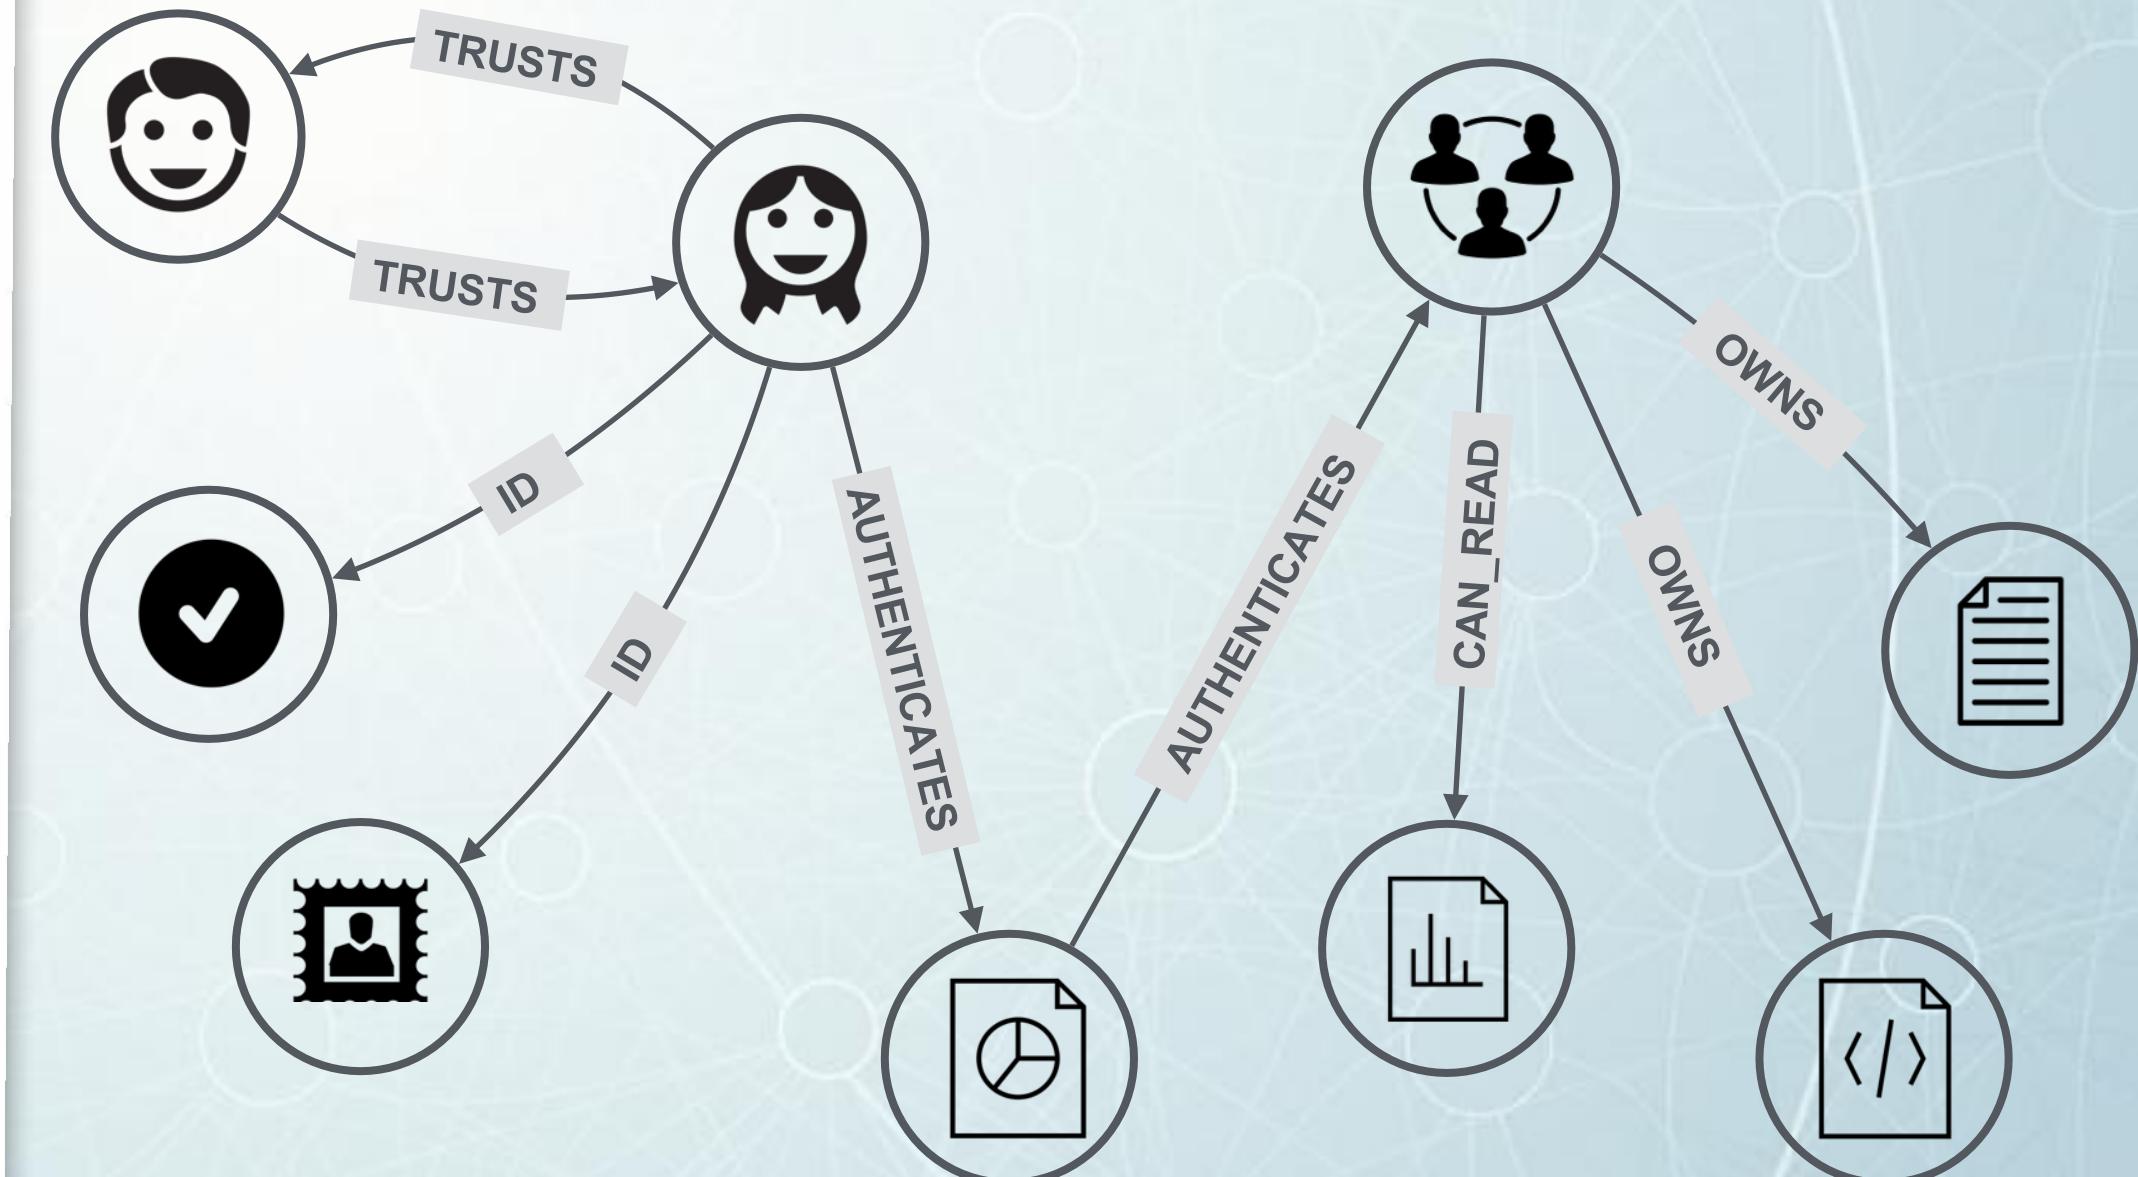
Fraud Detection

Graph Based Search

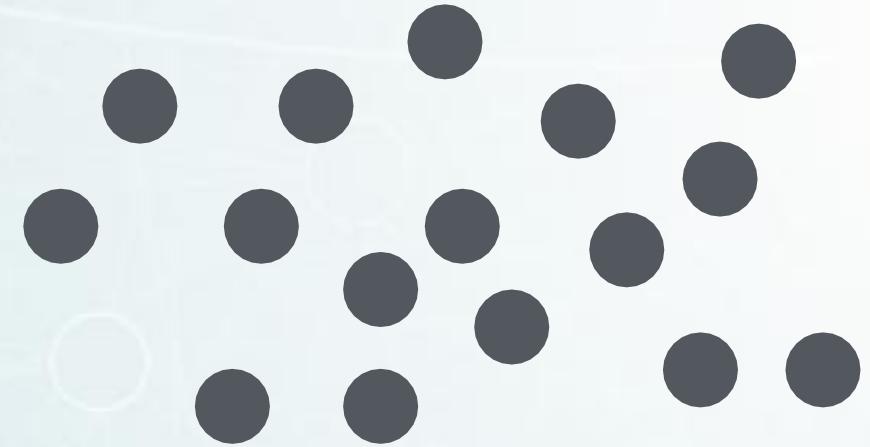
Network & IT-Operations



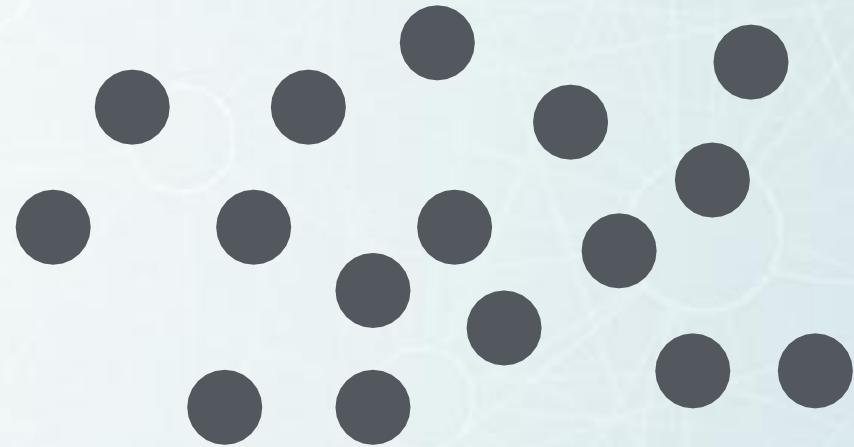
## GRAPH THINKING: Identity And Access Management



# A way of representing data

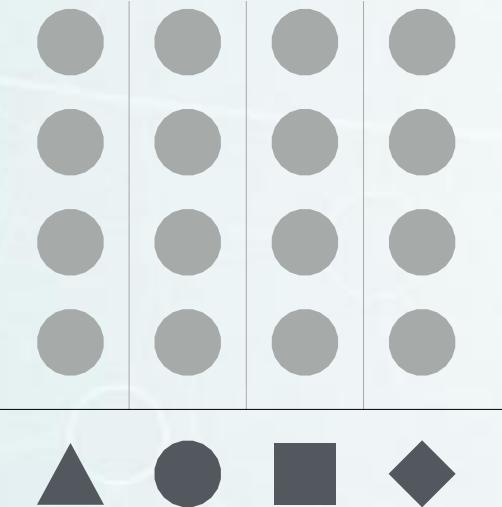


**DATA**

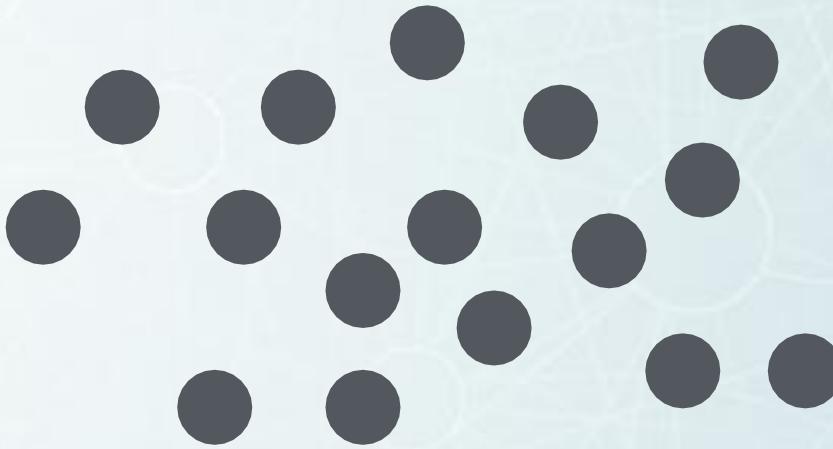


**DATA**

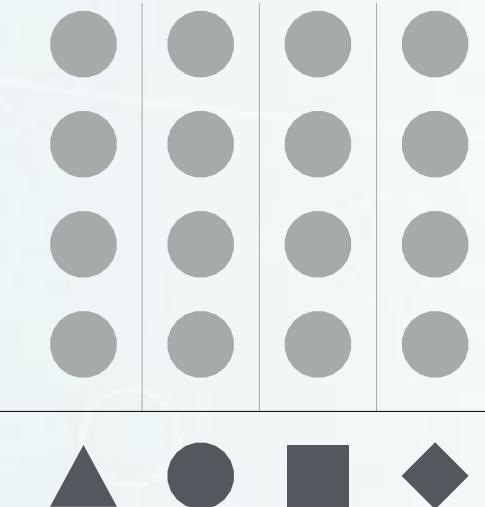
# A way of representing data



Relational Database



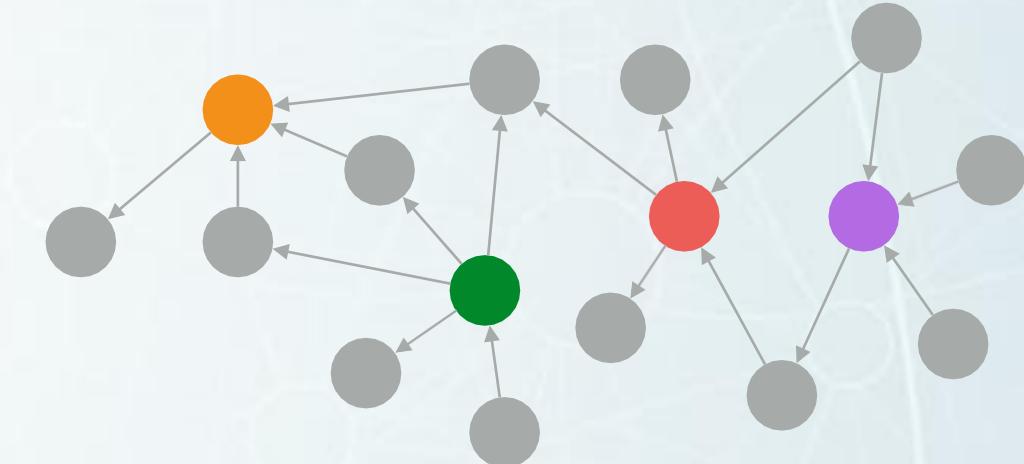
# A way of representing data



Relational Database



Graph Database



## Good for:

- Well-understood data structures that don't change too frequently
- Known problems involving discrete parts of the data, or minimal connectivity

## Good for:

- Dynamic systems: where the data topology is difficult to predict
- Dynamic requirements: that evolve with the business
- Problems where the relationships in data contribute meaning & value

# **THE PROPERTY GRAPH DATA MODEL**



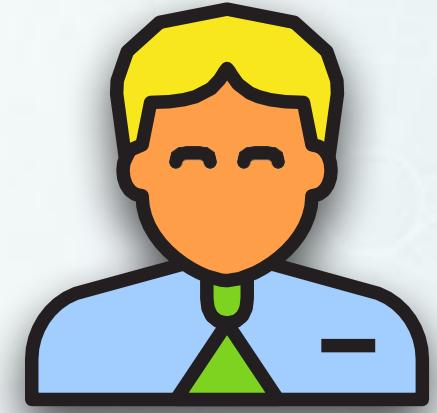
# Ann Loves Dan



Ann

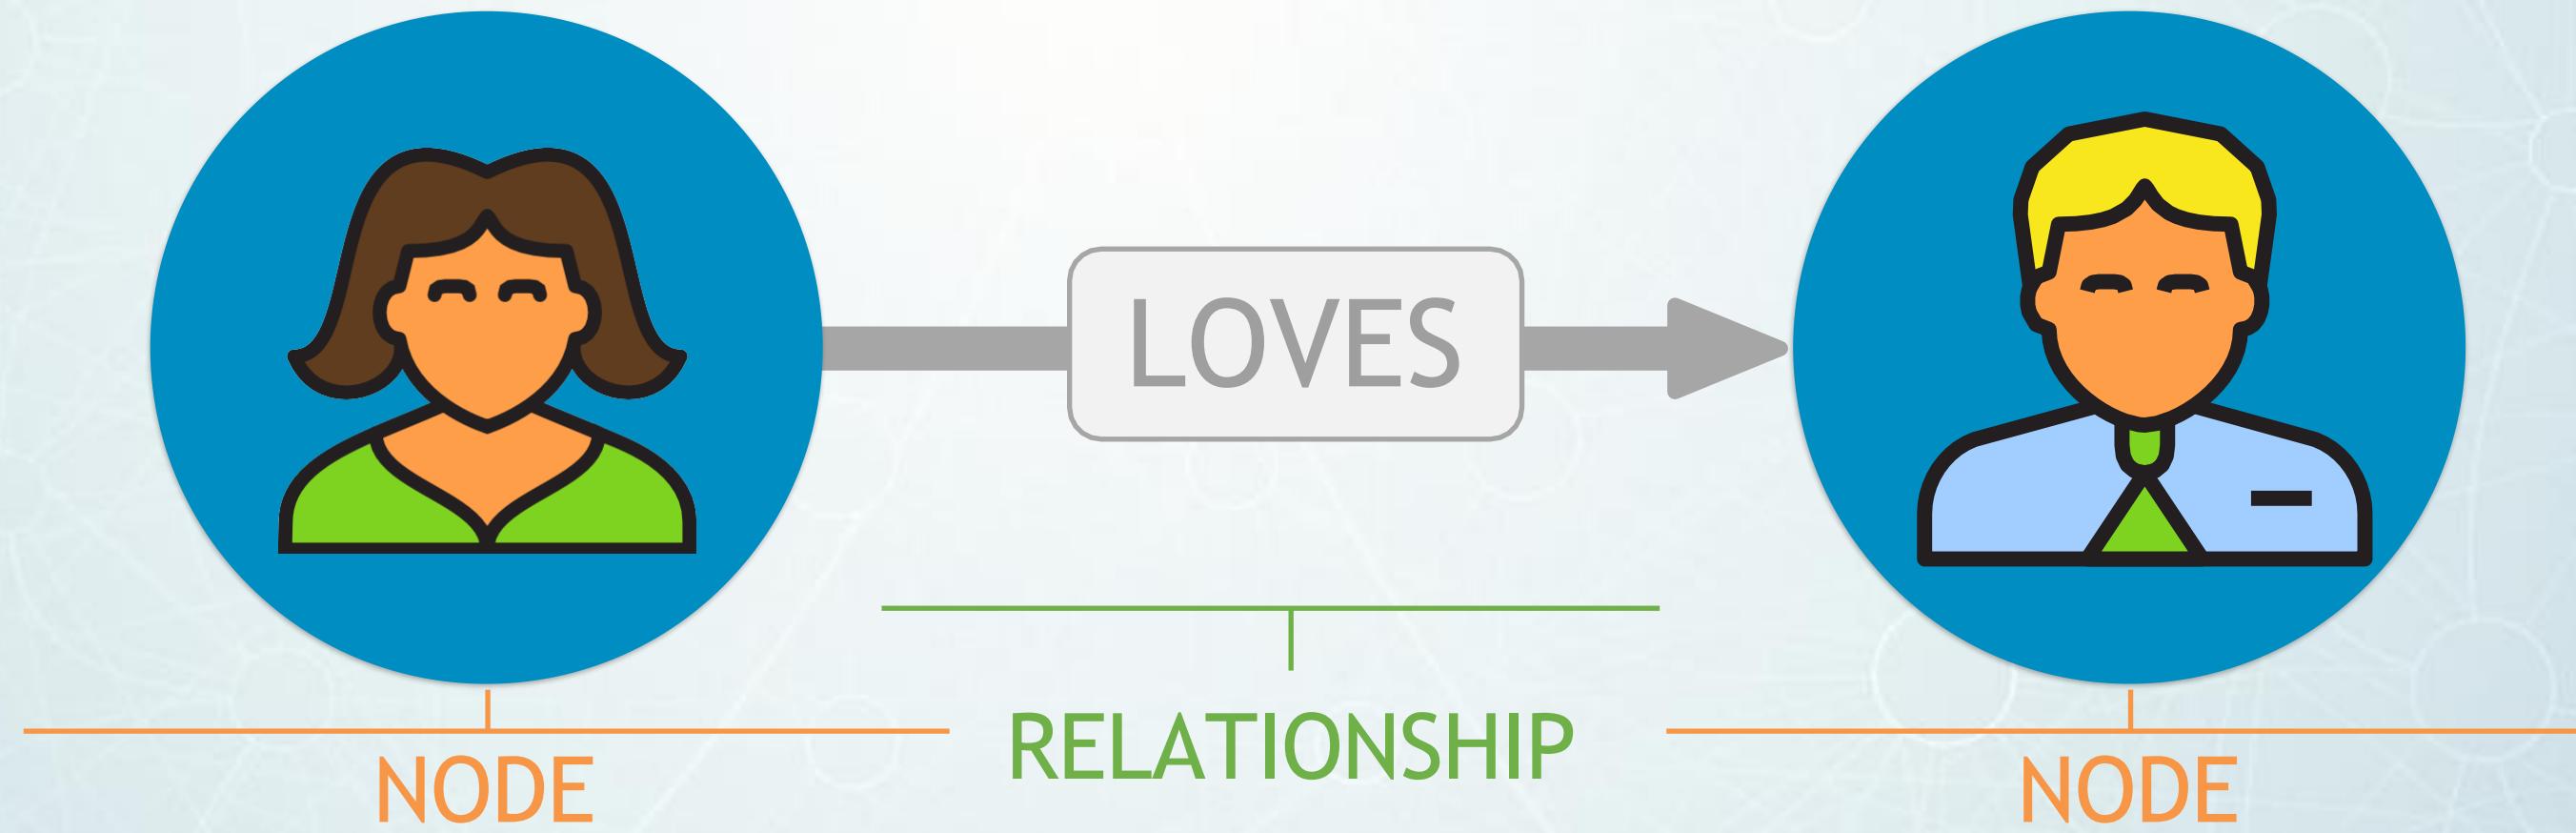


Loves

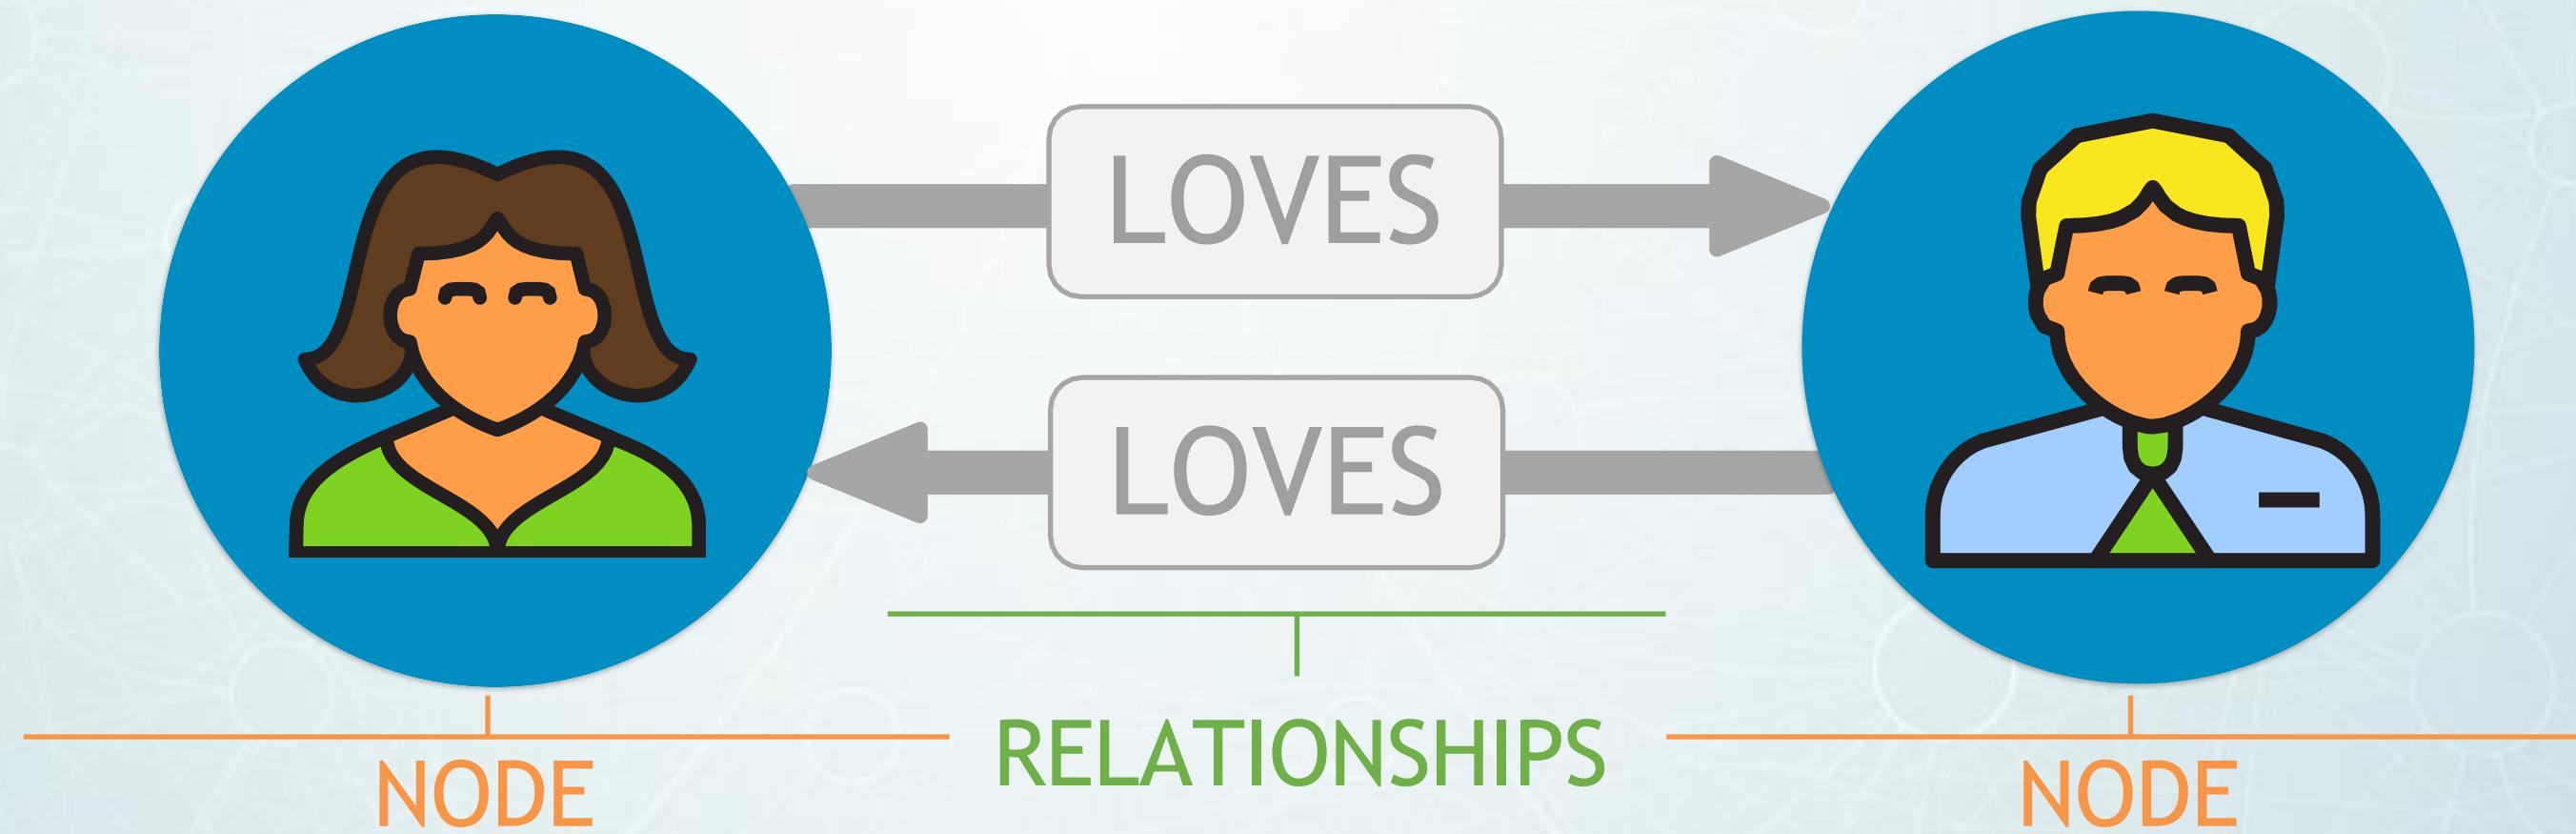


Dan

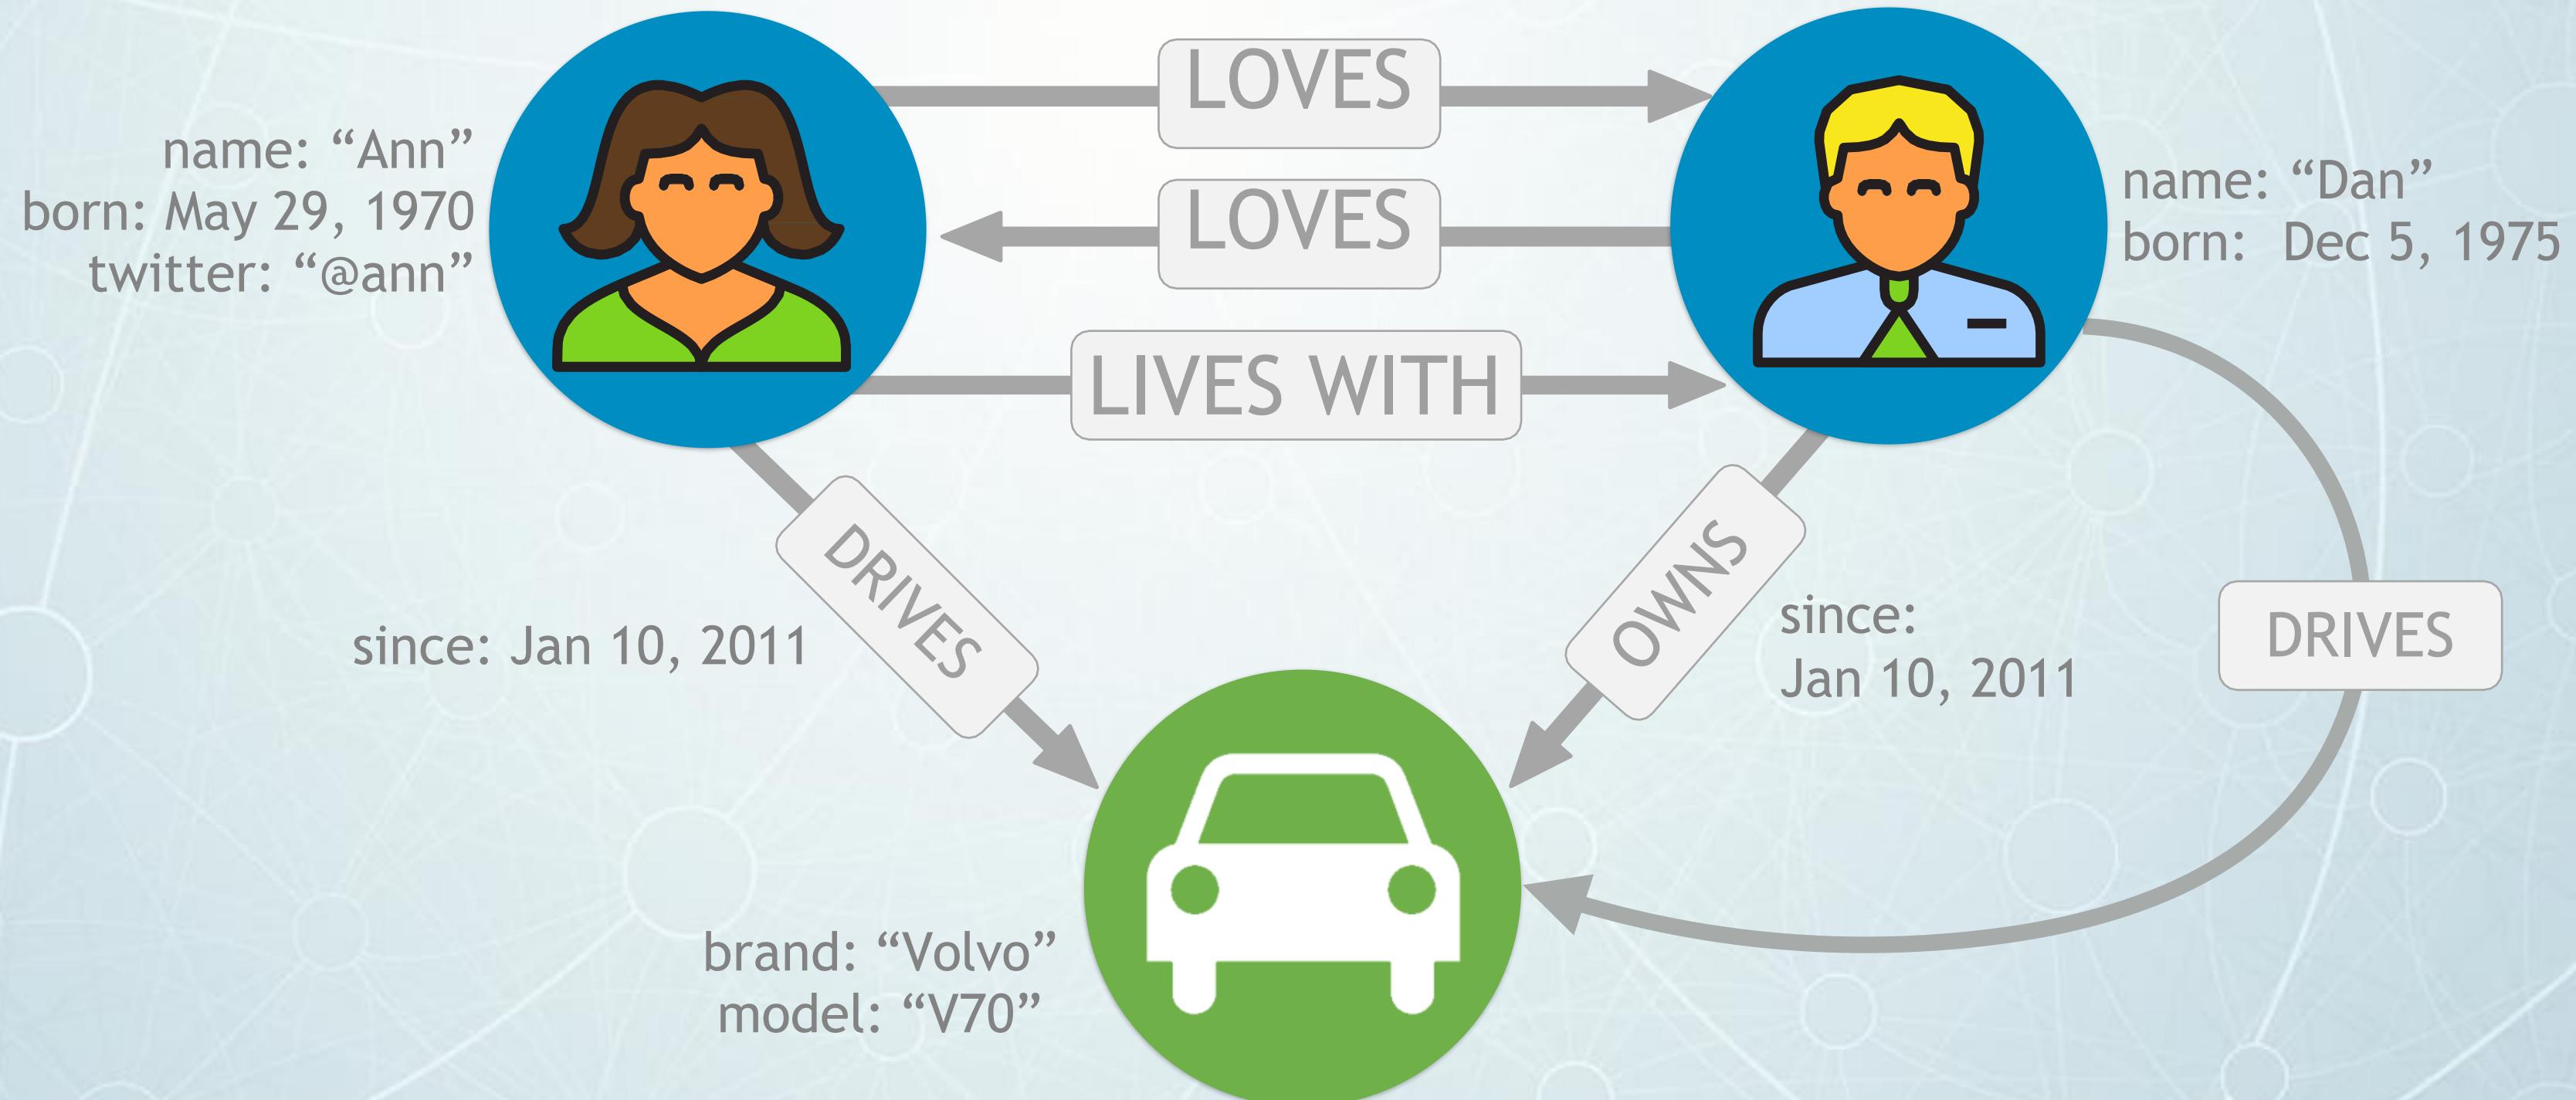
# Ann Loves Dan



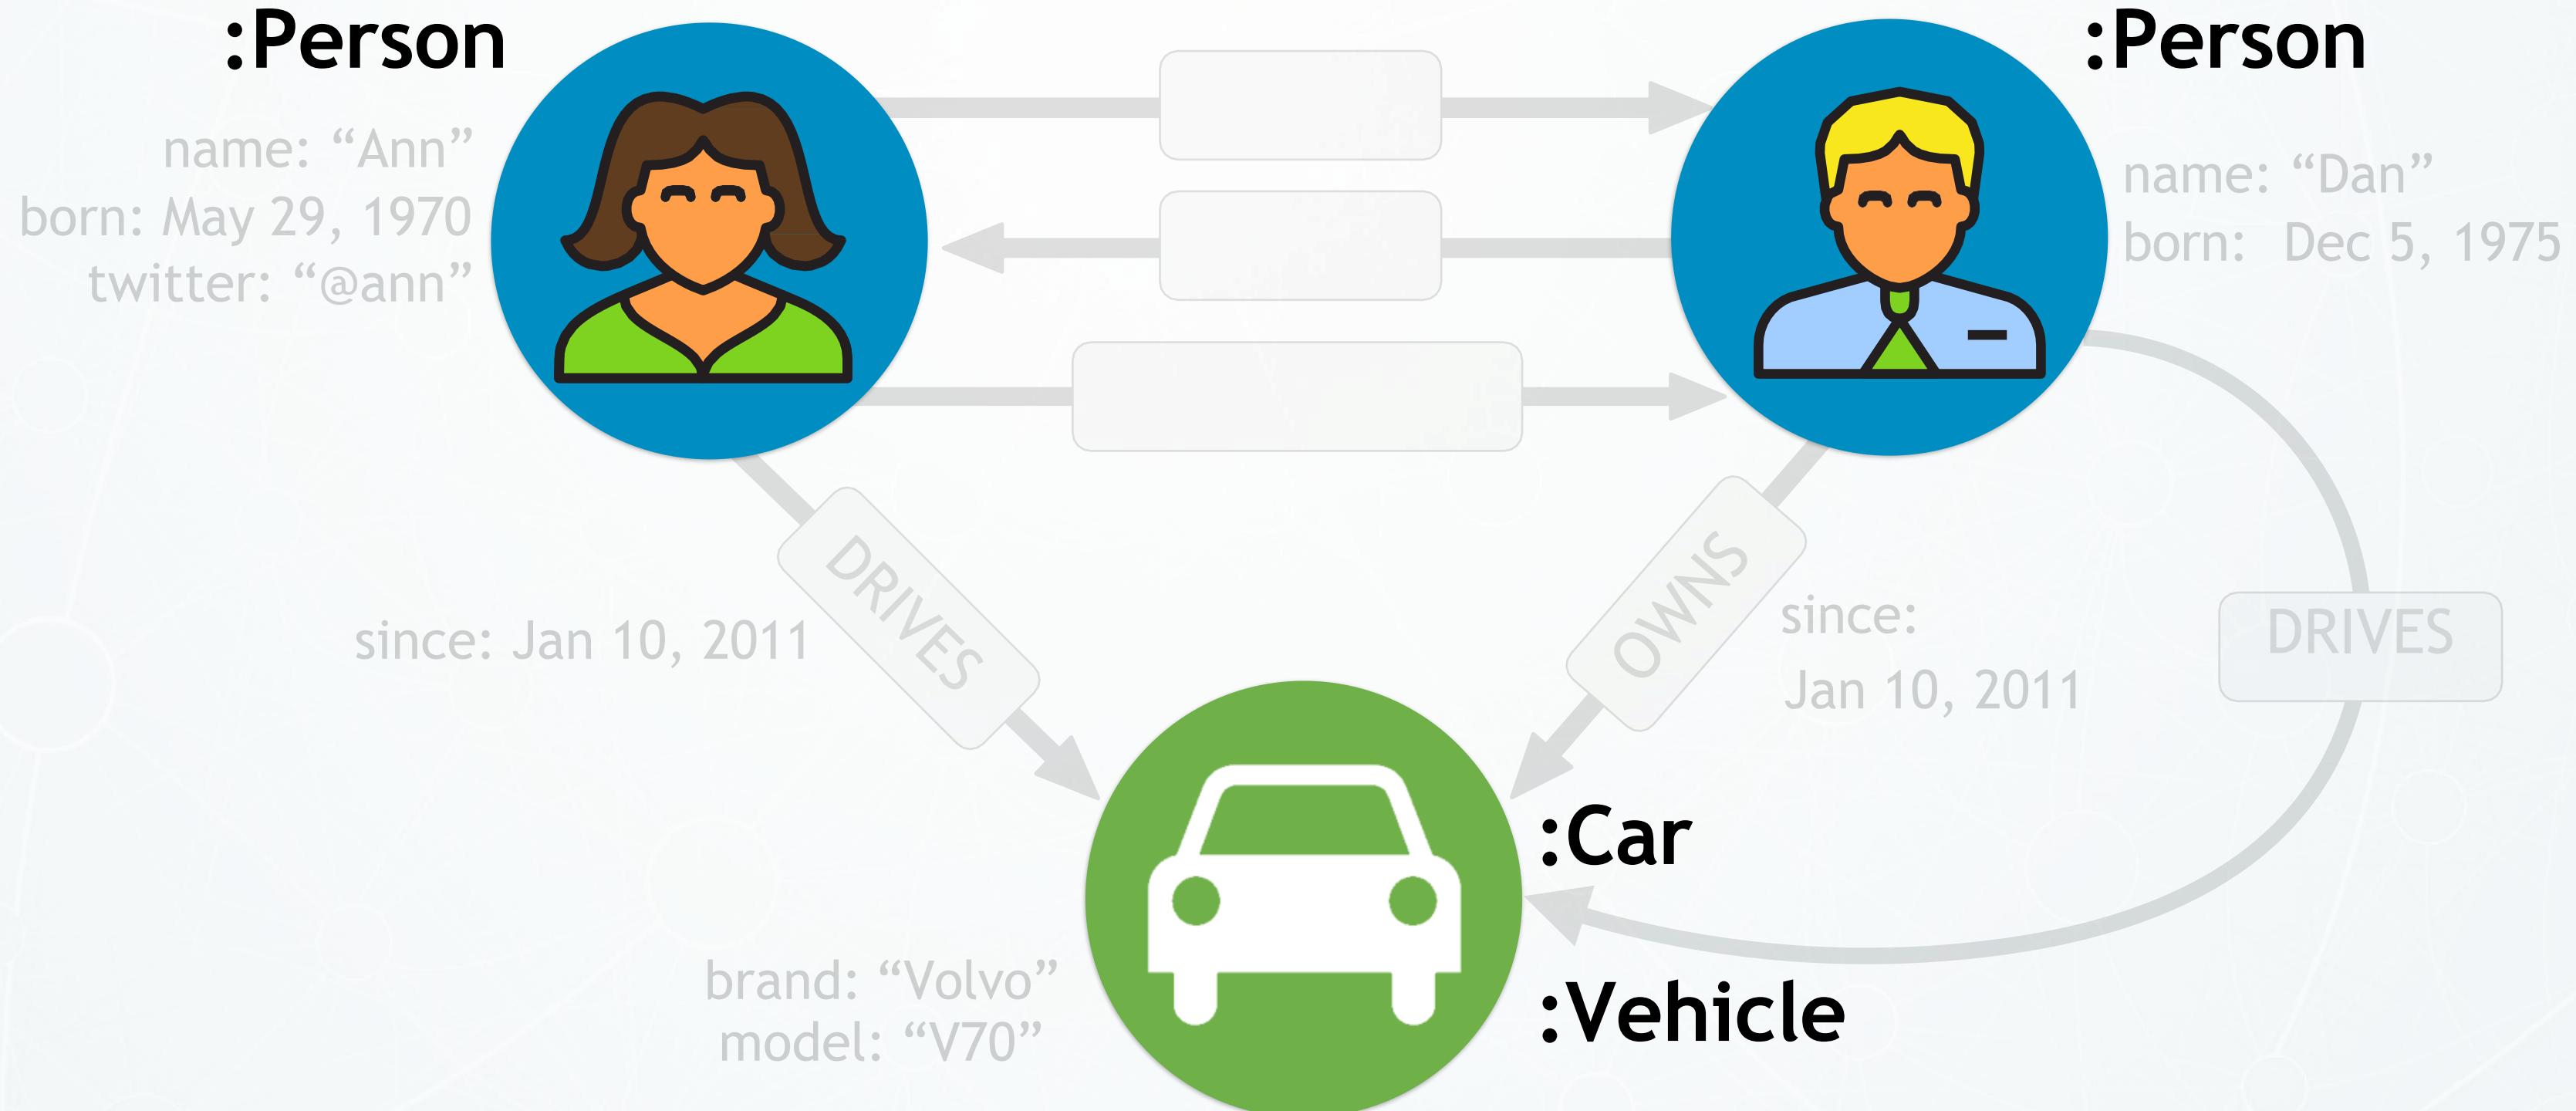
# Relationships are Directional



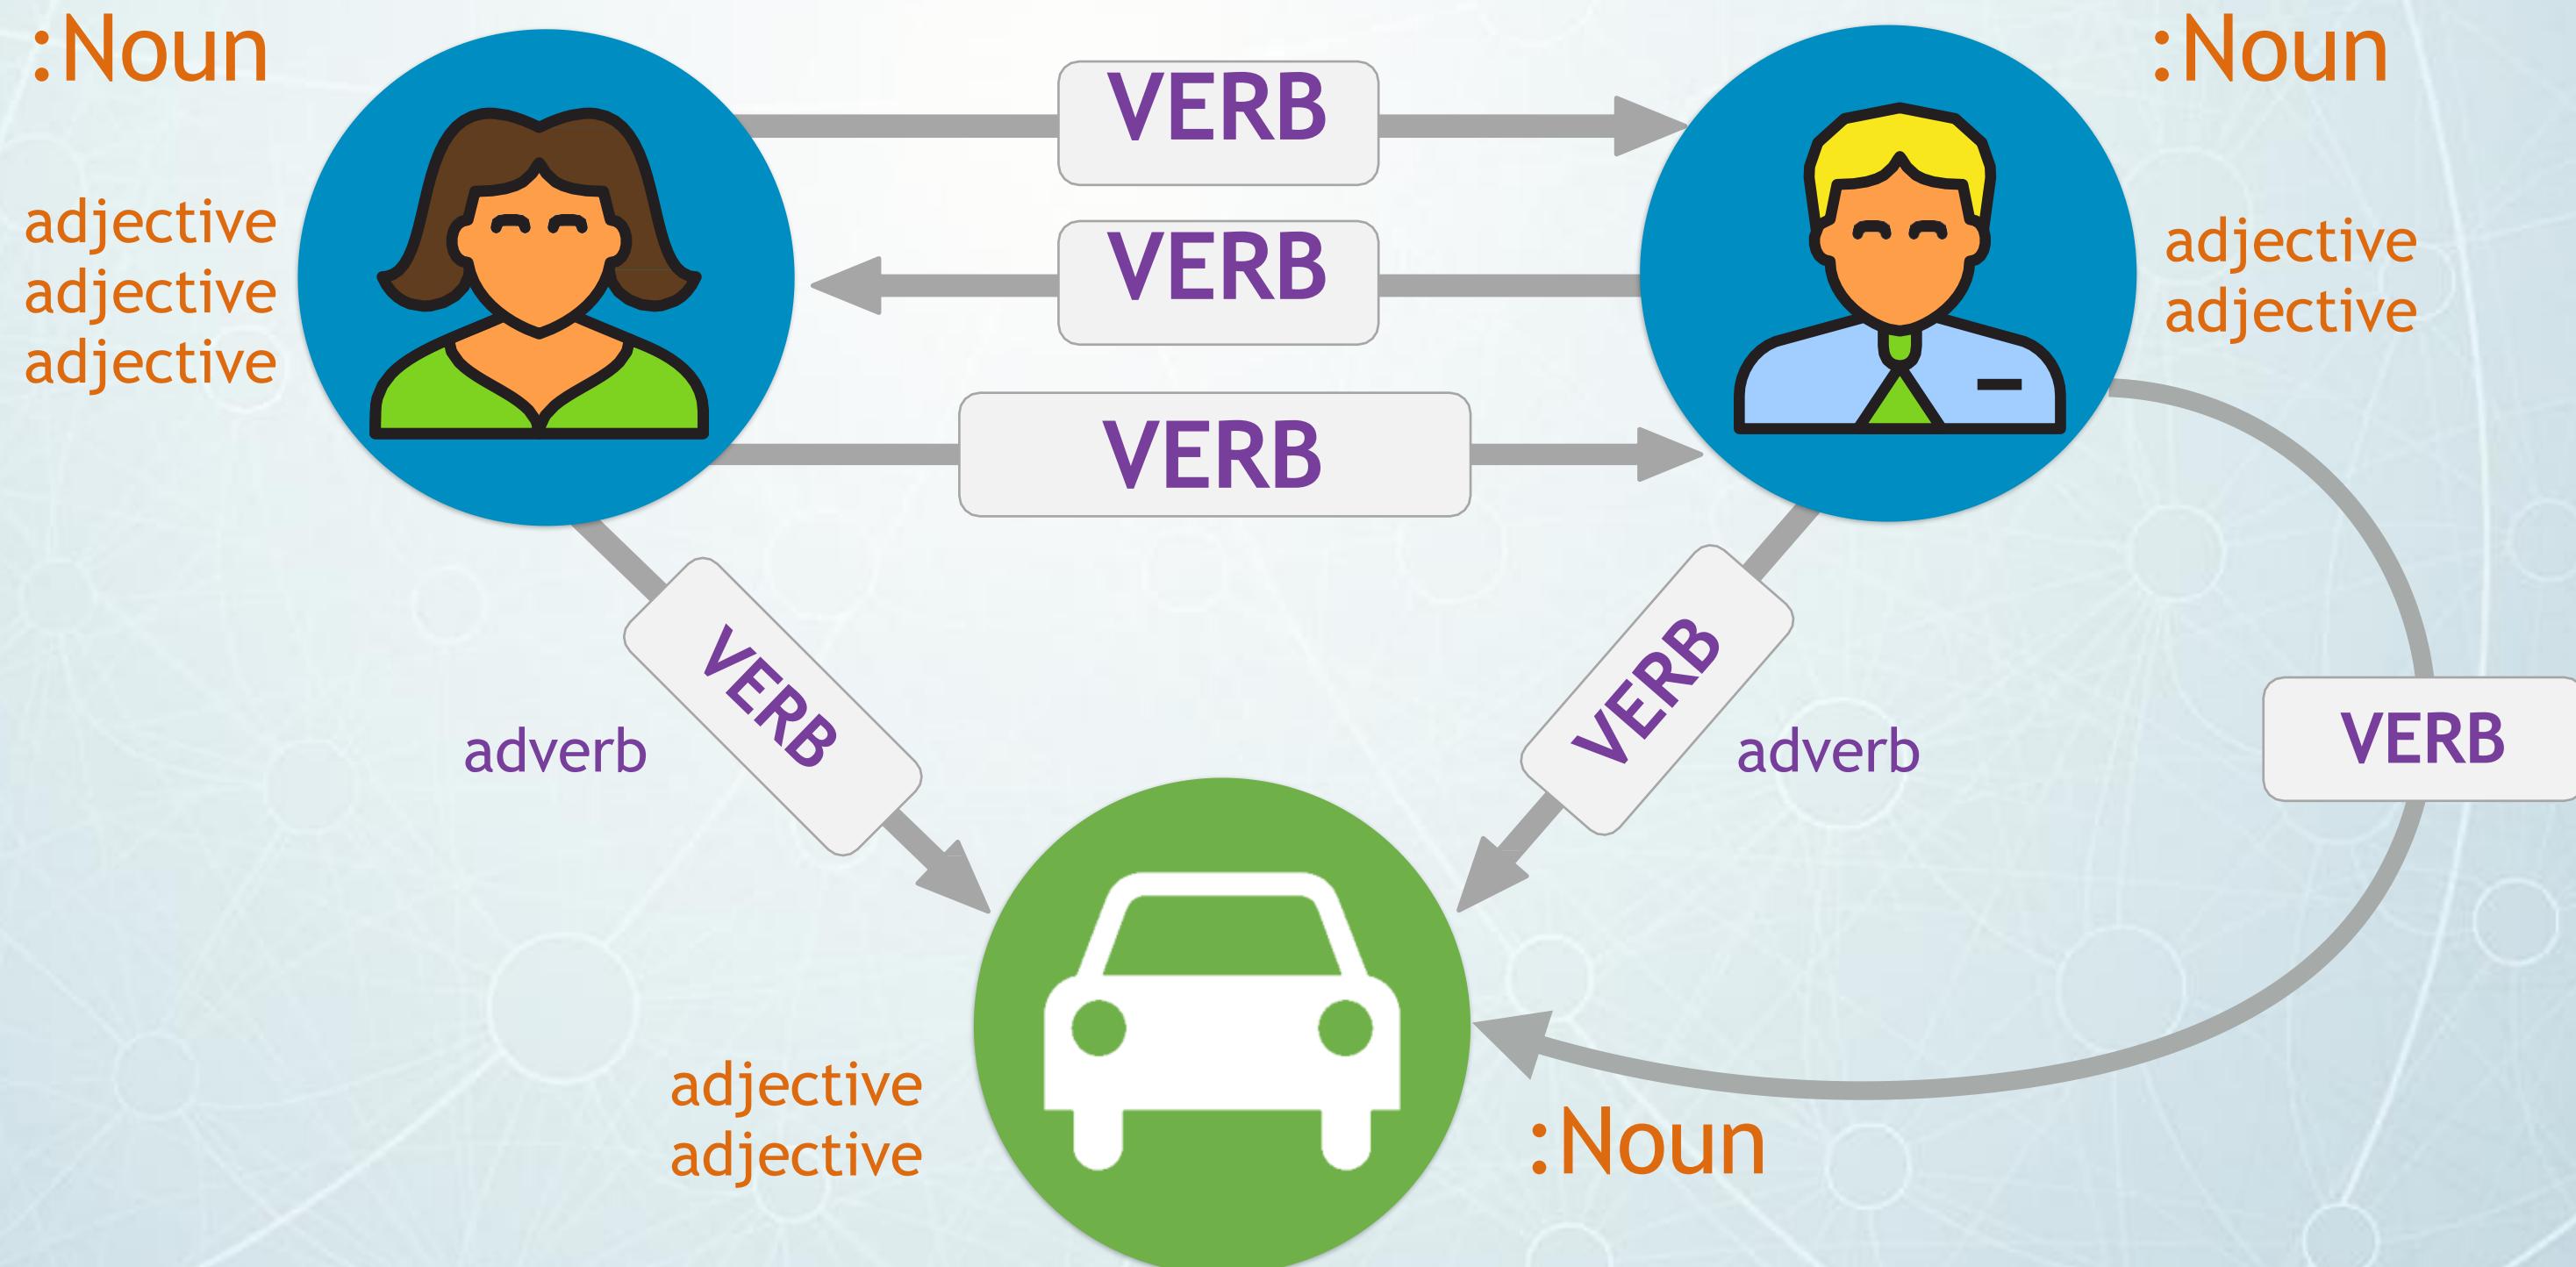
# Detailed Property Graph



# Labeled Property Graph



# Mapping to Languages



# Property Graph Model Components

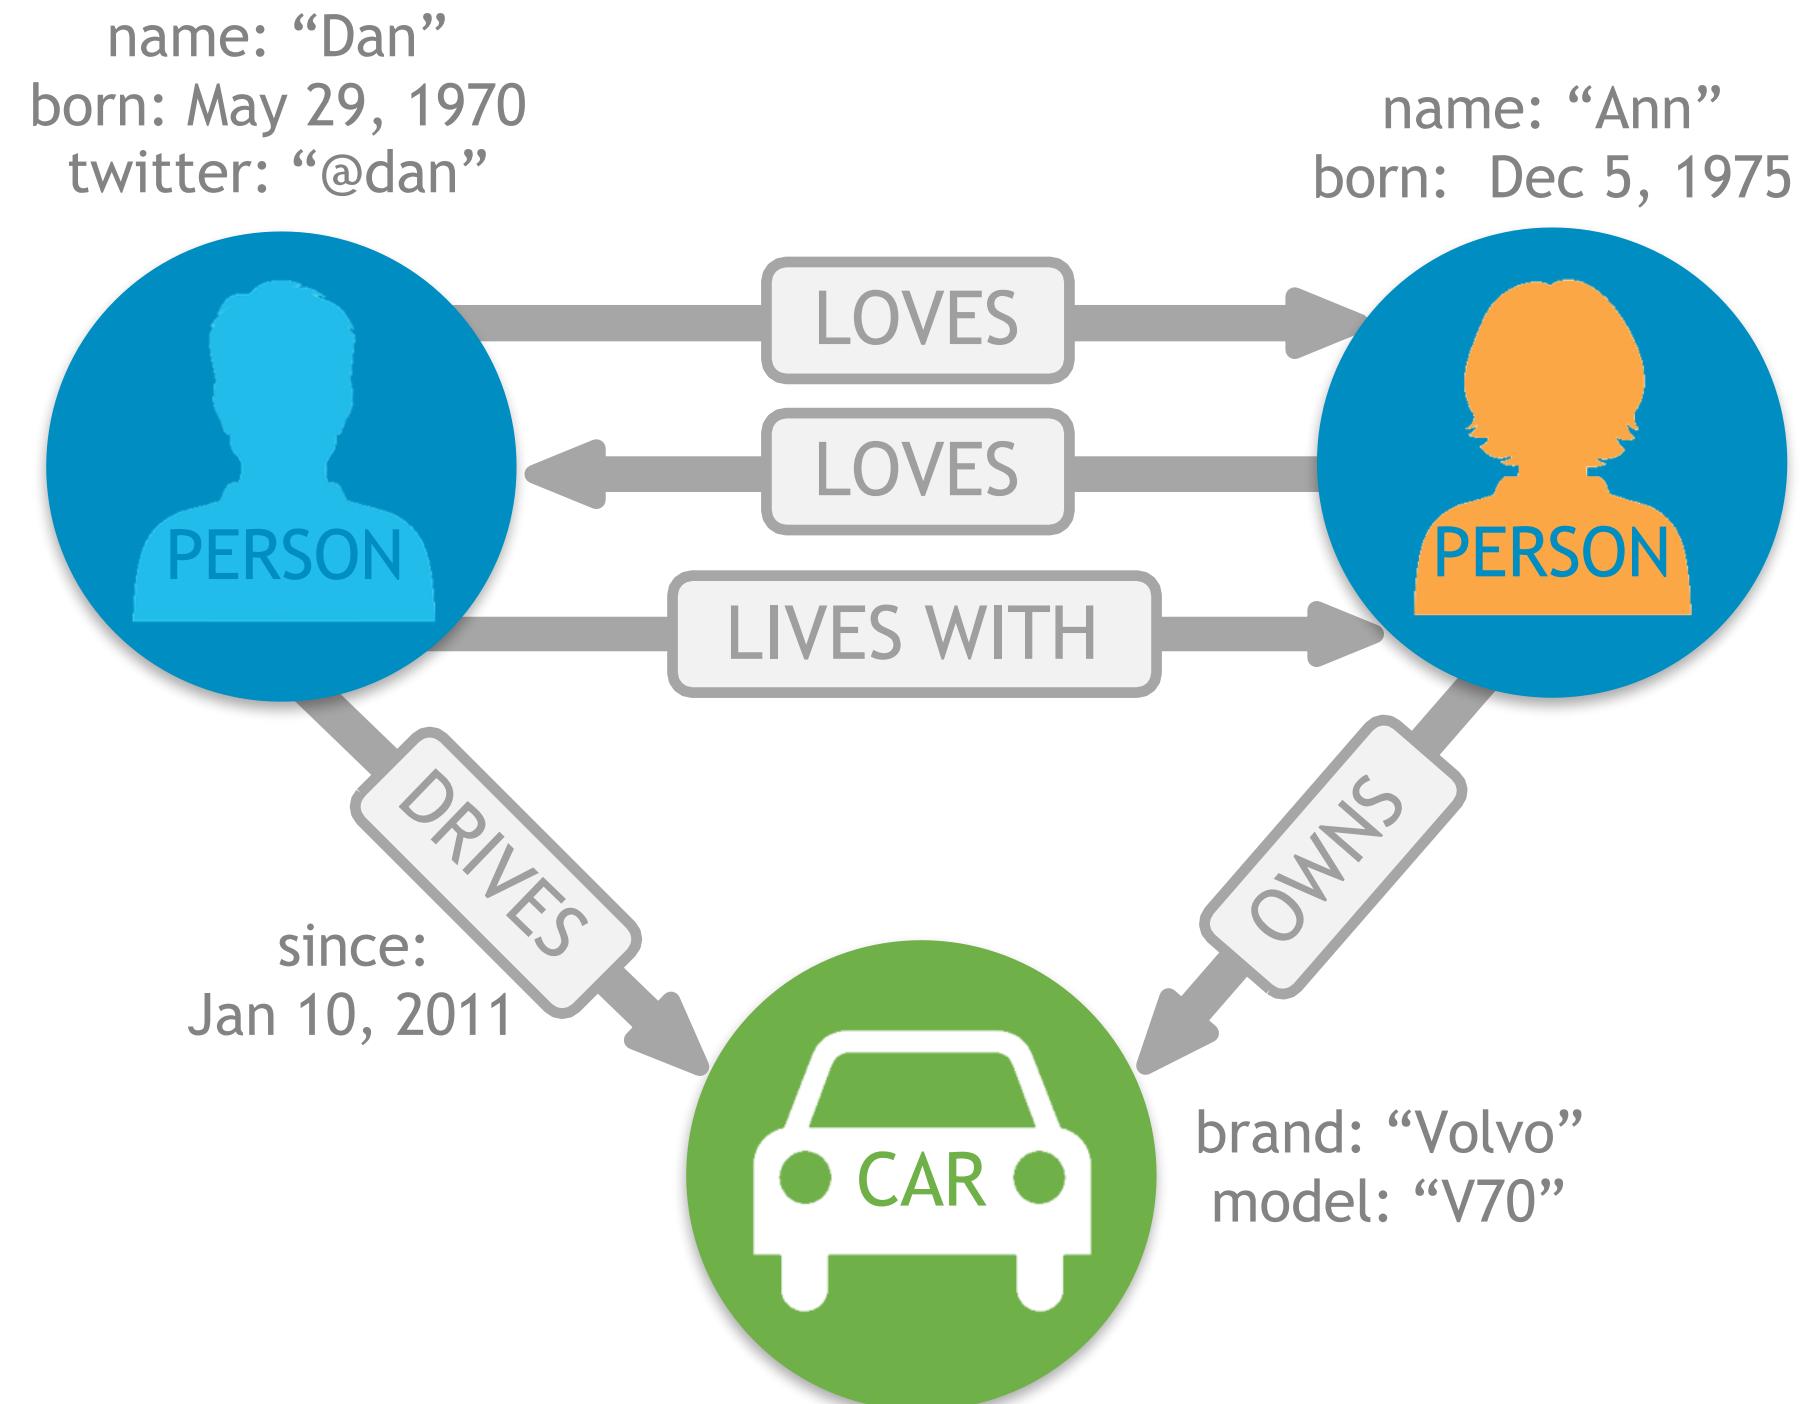


## Nodes

- The objects in the graph
- Can have name-value *properties*
- Can be *labeled*

## Relationships

- Relate nodes by type and direction
- Can have name-value *properties*



# WHY GRAPHS?



**Intuitivness**

**Speed**

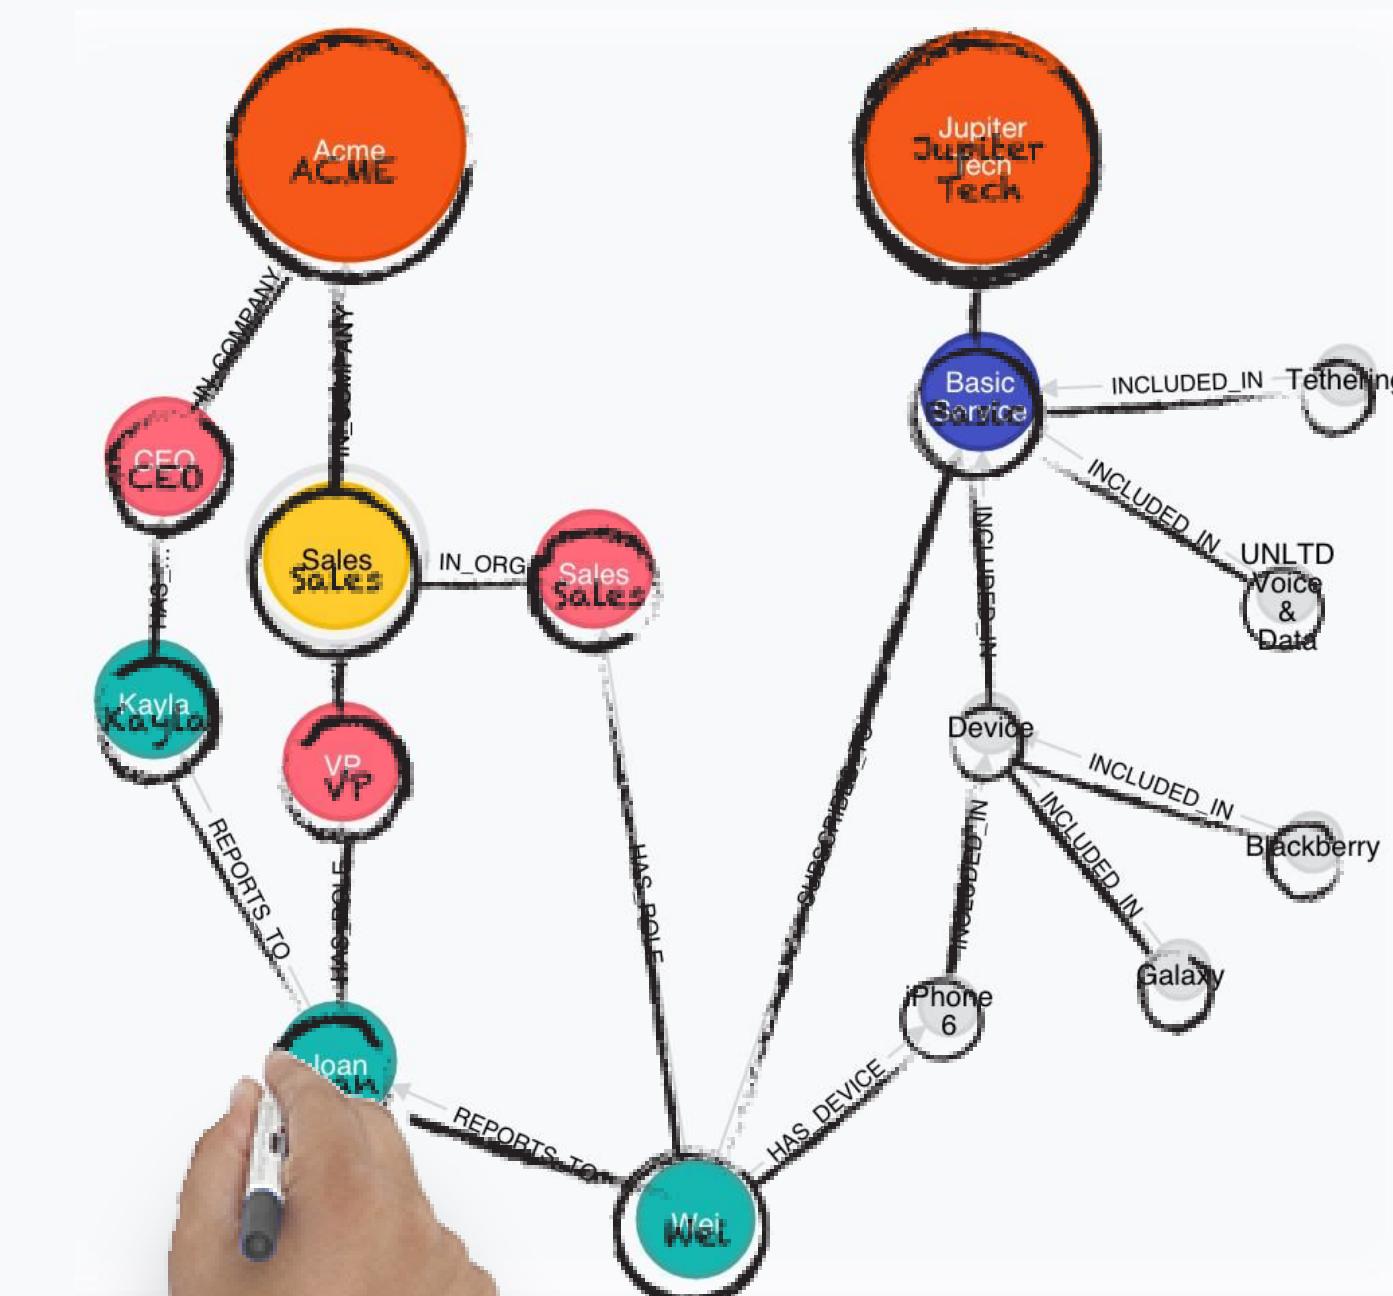
**Agility**

# Intuitiveness

Speed

Agility

# Intuitiveness



Intuitivness

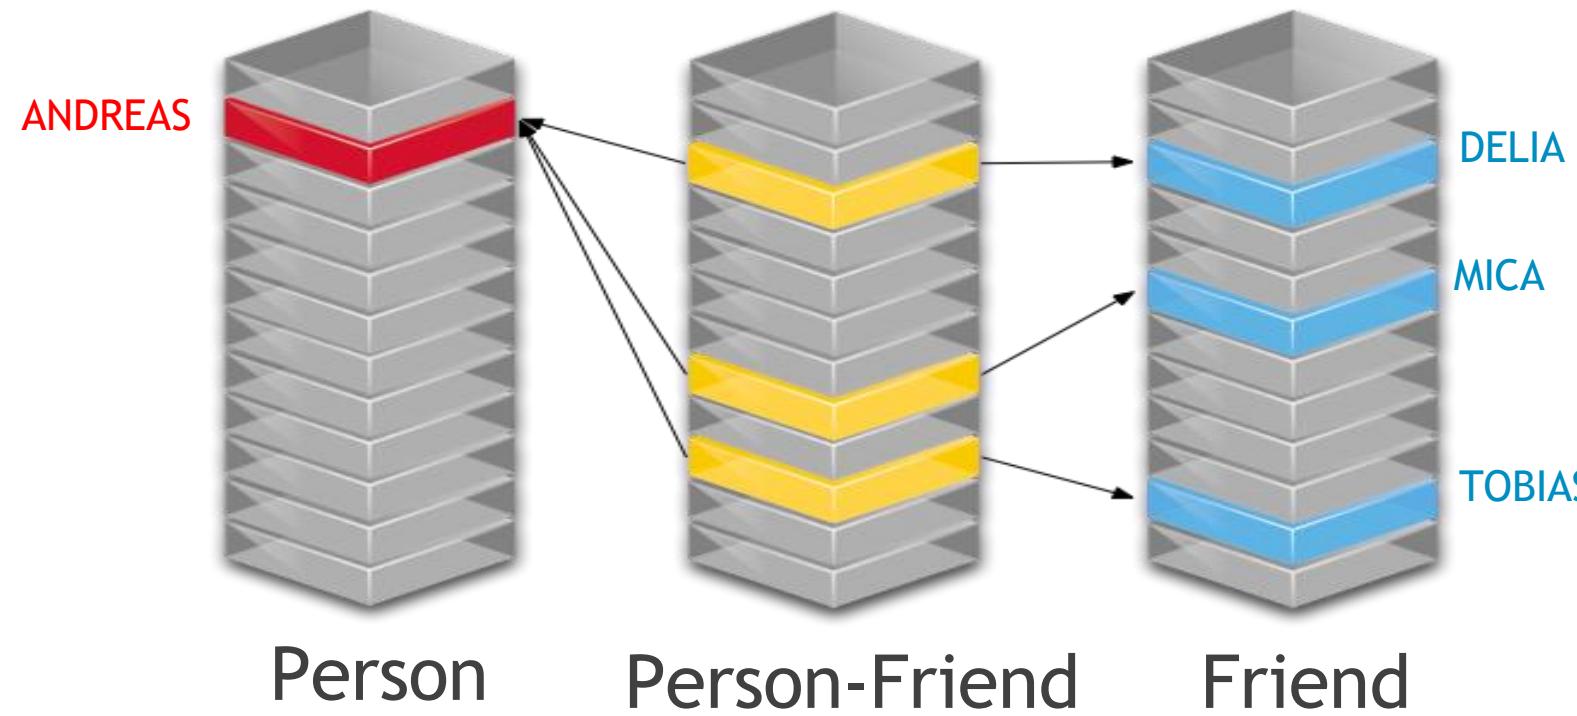
Speed

Agility

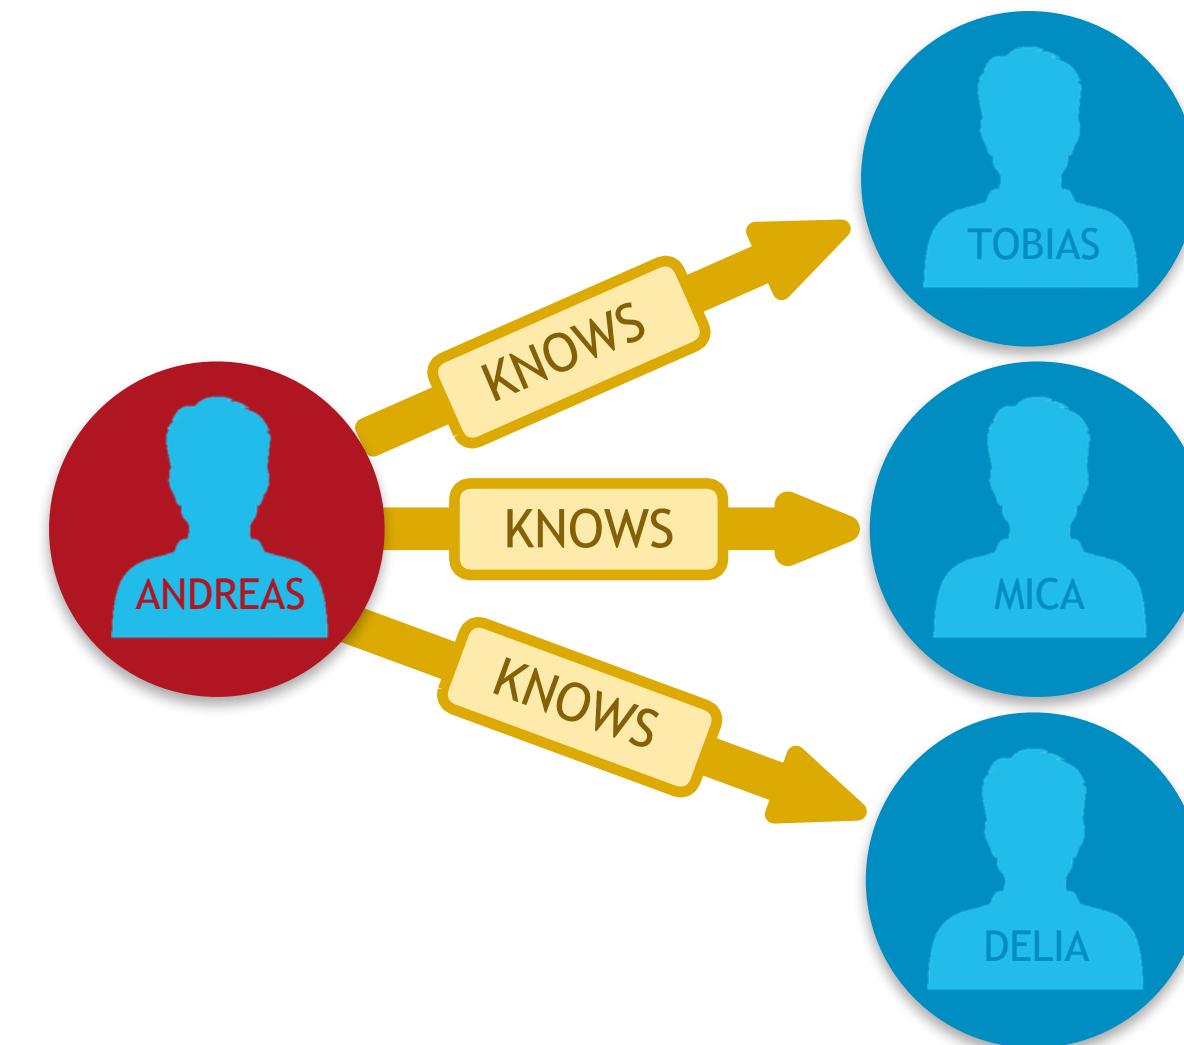
# Relational Versus Graph Models



## Relational Model



## Graph Model



Index free adjacency

# Speed



“We found Neo4j to be literally **thousands of times faster** than our prior MySQL solution, with queries that require 10-100 times less code. Today, Neo4j provides eBay with functionality that was previously impossible.”

- Volker Pacher, Senior Developer

**“Minutes to milliseconds” performance**  
*Queries up to 1000x faster than RDBMS or other NoSQL*



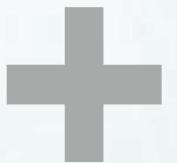
**Intuitivness**

**Speed**

**Agility**

# **Agility =**

**A Naturally Adaptive Model**



**A Query Language Designed  
for Connectedness**

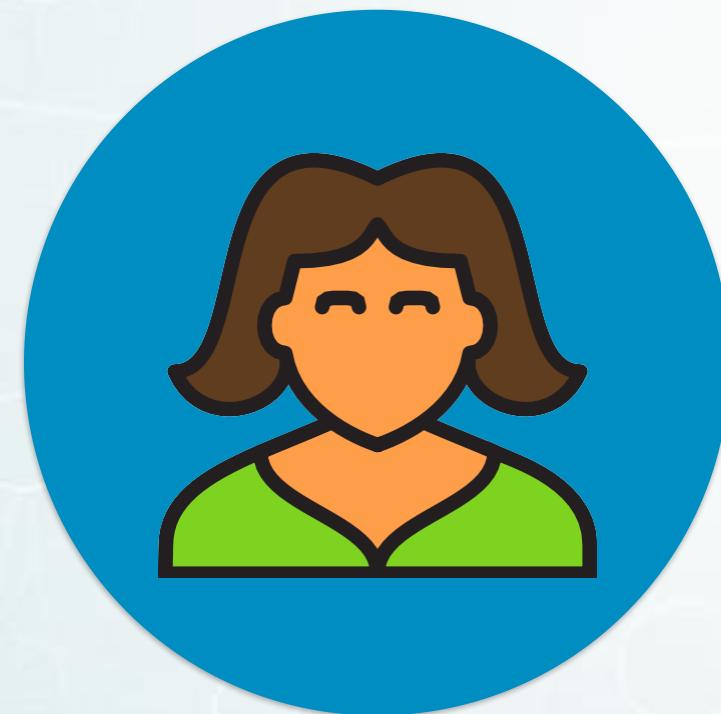
# CYPHER

SQL for graphs



# **(Very) Brief Cypher Tutorial**

# Creating the Data



NODE



NODE

```
CREATE (:Person { name:“Ann”}) - [:LOVES]-> (:Person { name:“Dan”})
```

# Representing Bi-Directionality



```
MATCH (:Person { name:“Ann”} ) - [:FB_FRIENDS] -> (:Person { name:“Dan”} )
```

```
MATCH (:Person { name:“Ann”} ) - [:FB_FRIENDS] -<- (:Person { name:“Dan”} )
```

# Cypher

## Typical Complex SQL Join

```
(SELECT T.directReportees AS directReportees, sum(T.count) AS count
FROM (
SELECT manager.pid AS directReportees, 0 AS count
  FROM person_reportee manager
 WHERE manager.pid = (SELECT id FROM person WHERE name = "fName lName")
UNION
  SELECT manager.pid AS directReportees, count(manager.directly_manages) AS count
  FROM person_reportee manager
 WHERE manager.pid = (SELECT id FROM person WHERE name = "fName lName")
 GROUP BY directReportees
UNION
SELECT manager.pid AS directReportees, count(reportee.directly_manages) AS count
  FROM person_reportee manager
 JOIN person_reportee reportee
ON manager.directly_manages = reportee.pid
 WHERE manager.pid = (SELECT id FROM person WHERE name = "fName lName")
 GROUP BY directReportees
UNION
SELECT manager.pid AS directReportees, count(L2Reportees.directly_manages) AS count
  FROM person_reportee manager
 JOIN person_reportee L1Reportees
ON manager.directly_manages = L1Reportees.pid
JOIN person_reportee reportee
ON manager.directly_manages = reportee.pid
 WHERE manager.pid = (SELECT id FROM person WHERE name = "fName lName")
 GROUP BY directReportees
) AS T
GROUP BY directReportees)
UNION
(SELECT T.directReportees AS directReportees, sum(T.count) AS count
FROM (
SELECT manager.directly_manages AS directReportees, 0 AS count
  FROM person_reportee manager
 WHERE manager.pid = (SELECT id FROM person WHERE name = "fName lName")
UNION
  SELECT reportee.pid AS directReportees, count(reportee.directly_manages) AS count
  FROM person_reportee manager
 JOIN person_reportee reportee
ON manager.directly_manages = reportee.pid
 WHERE manager.pid = (SELECT id FROM person WHERE name = "fName lName")
 GROUP BY directReportees
) )
```

```
SELECT depth1Reportees.pid AS directReportees,
count(depth2Reportees.directly_manages) AS count
  FROM person_reportee manager
 JOIN person_reportee L1Reportees
ON manager.directly_manages = L1Reportees.pid
 JOIN person_reportee L2Reportees
ON L1Reportees.directly_manages = L2Reportees.pid
 WHERE manager.pid = (SELECT id FROM person WHERE name = "fName lName")
 GROUP BY directReportees
) AS T
GROUP BY directReportees)
UNION
(SELECT T.directReportees AS directReportees, sum[T.count] AS count
  FROM(
    SELECT reportee.directly_manages AS directReportees, 0 AS count
  FROM person_reportee manager
 JOIN person_reportee reportee
ON manager.directly_manages = reportee.pid
 WHERE manager.pid = (SELECT id FROM person WHERE name = "fName lName")
 GROUP BY directReportees
) AS T
GROUP BY directReportees)
UNION
SELECT L2Reportees.pid AS directReportees, count(L2Reportees.directly_manages)
AS count
  FROM person_reportee manager
 JOIN person_reportee L1Reportees
ON manager.directly_manages = L1Reportees.pid
 JOIN person_reportee L2Reportees
ON L1Reportees.directly_manages = L2Reportees.pid
 WHERE manager.pid = (SELECT id FROM person WHERE name = "fName lName")
 GROUP BY directReportees
) AS T
GROUP BY directReportees)
UNION
(SELECT L2Reportees.directly_manages AS directReportees, 0 AS count
  FROM person_reportee manager
 JOIN person_reportee L1Reportees
ON manager.directly_manages = L1Reportees.pid
 JOIN person_reportee L2Reportees
ON L1Reportees.directly_manages = L2Reportees.pid
 WHERE manager.pid = (SELECT id FROM person WHERE name = "fName lName")
)
```

## The Same Query using Cypher

```
MATCH (boss) - [:MANAGES*0..3] -> (sub),
      (sub) - [:MANAGES*1..3] -> (report)
WHERE boss.name = "John Doe"
RETURN sub.name AS Subordinate,
       count(report) AS Total
```

## Project Impact

### Less time writing queries

- More time understanding the answers
- Leaving time to ask the next question

### Less time debugging queries:

- More time writing the next piece of code
- Improved quality of overall code base

### Code that's easier to read:

- Faster ramp-up for new project members
- Improved maintainability & troubleshooting

## What is openCypher?

openCypher is an open source project to bring a new public implementation of the industry's most widely adopted graph query language: **Cypher**.

### Focus on Your Domain

Graphs naturally describe your domain, and Cypher lets you focus on that domain instead of getting lost in the mechanics of data access. Both expressive and efficient, Cypher is intuitive and immediately familiar, without the steep learning curve of learning a new language.

### Human Readable

Cypher is a human-readable query language that makes complex things possible. A combination of English prose and intuitive iconography, Cypher is accessible to developers and operations professionals alike.

### Complete Open Source Access

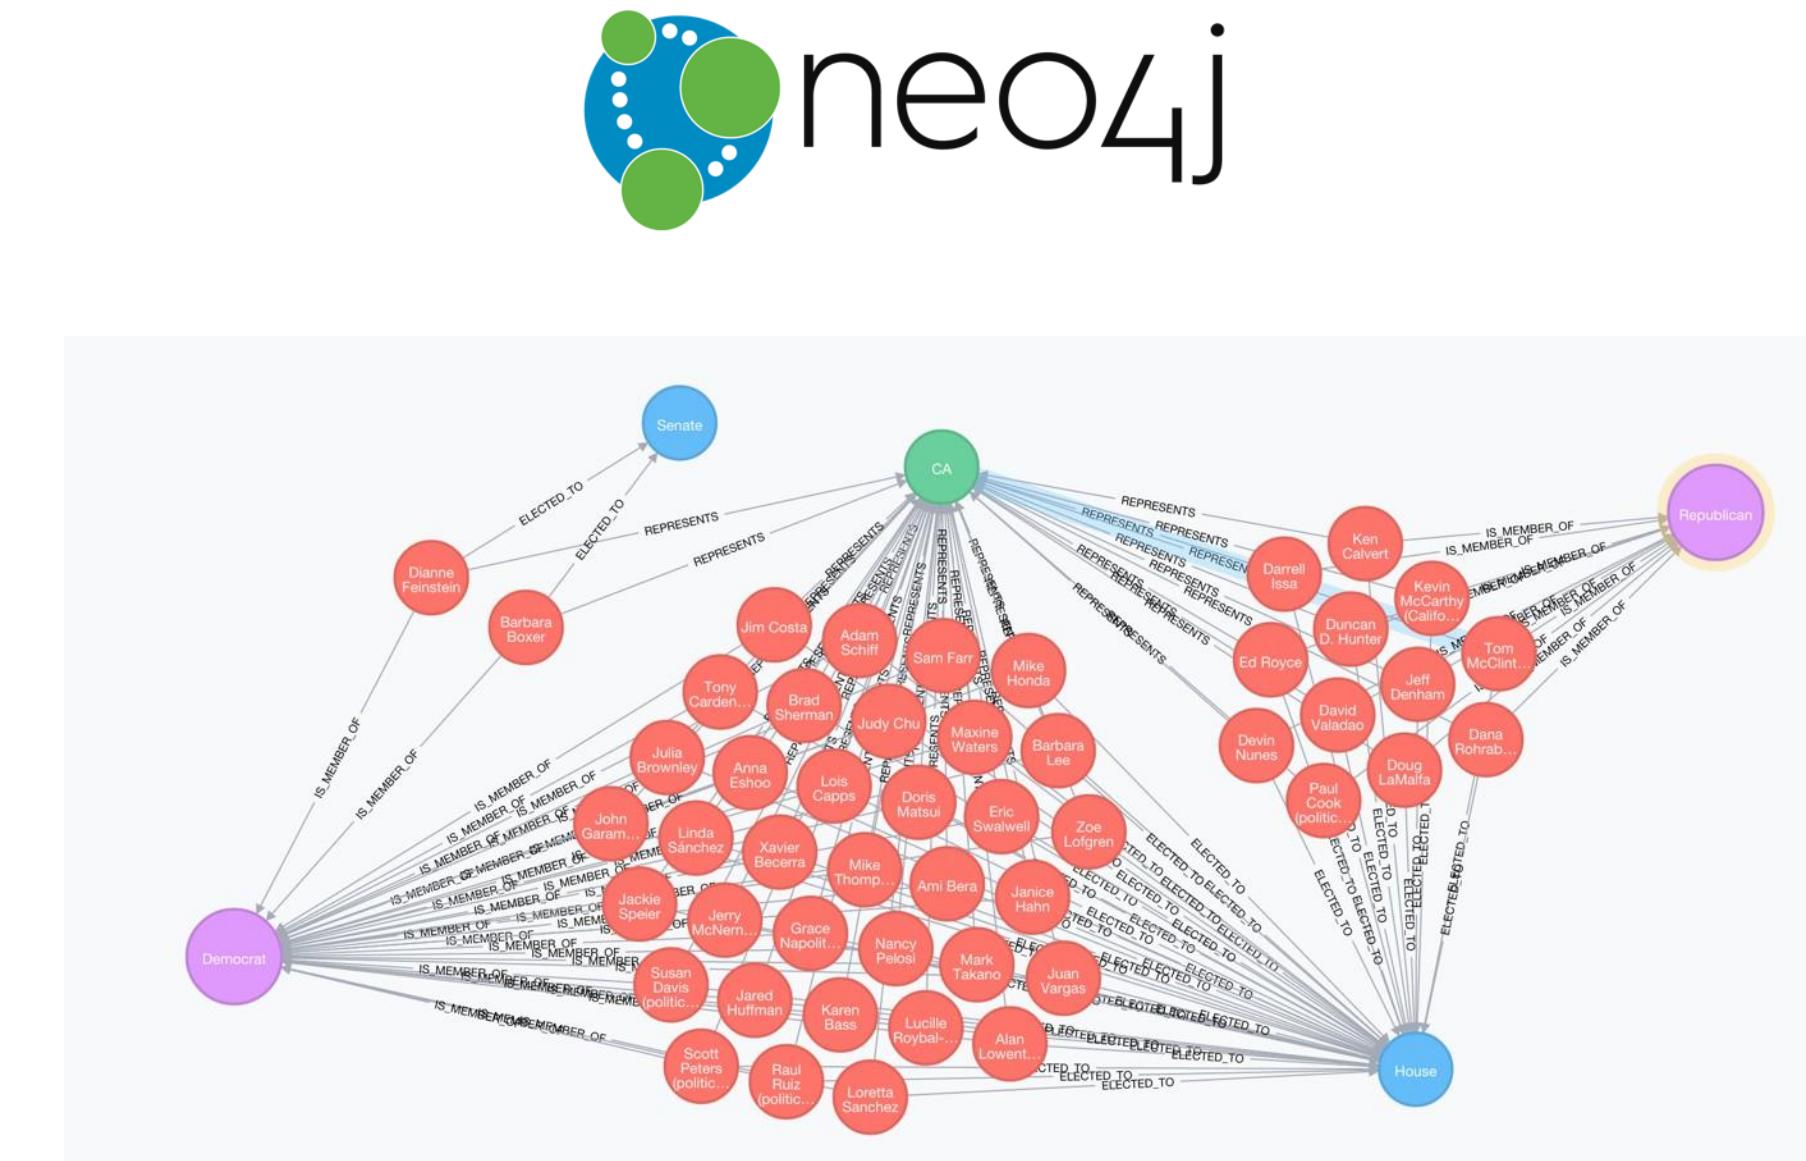
The openCypher project means you can use Cypher as your query language for graph processing capabilities within any product or application. Everyone from Oracle and Apache Spark to Tableau and Structr can (and do) use Cypher's capabilities – now you can too.

```
MATCH (cypher:QueryLanguage)-[:QUERIES]->(graphs)
MATCH (cypher)<-[USES]-(u:User) WHERE u.name IN ['Oracle', 'Apache Spark', 'Tableau', 'Structr']
MATCH (openCypher)-[:MAKES_AVAILABLE]->(cypher)
RETURN cypher.attributes
-----
['awesome',...]
```

<http://www.opencypher.org/>

## Graph Database

- Property graph data model
  - Nodes and relationships
- Native graph processing
- (open)Cypher query language



# Neo4j - Key Product Features



## Native Graph Storage

Ensures data consistency and performance

## Native Graph Processing

Millions of hops per second, in real time

## “Whiteboard Friendly” Data Modeling

Model data as it naturally occurs

## High Data Integrity

Fully ACID transactions

## Powerful, Expressive Query Language

Requires 10x to 100x less code than SQL

## Scalability and High Availability

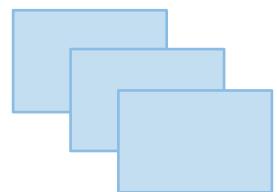
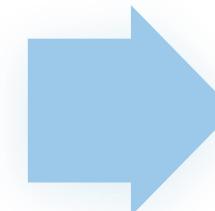
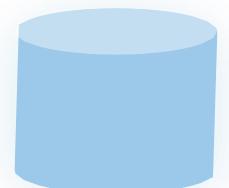
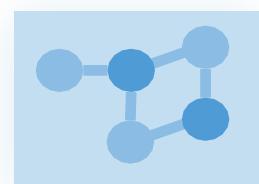
Vertical and horizontal scaling optimized for graphs

## Built-in ETL

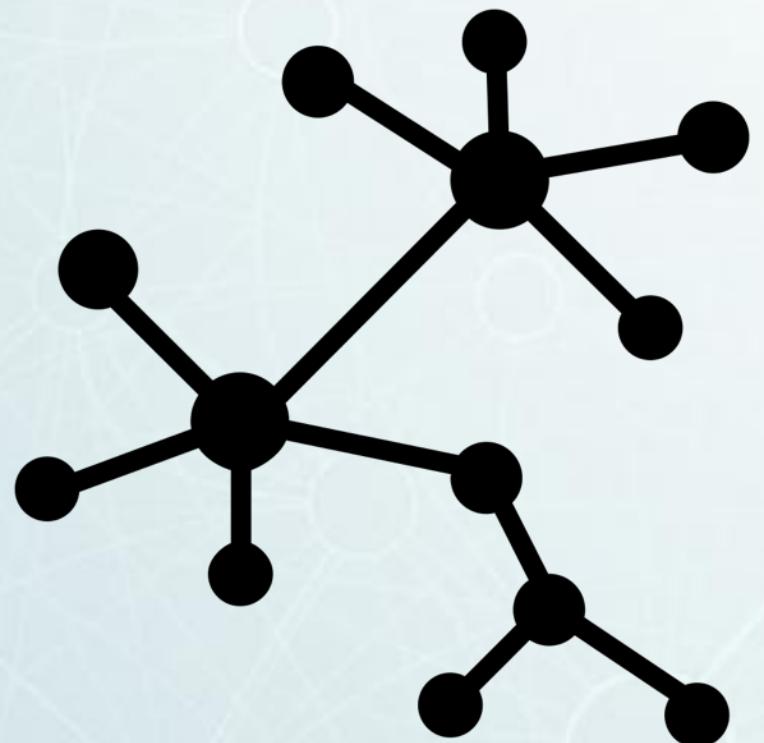
Seamless import from other databases

## Integration

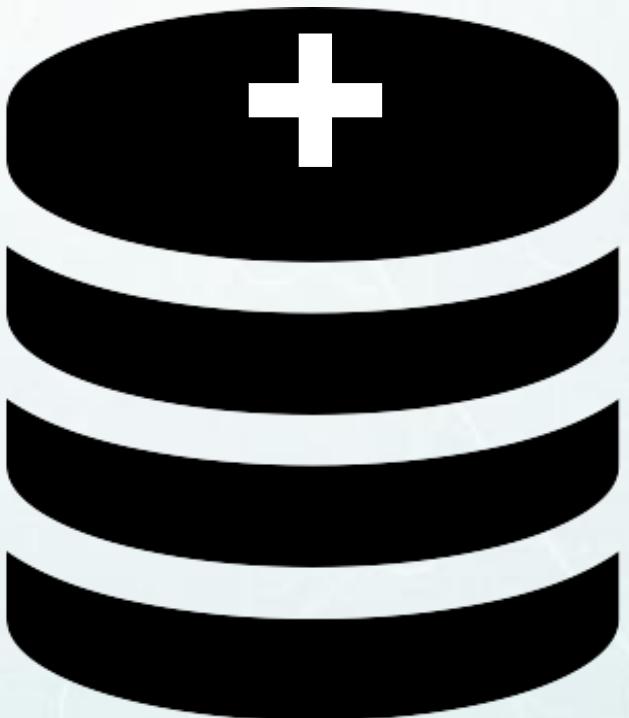
Drivers and APIs for popular languages



# How do you use Neo4j?



CREATE MODEL



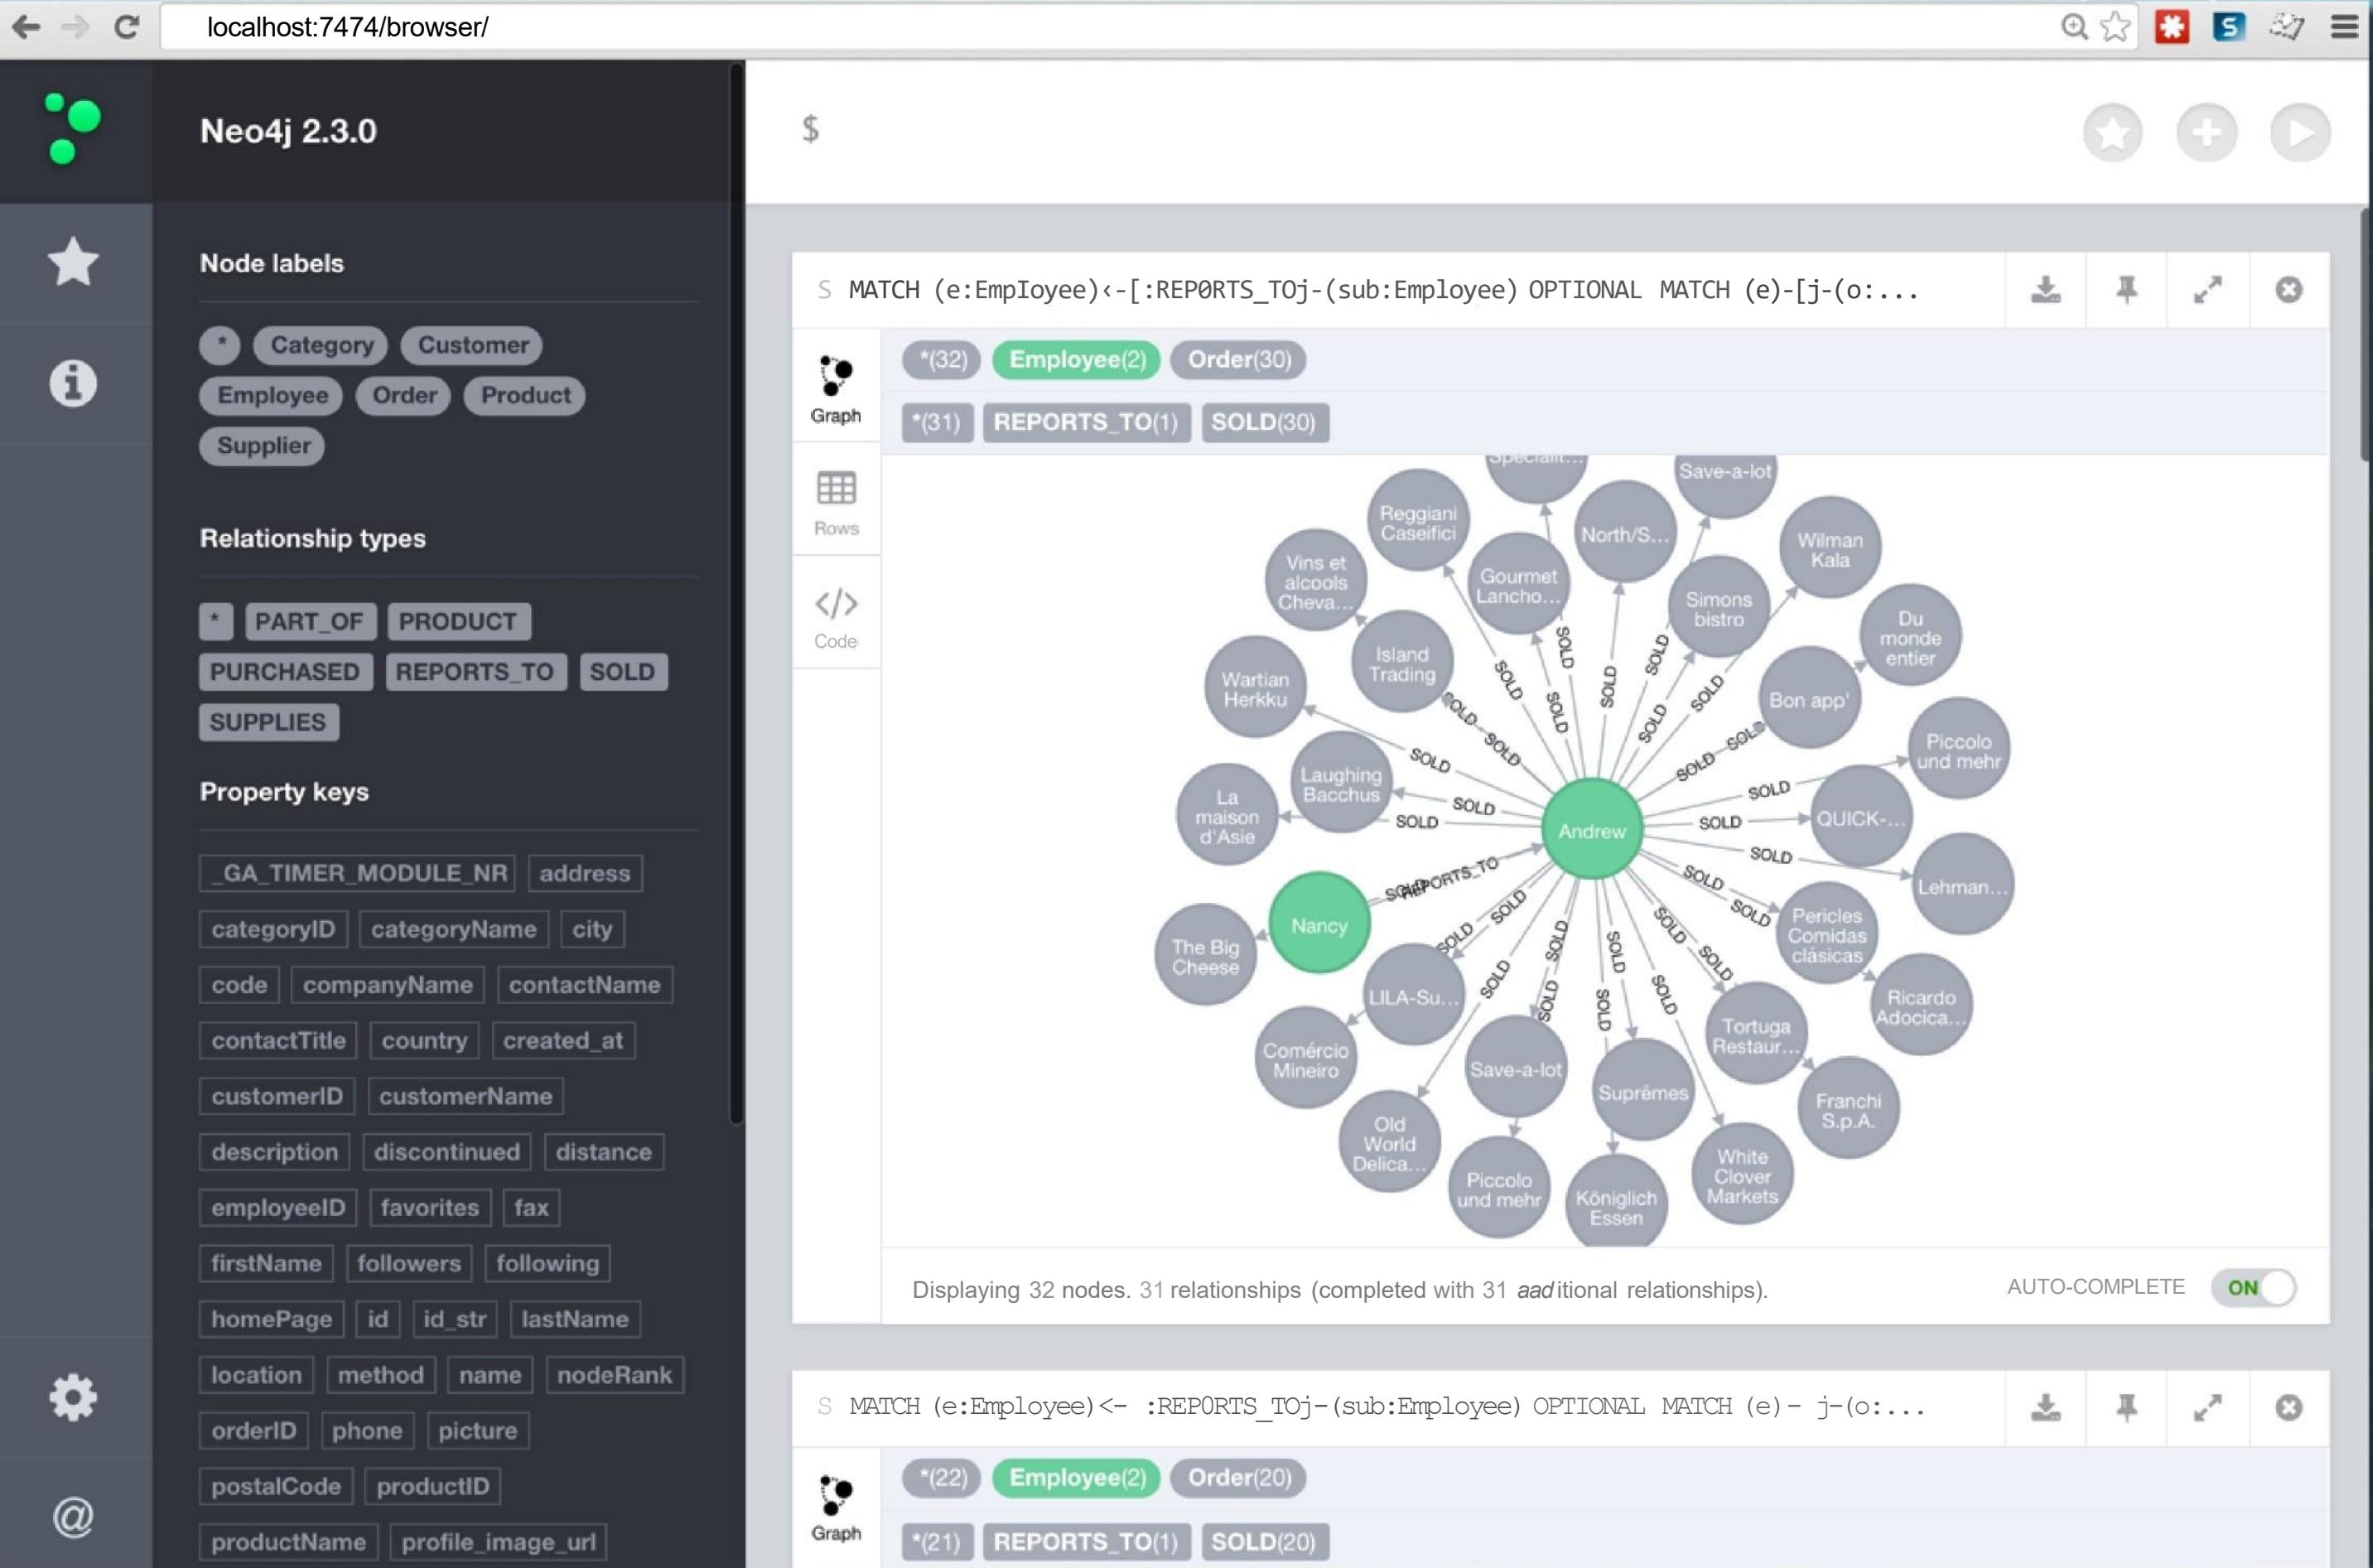
LOAD DATA



QUERY DATA

# How do you use Neo4j?





# How do you use Neo4j?



# Language Drivers

**JS**



**nv** **php**

# Language Drivers

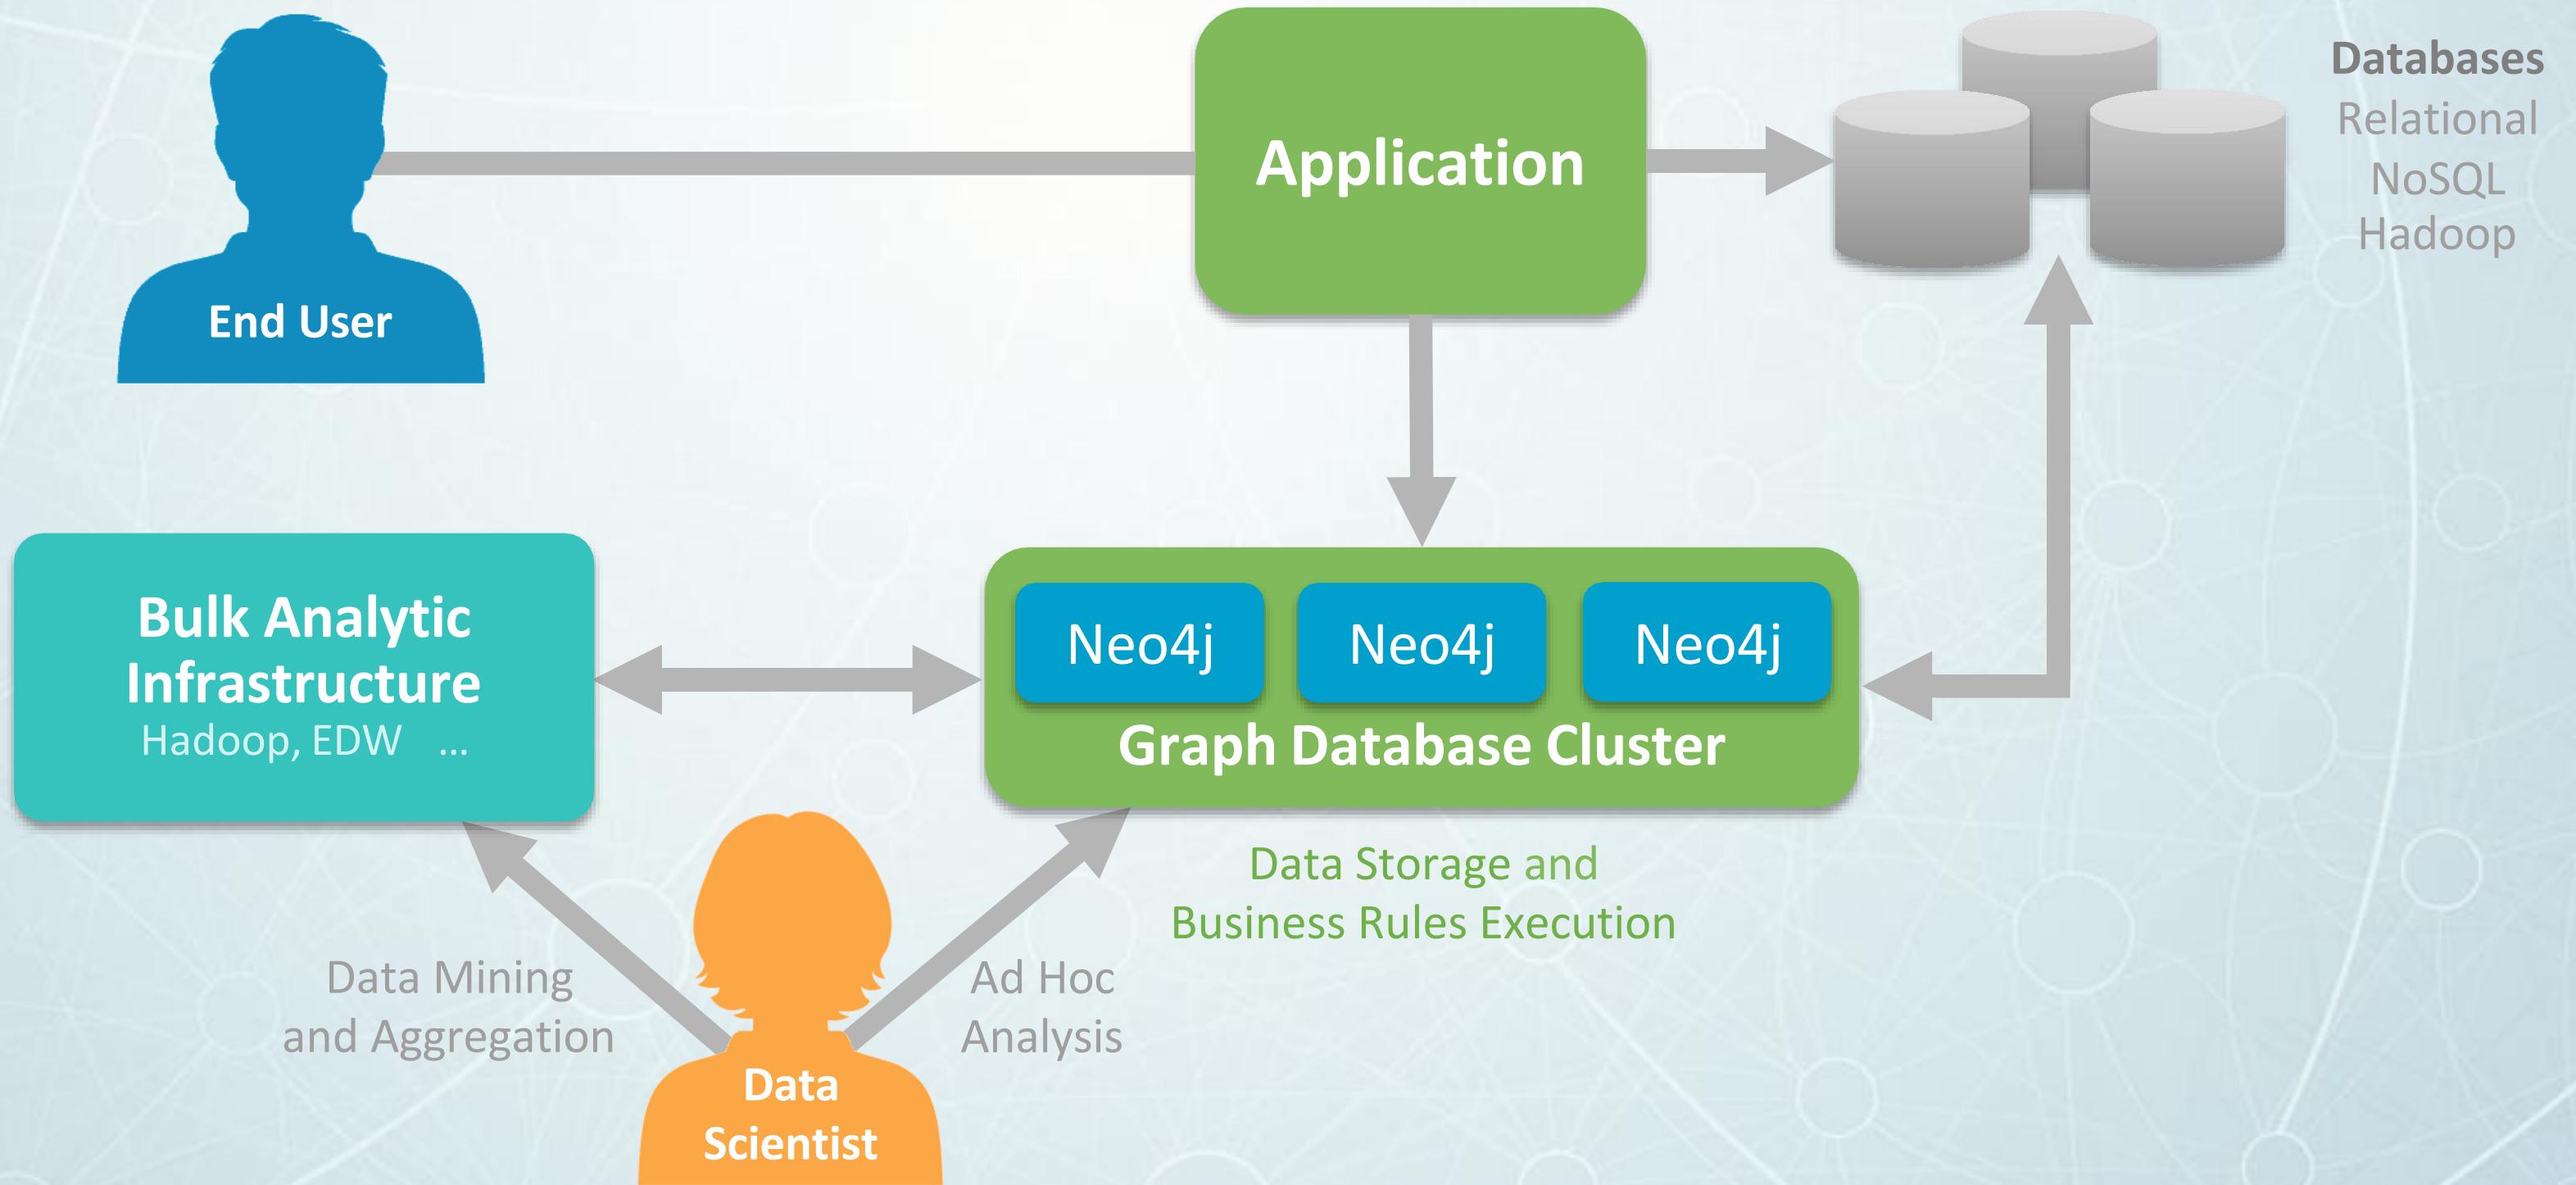
>>=



# Native Server-Side Extensions



# Architectural Options



# Day in the Life of a RDBMS Developer



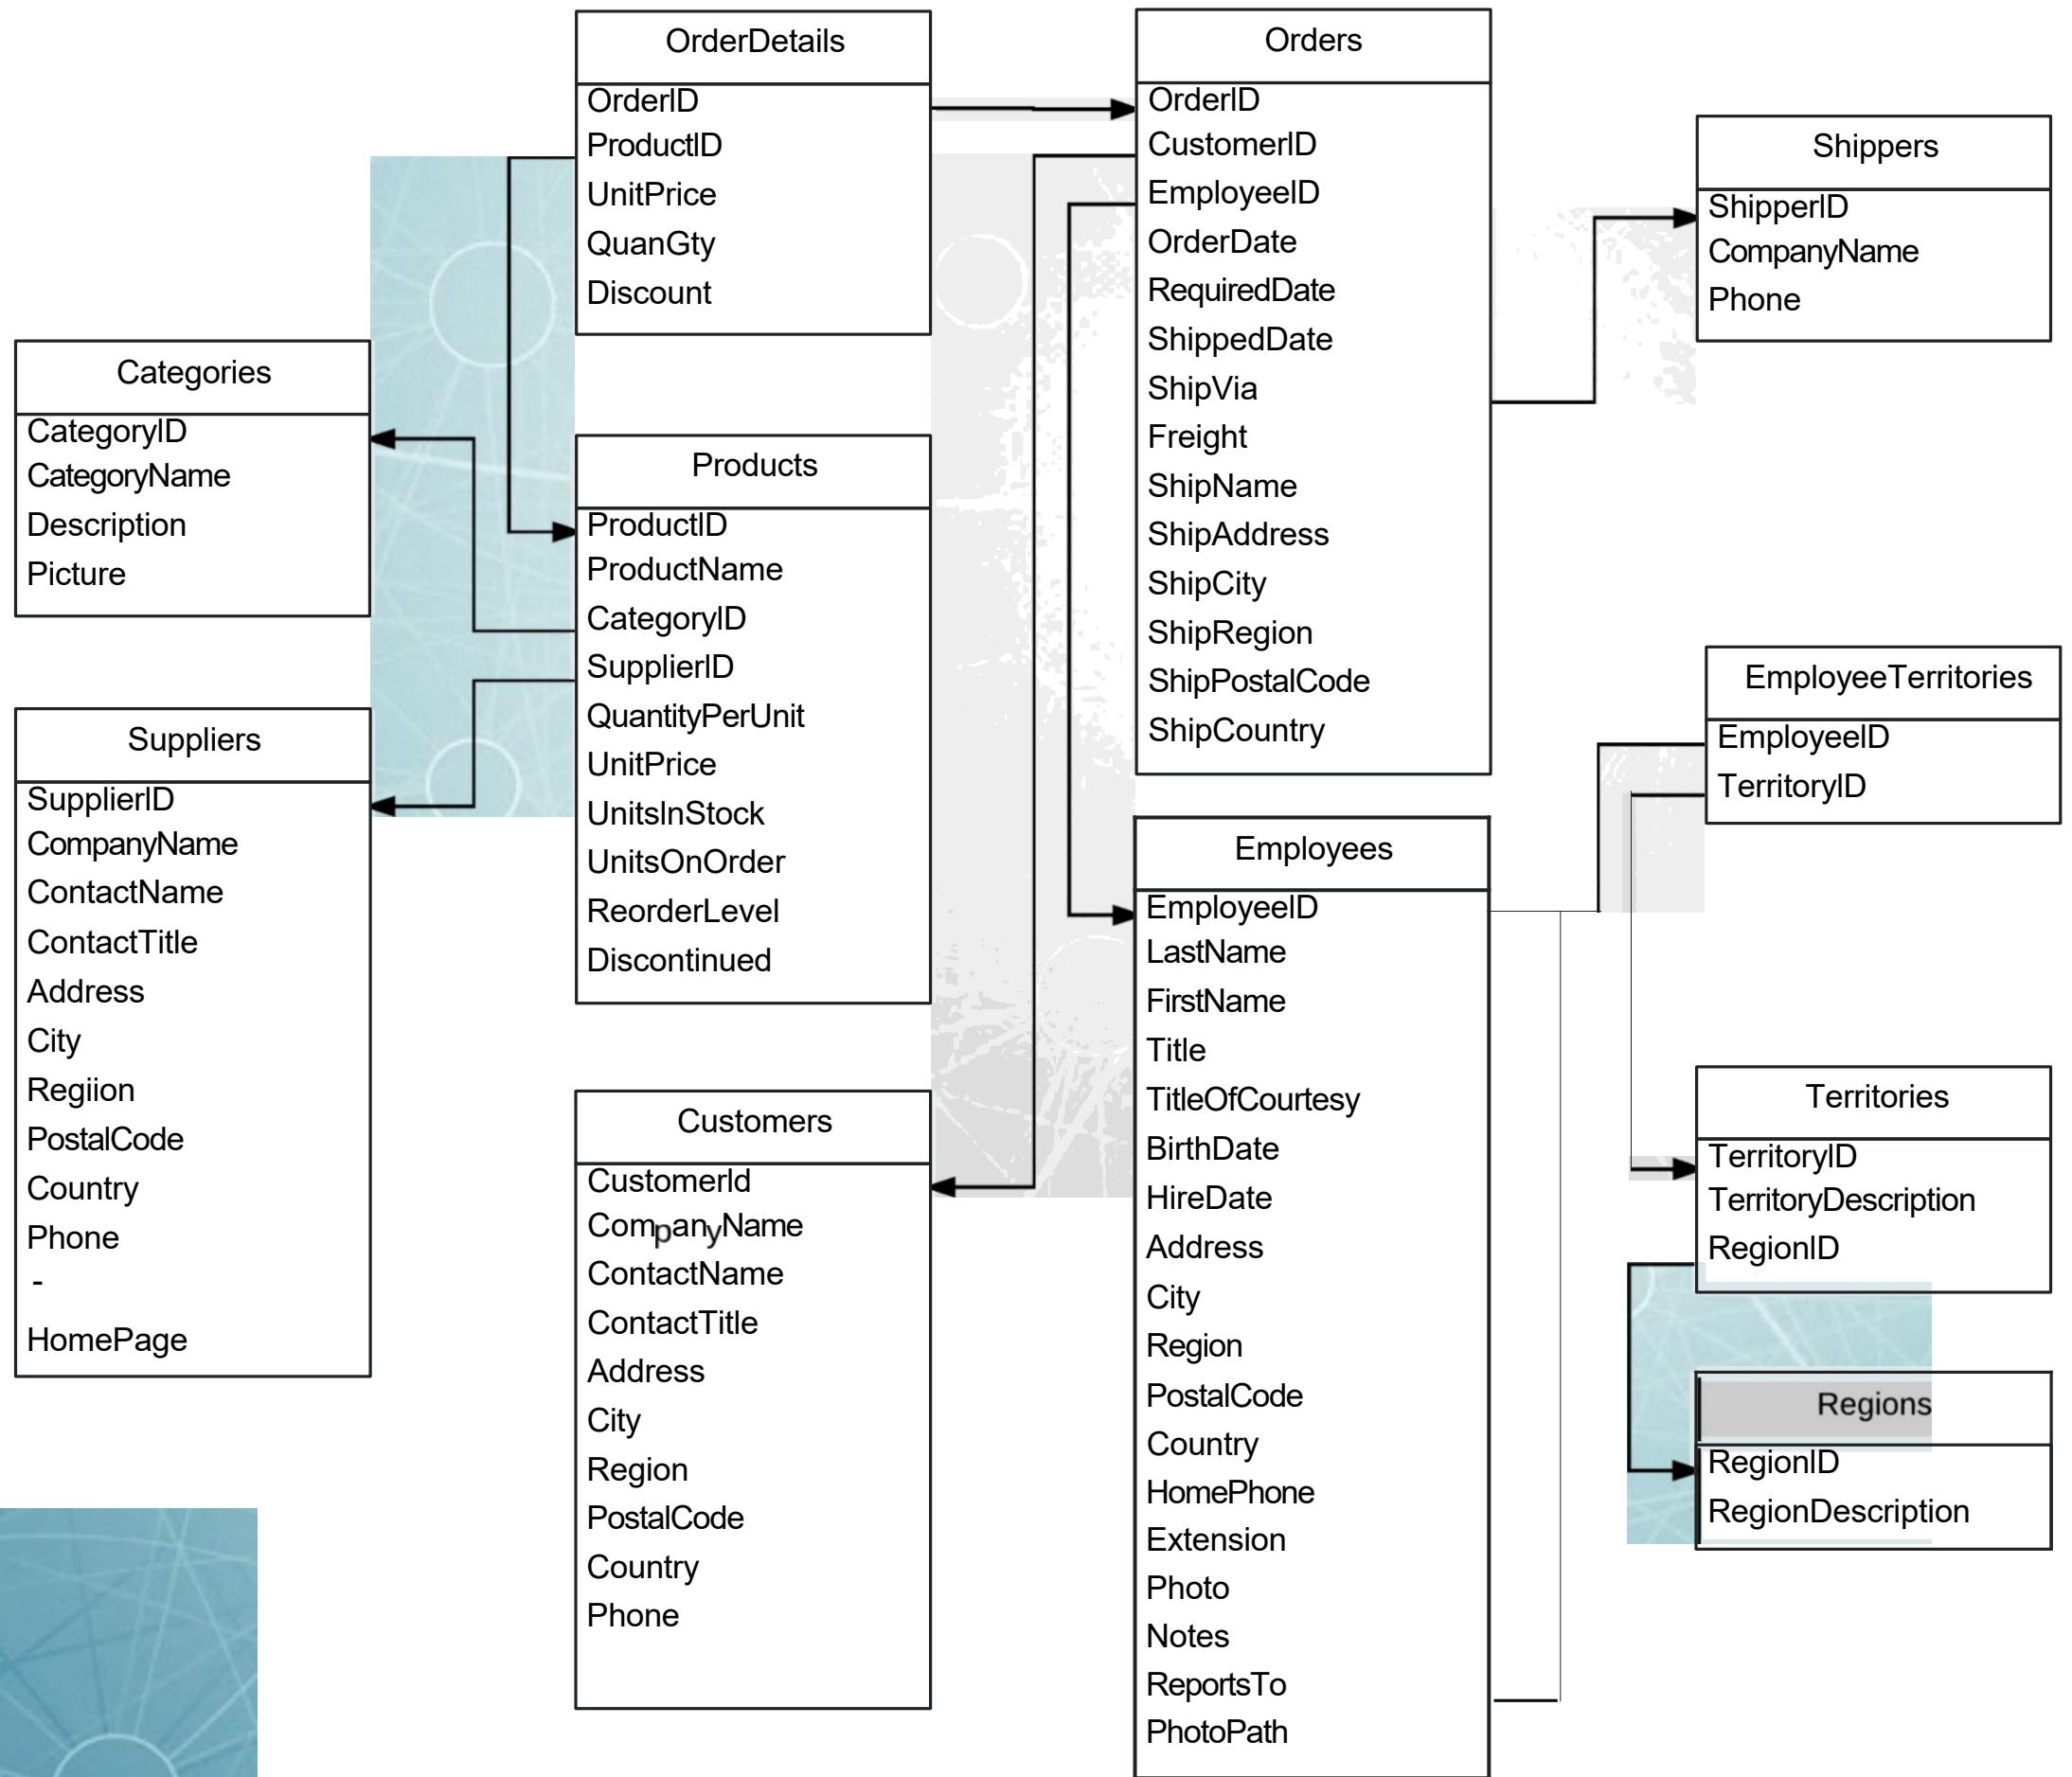
# **SQL Pains**

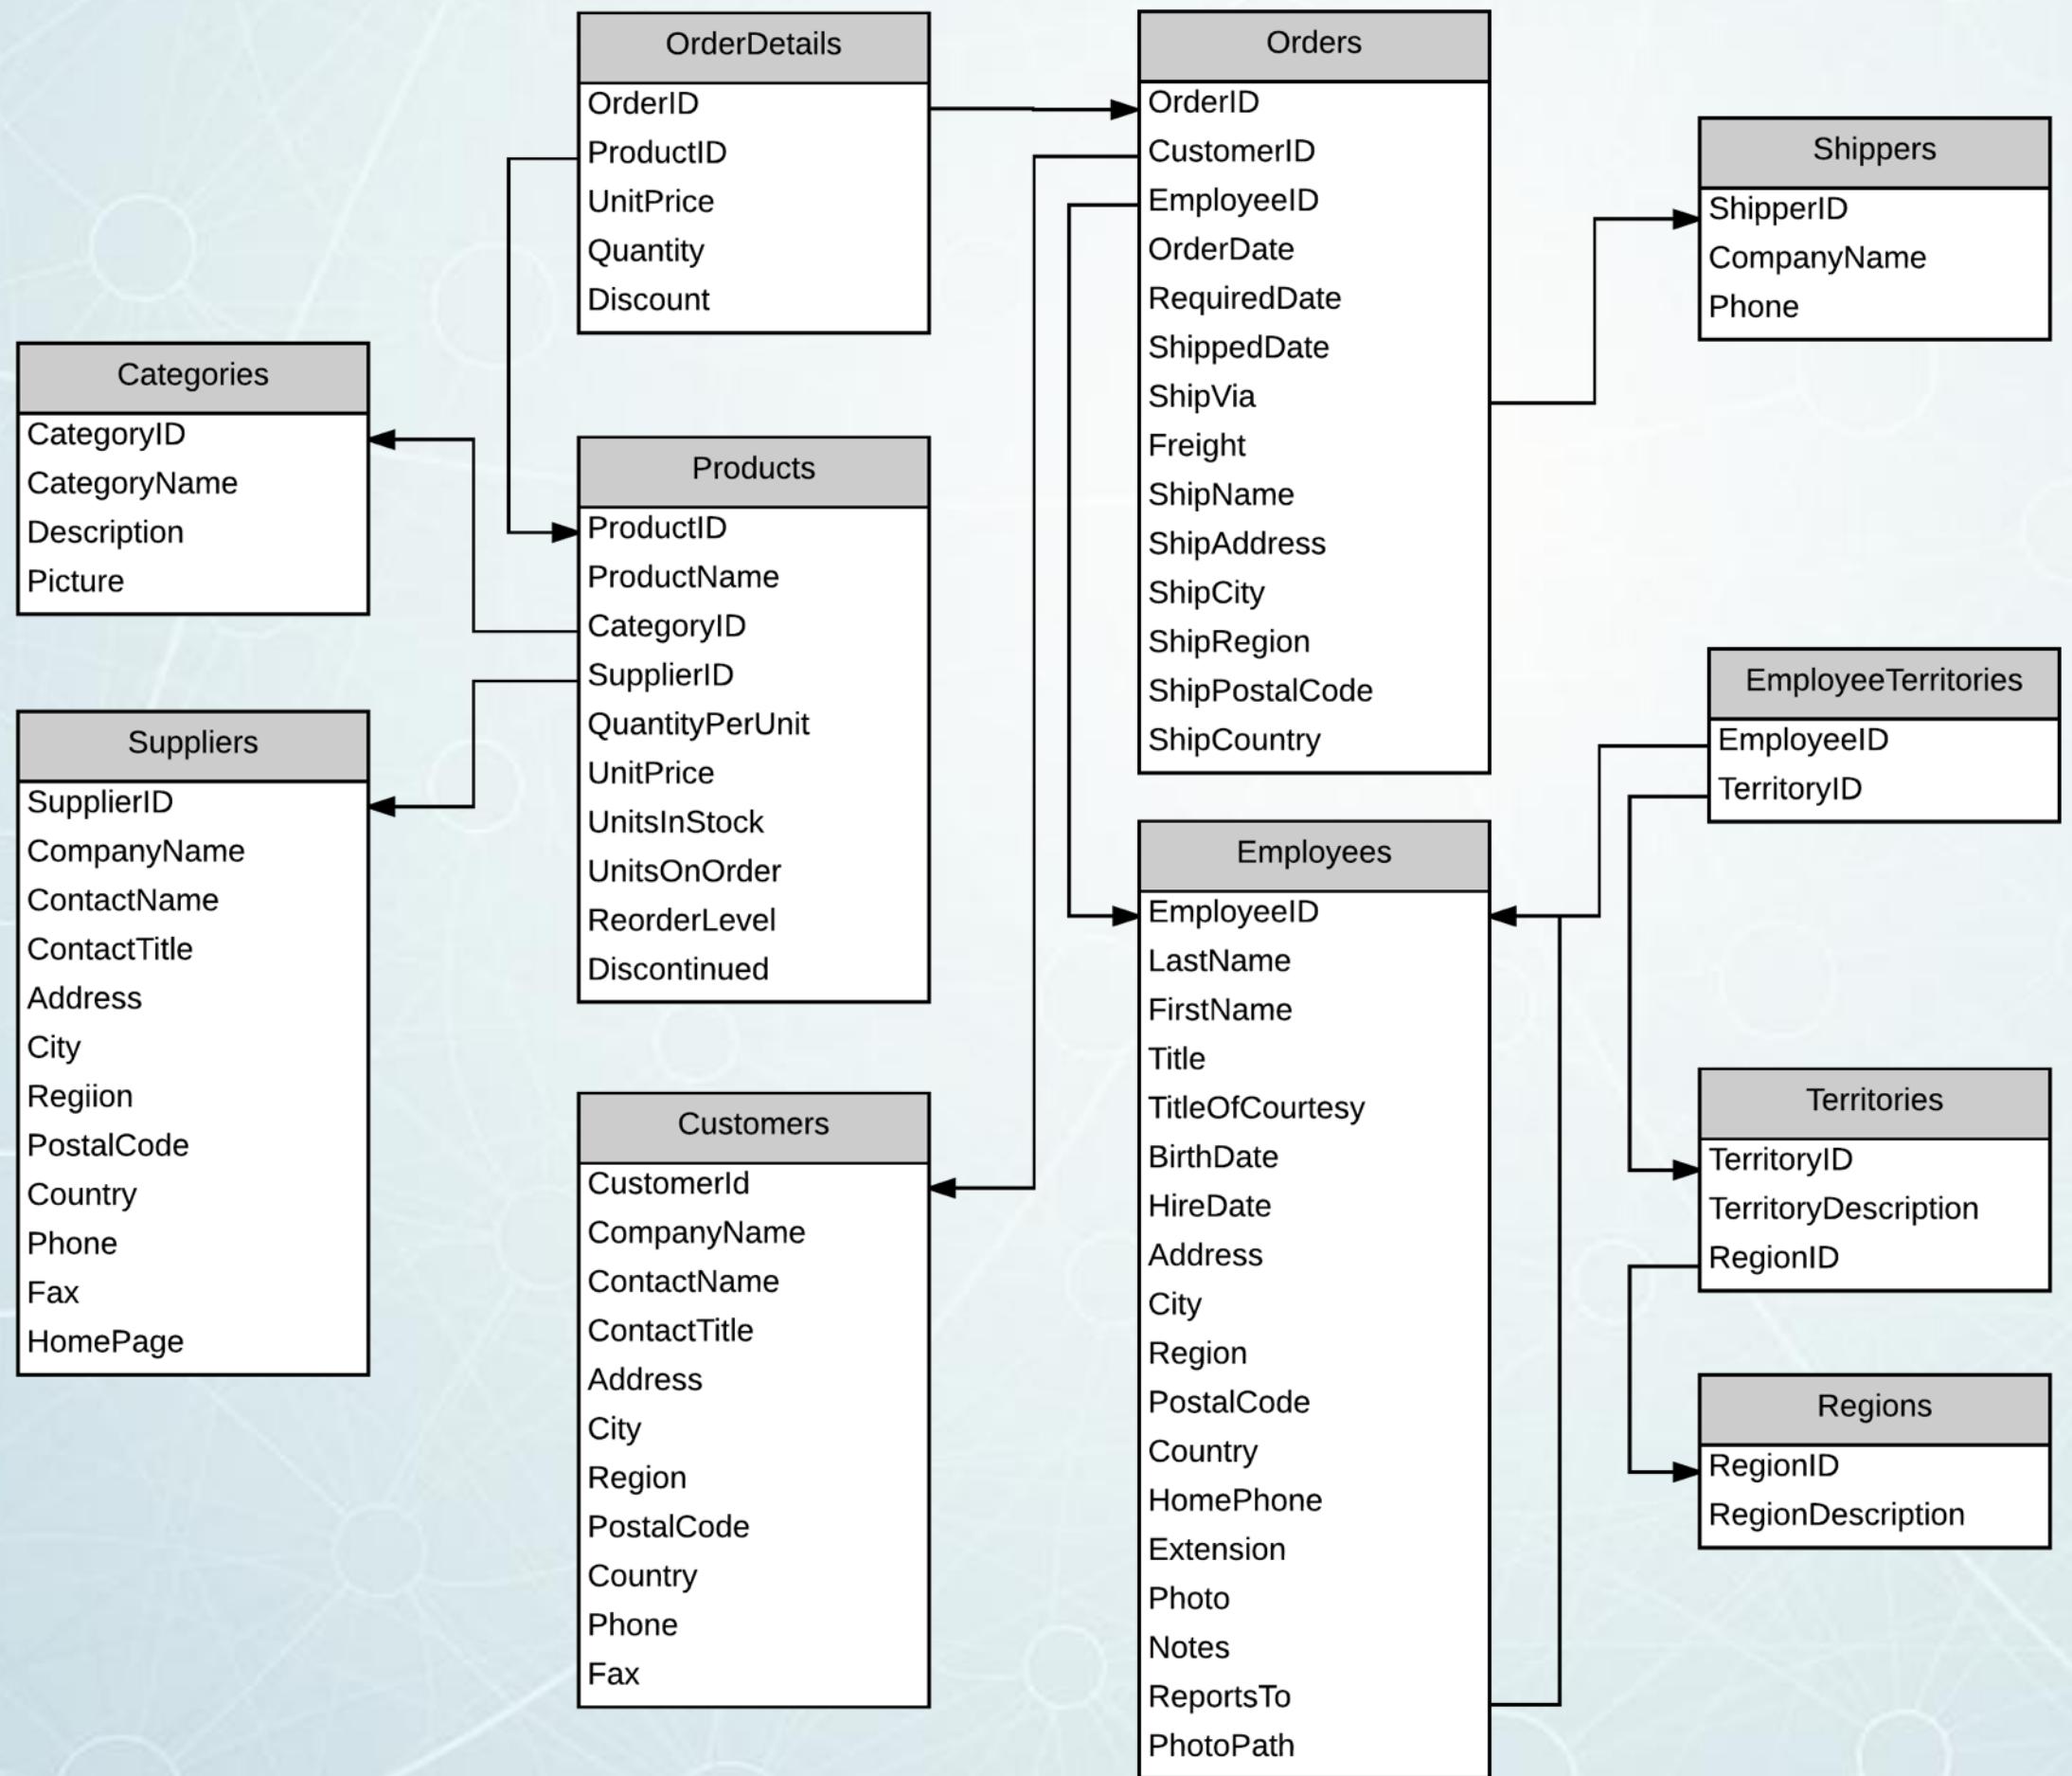
- Complex to model and store relationships
- Performance degrades with increases in data
- Queries get long and complex
- Maintenance is painful

# Graph Gains

- Easy to model and store relationships
- Performance of relationship traversal remains constant with growth in data size
- Queries are shortened and more readable
- Adding additional properties and relationships can be done on the fly - no migrations

# **FROM RDBMS TO GRAPHS**





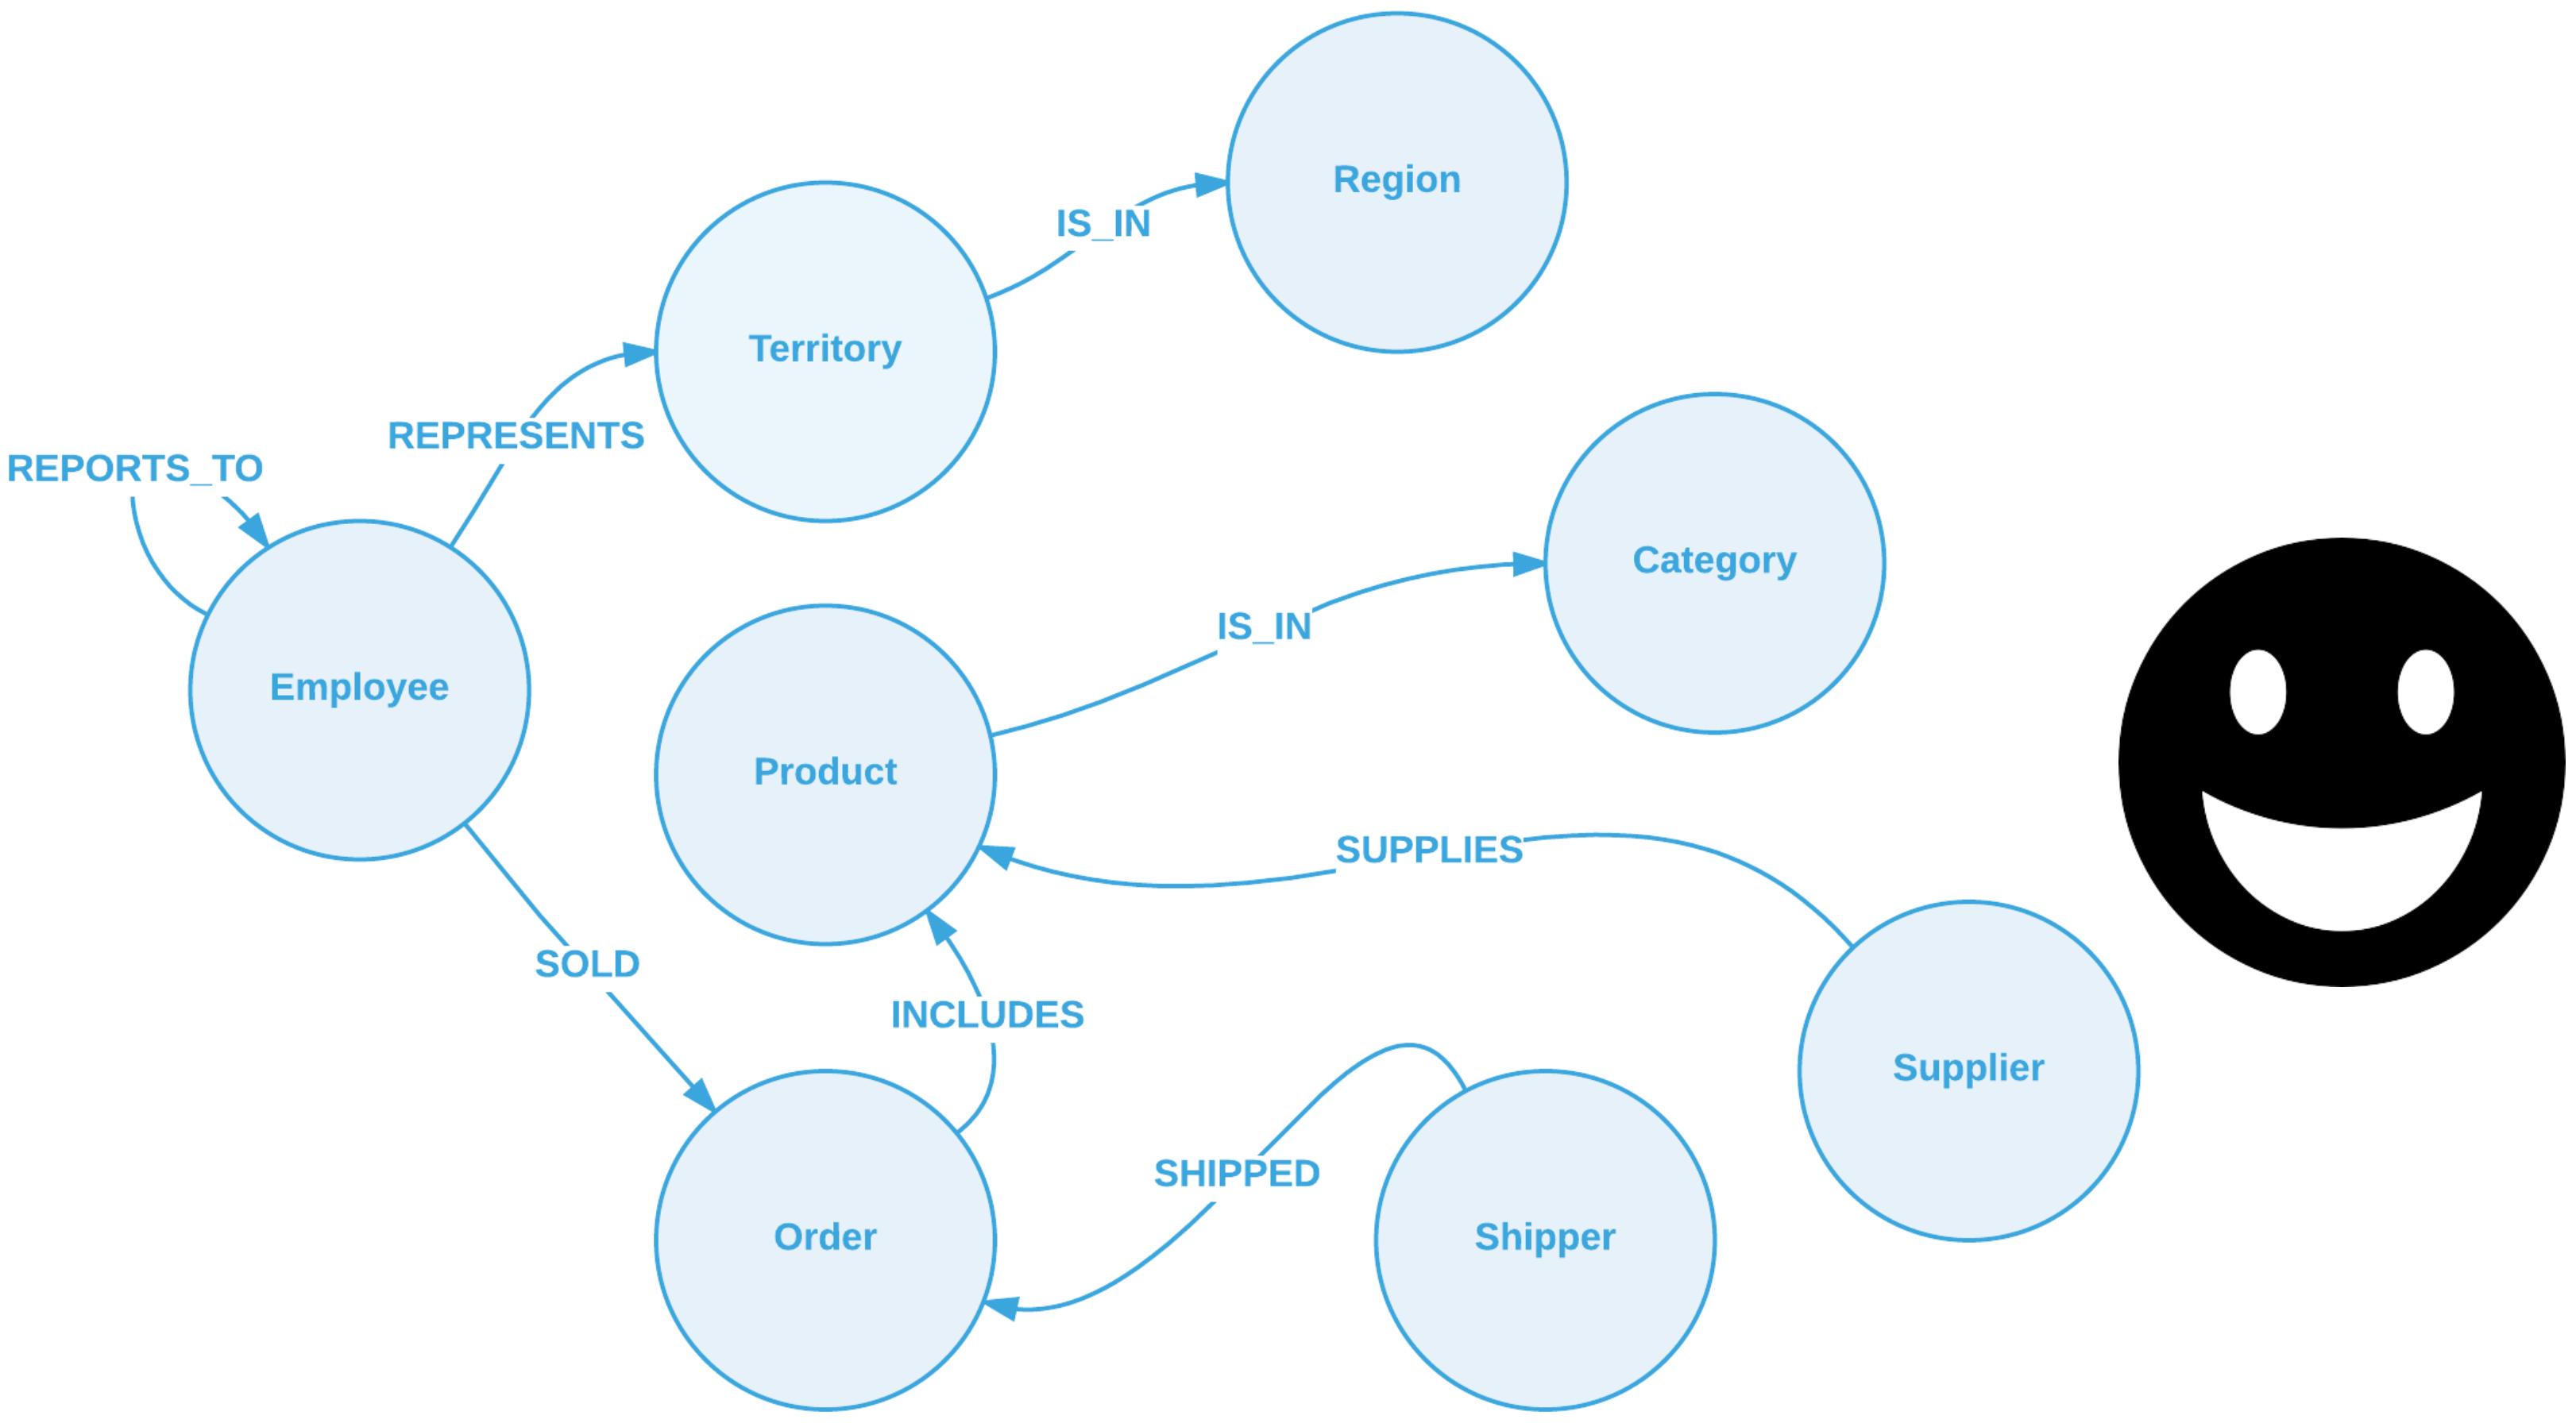
# Northwind



# Northwind - the canonical RDBMS Example

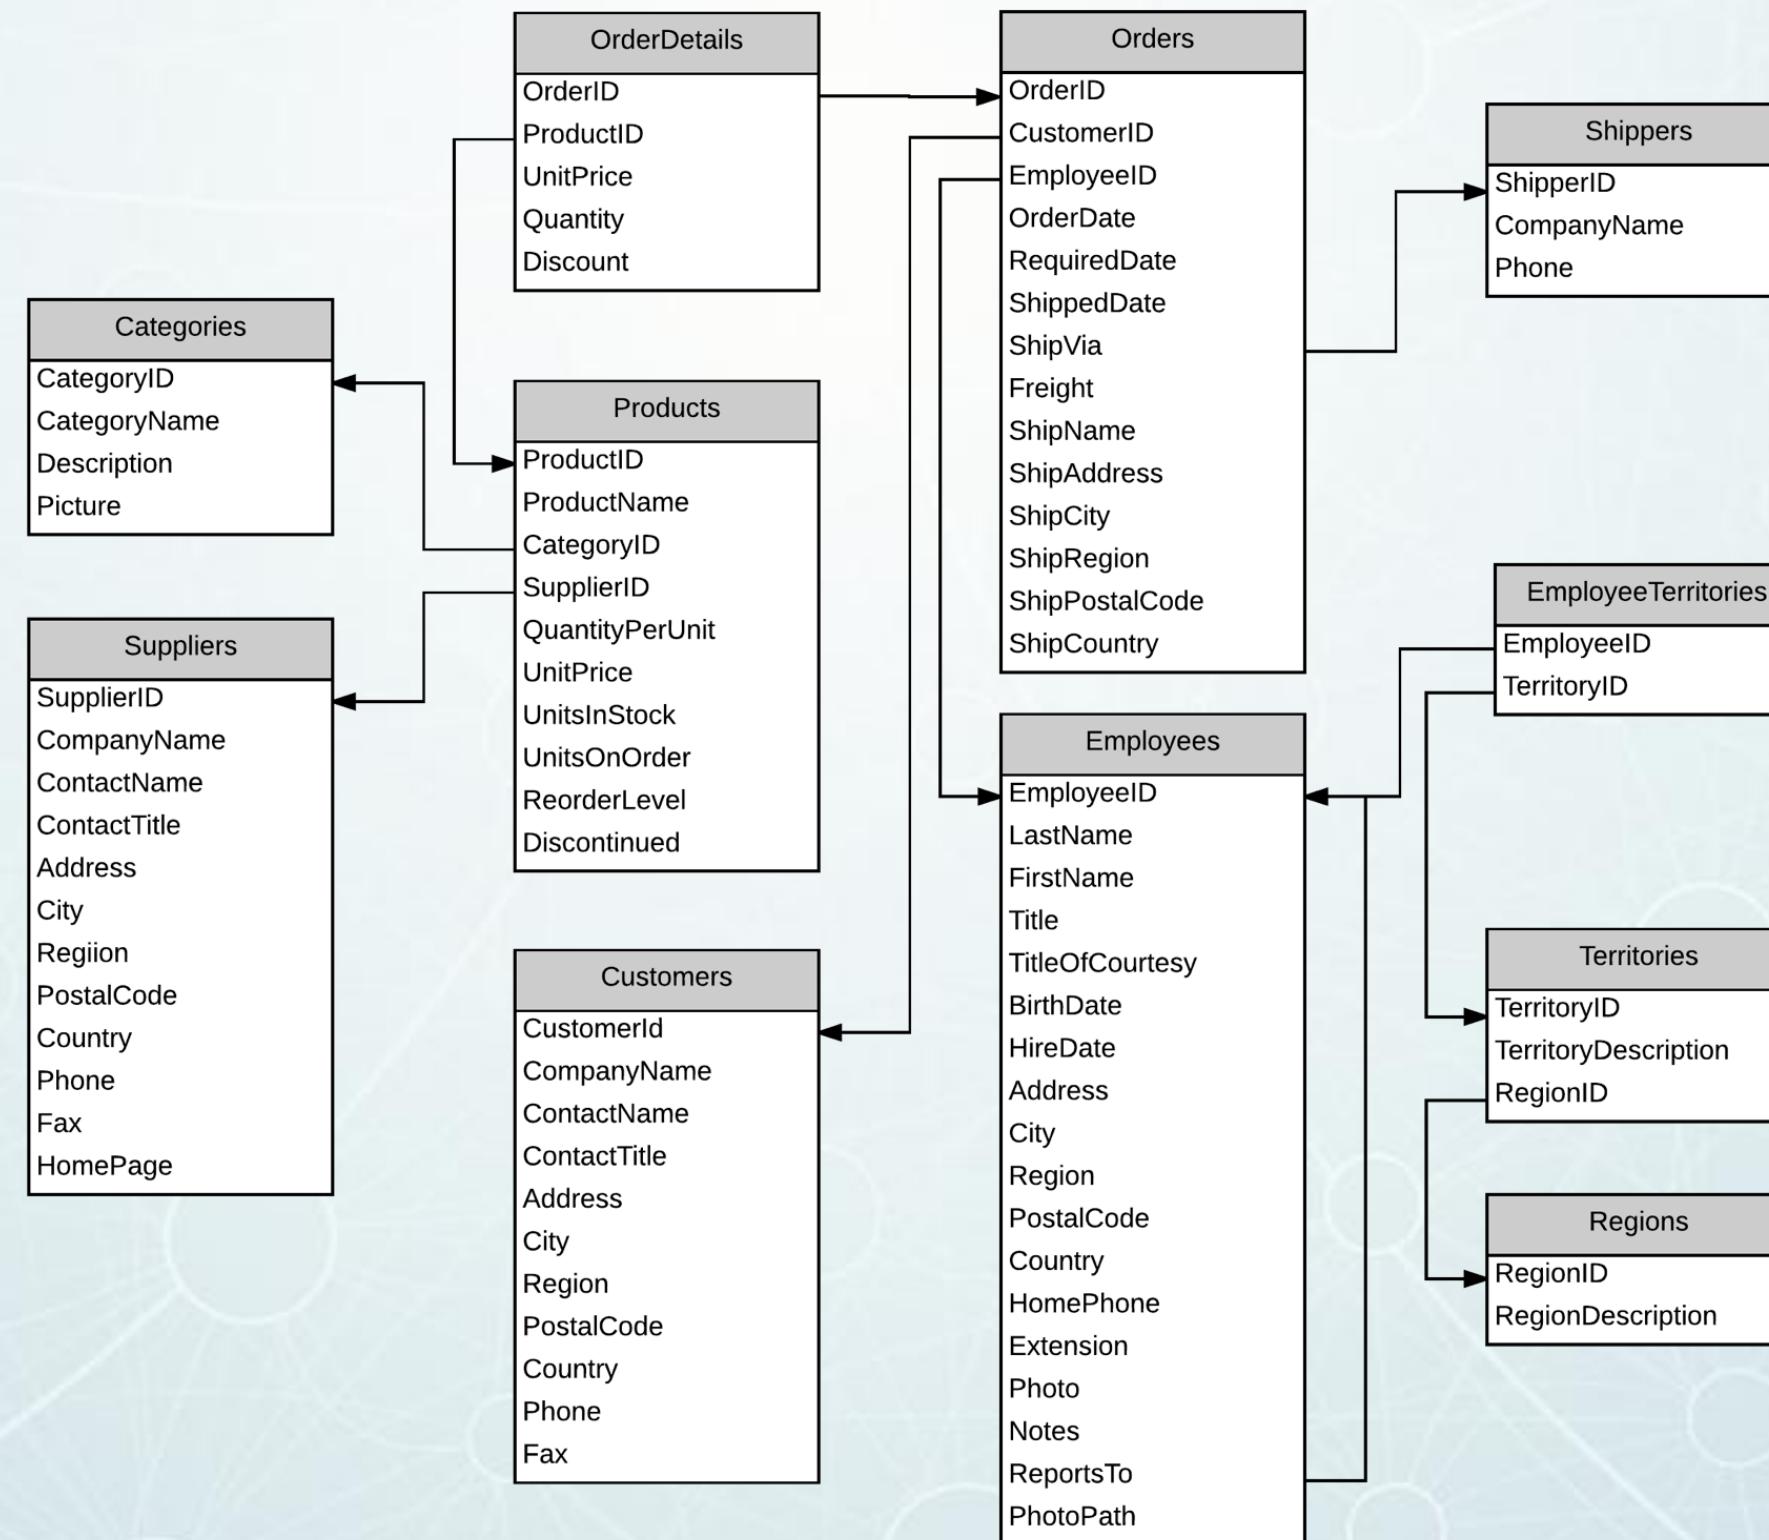


**( $\triangle$ )-[:TO]>(Graph)**

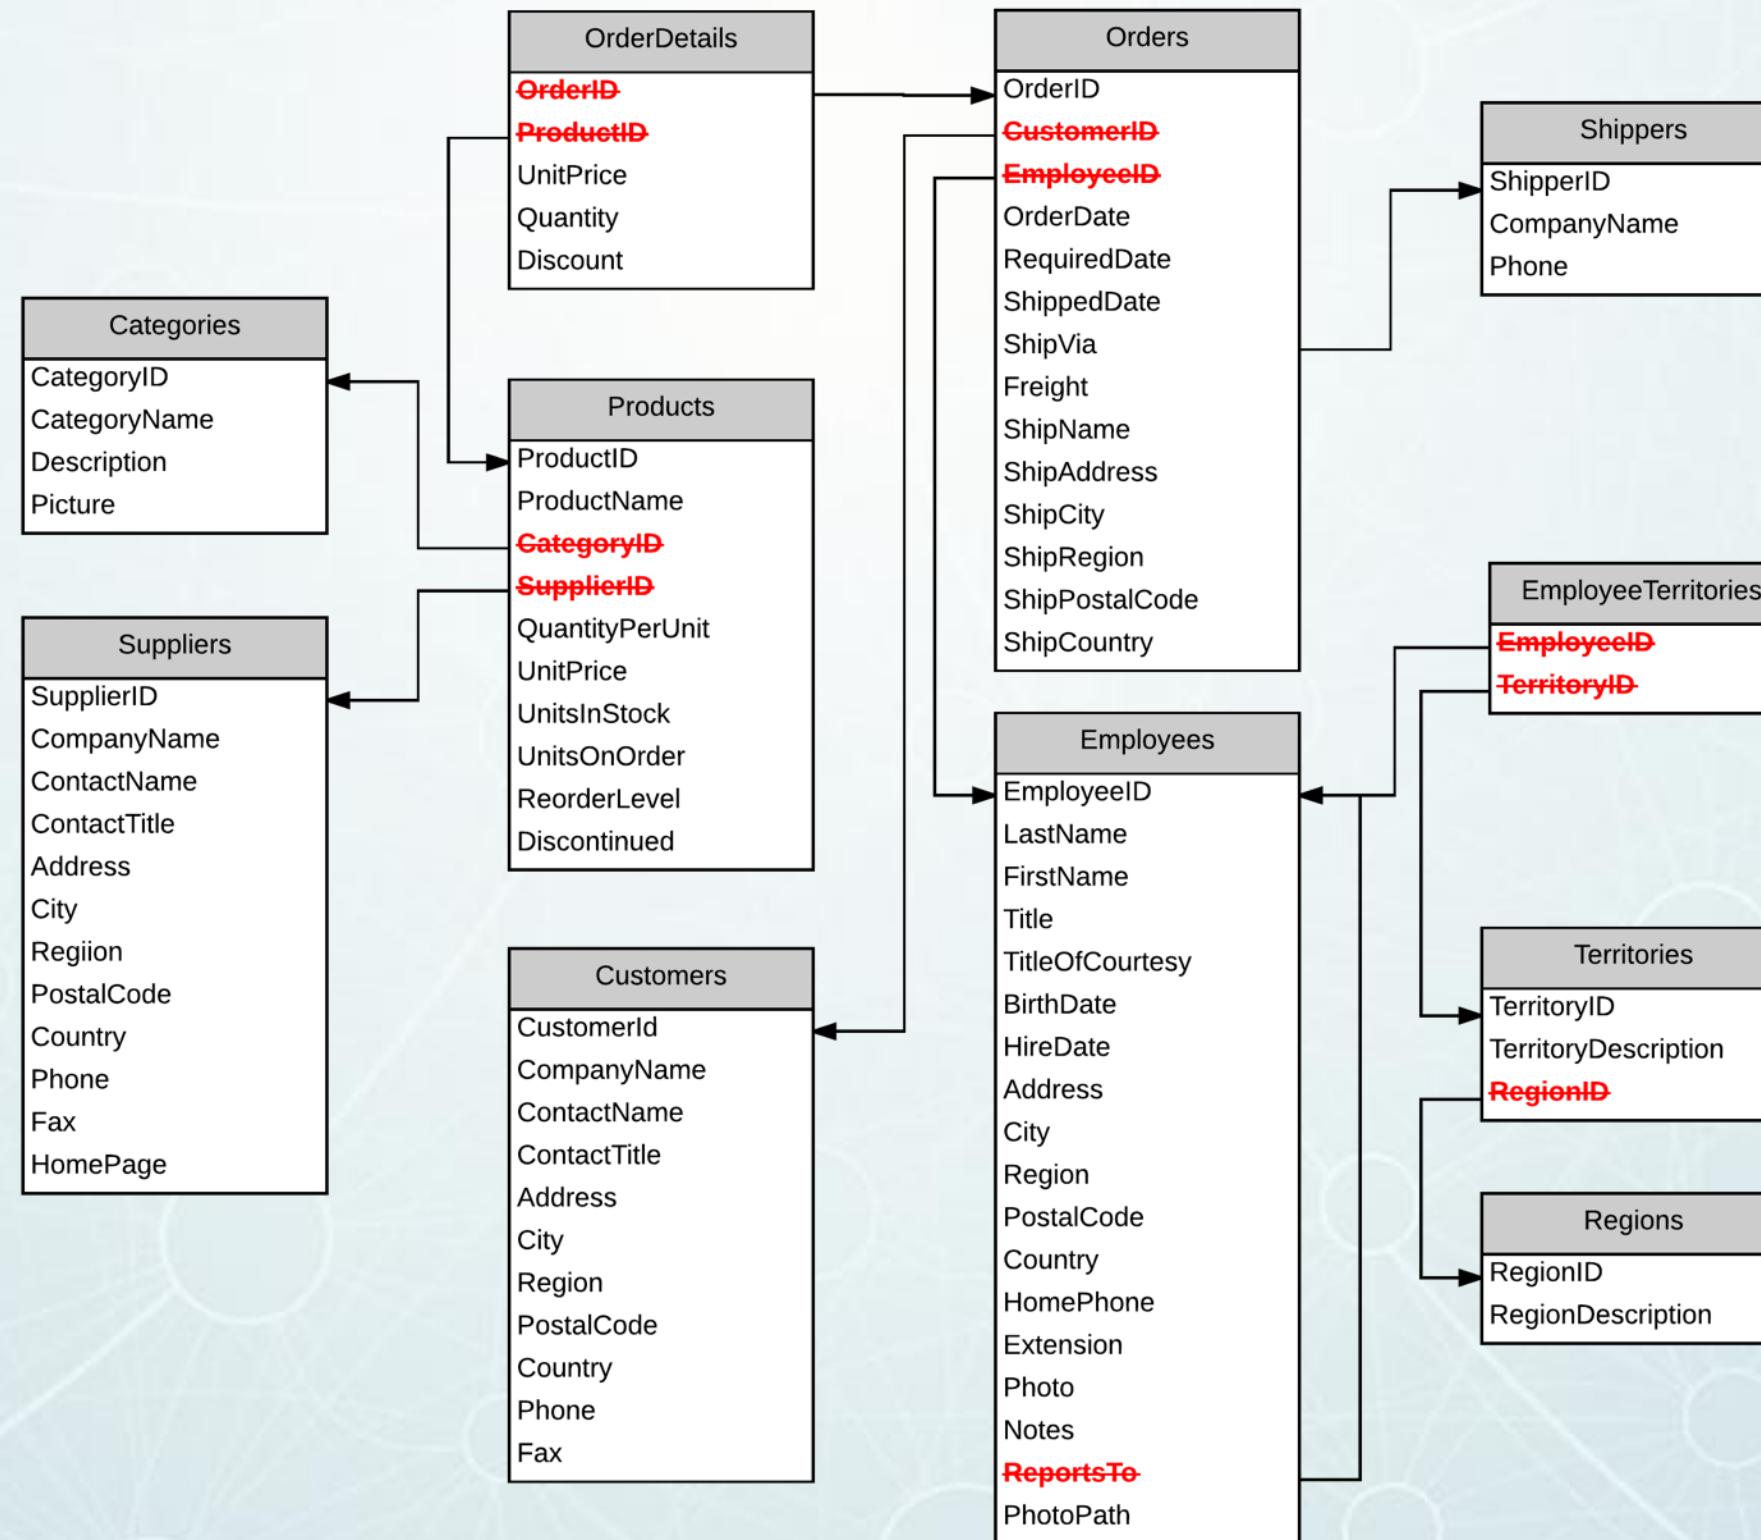


(↑)-[:IS\_BETTER\_AS]->(Graph)

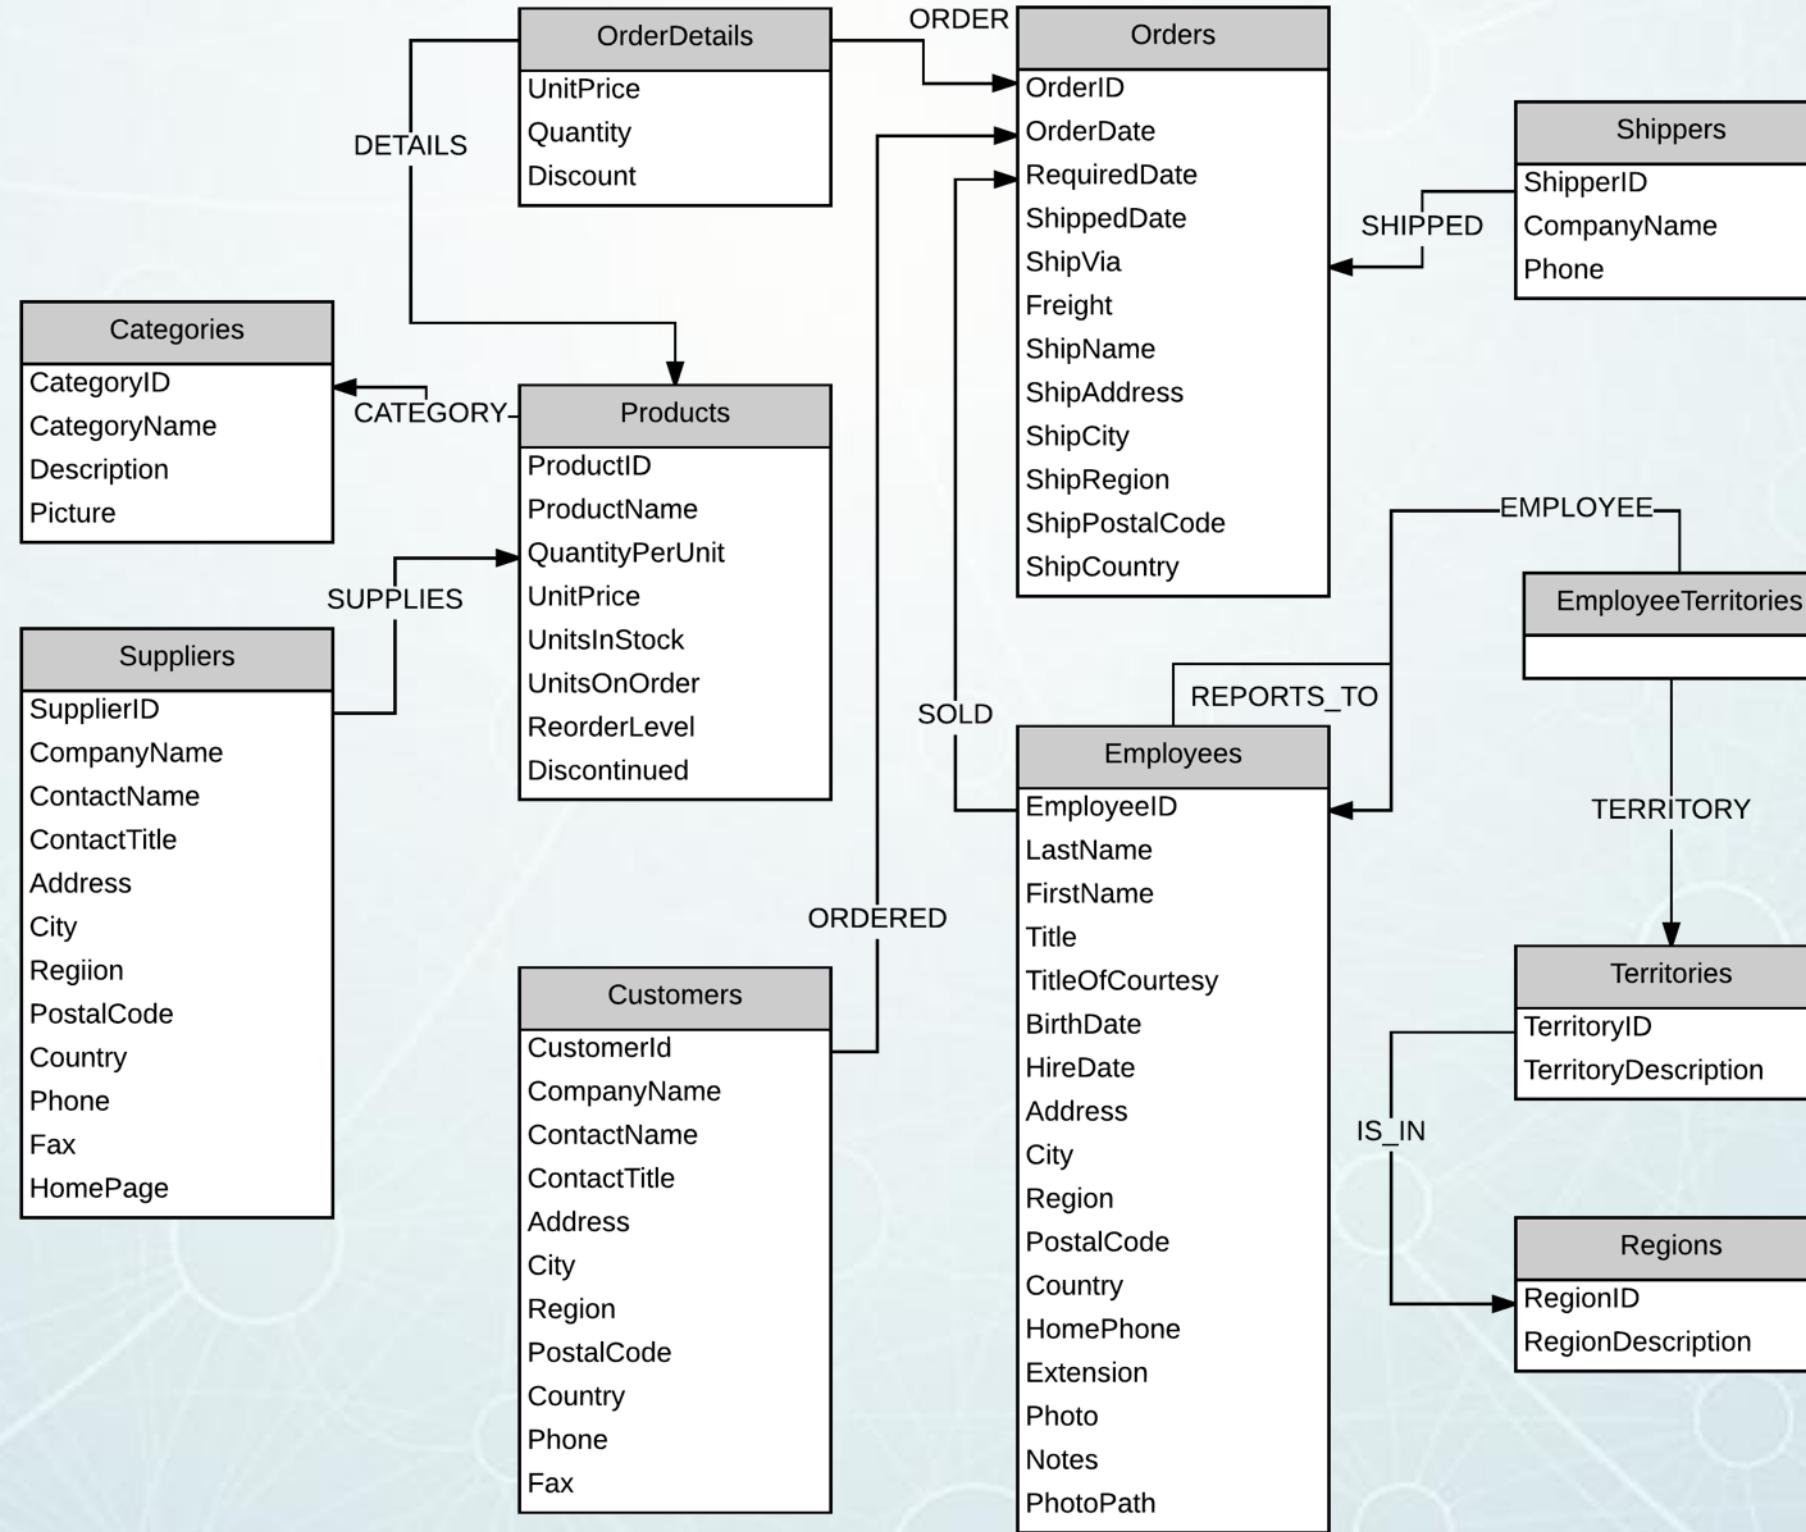
# Starting with the ER Diagram



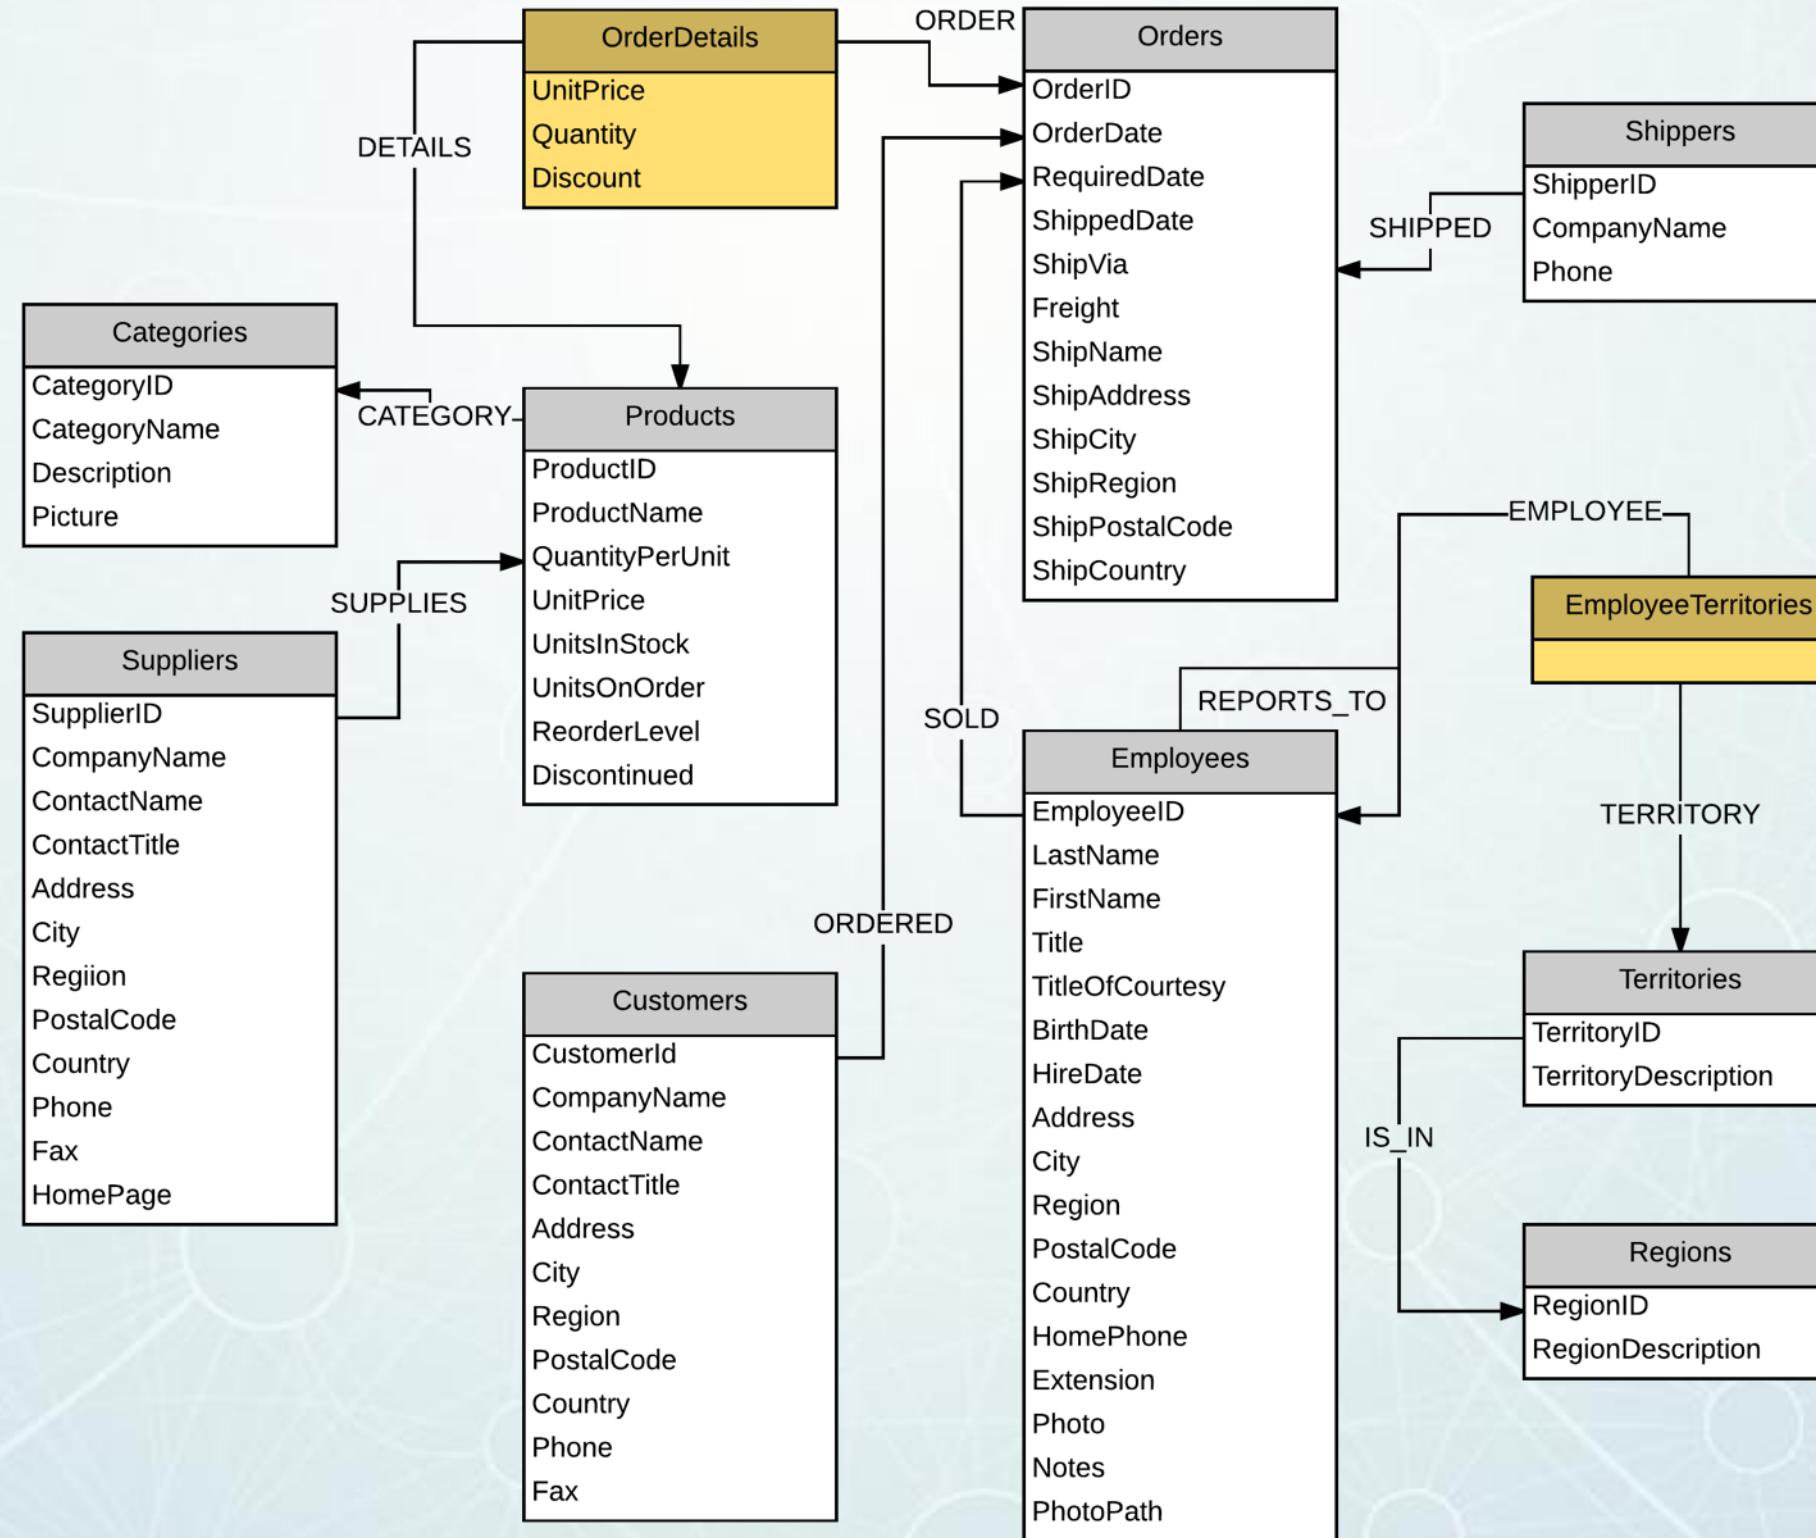
# Locate the Foreign Keys



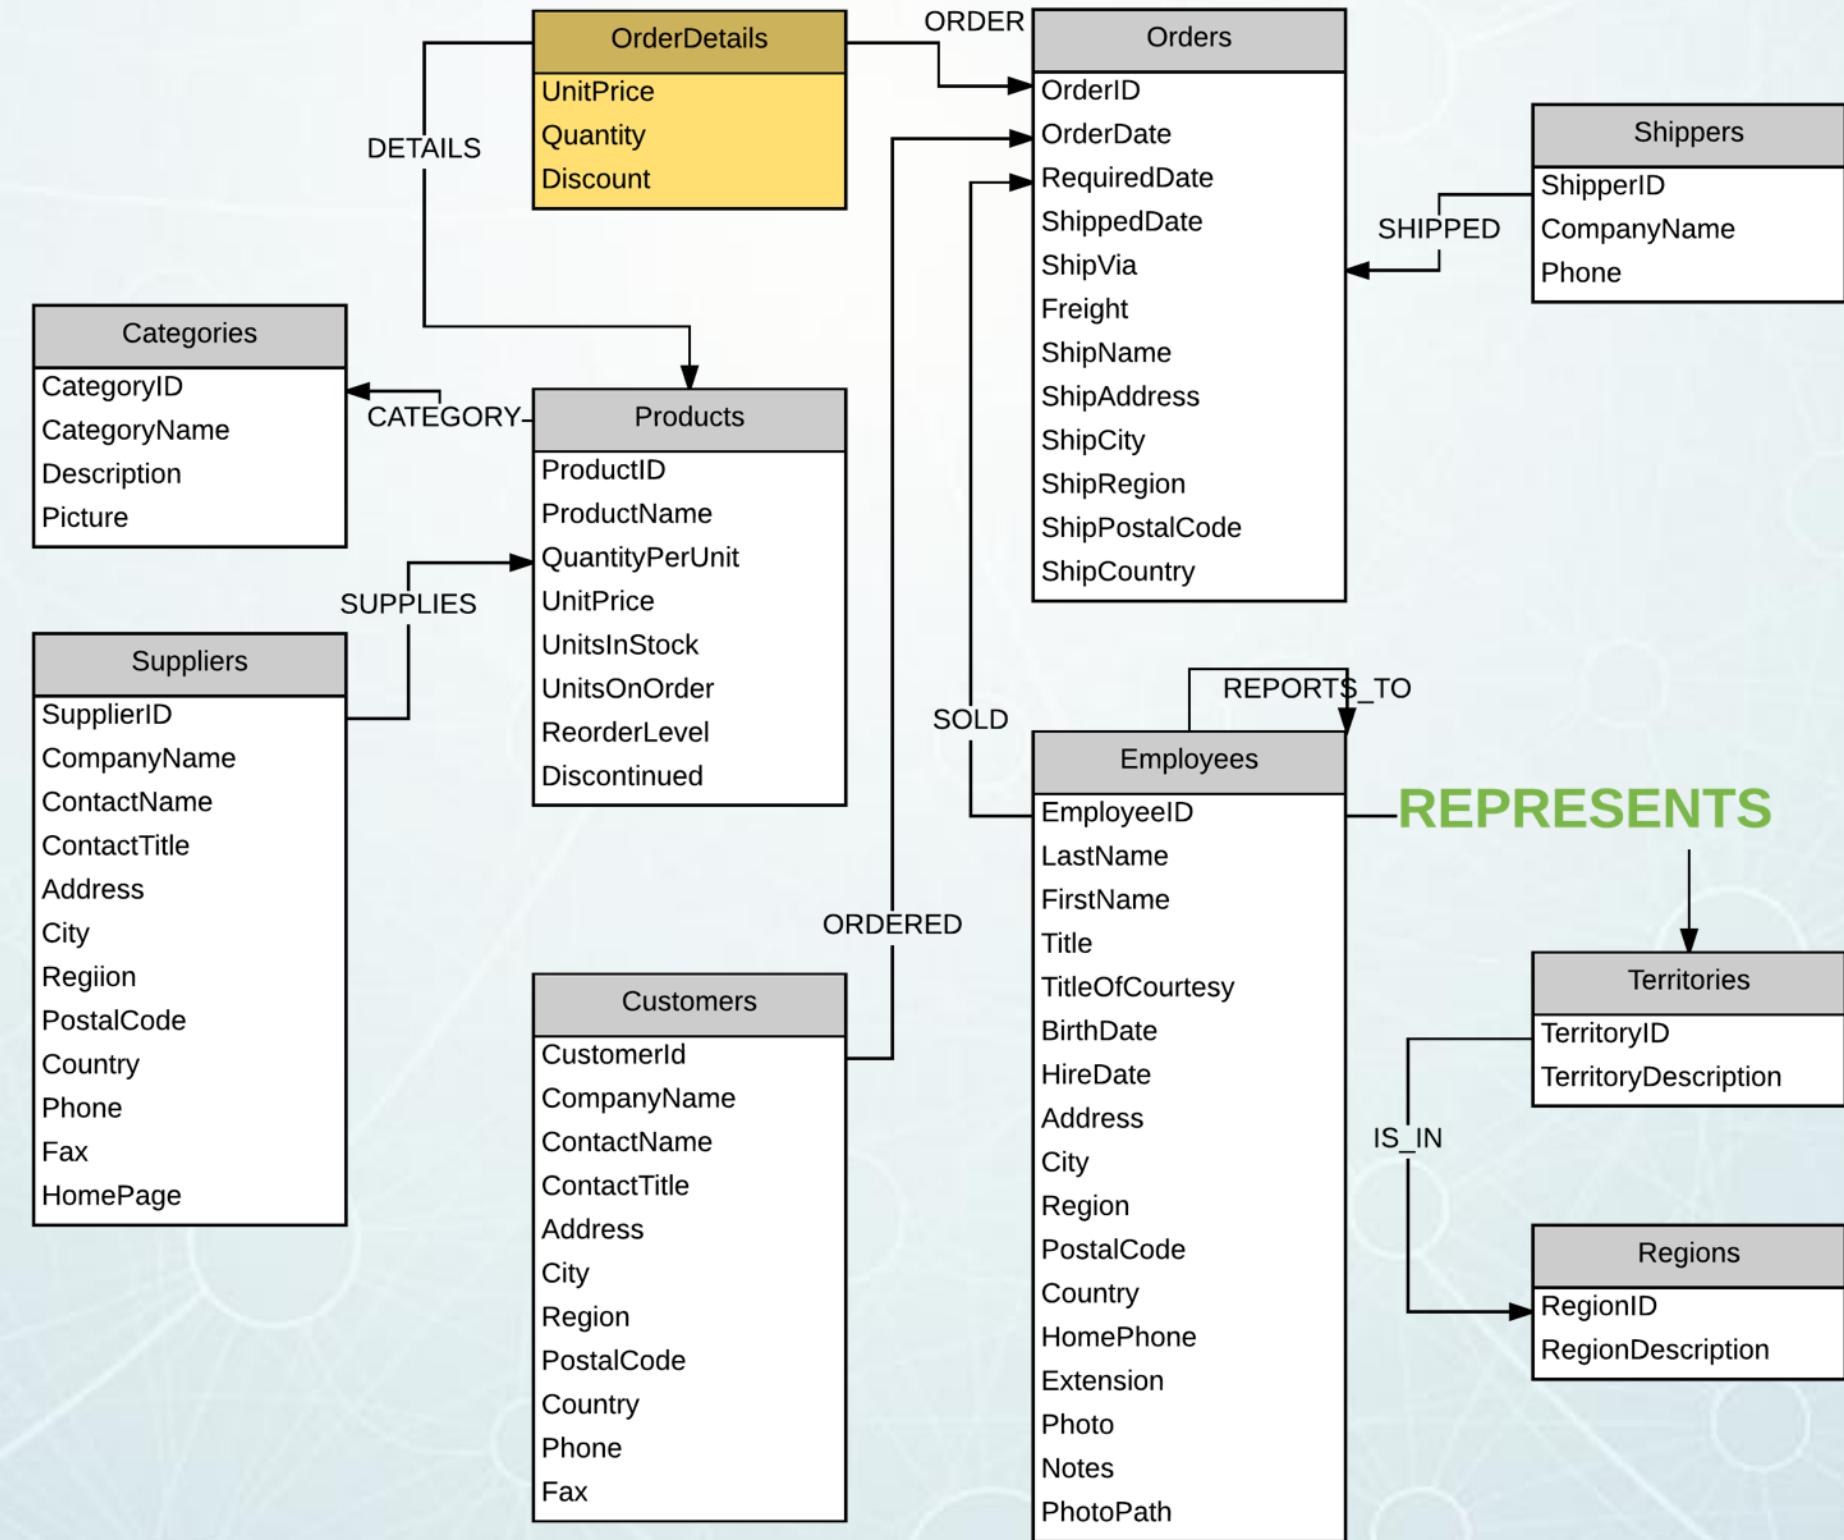
# Drop the Foreign Keys



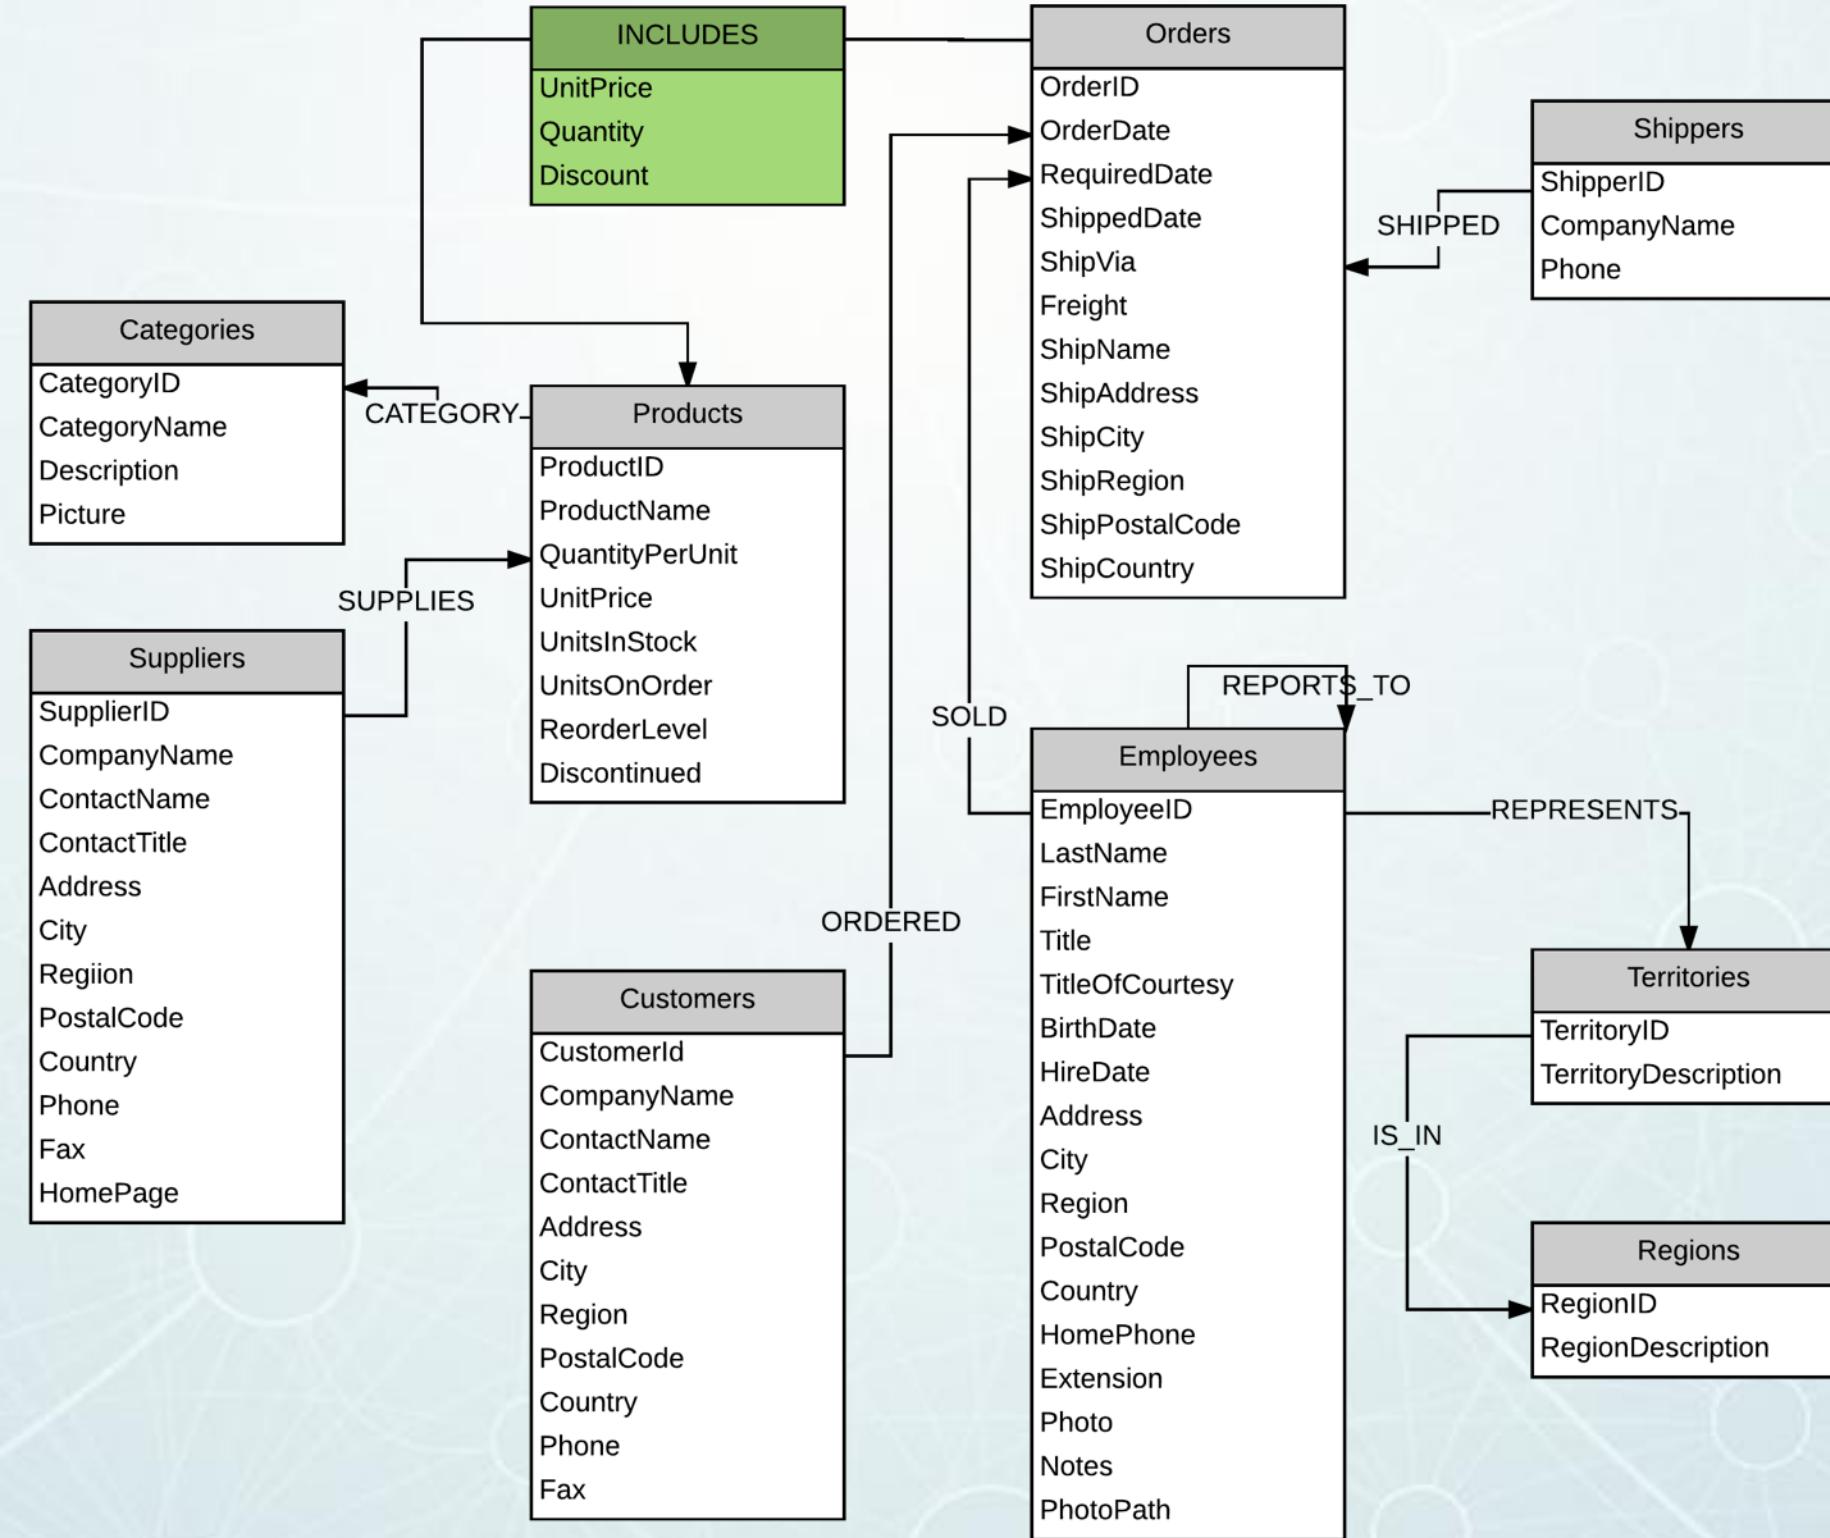
# Find the JOIN Tables



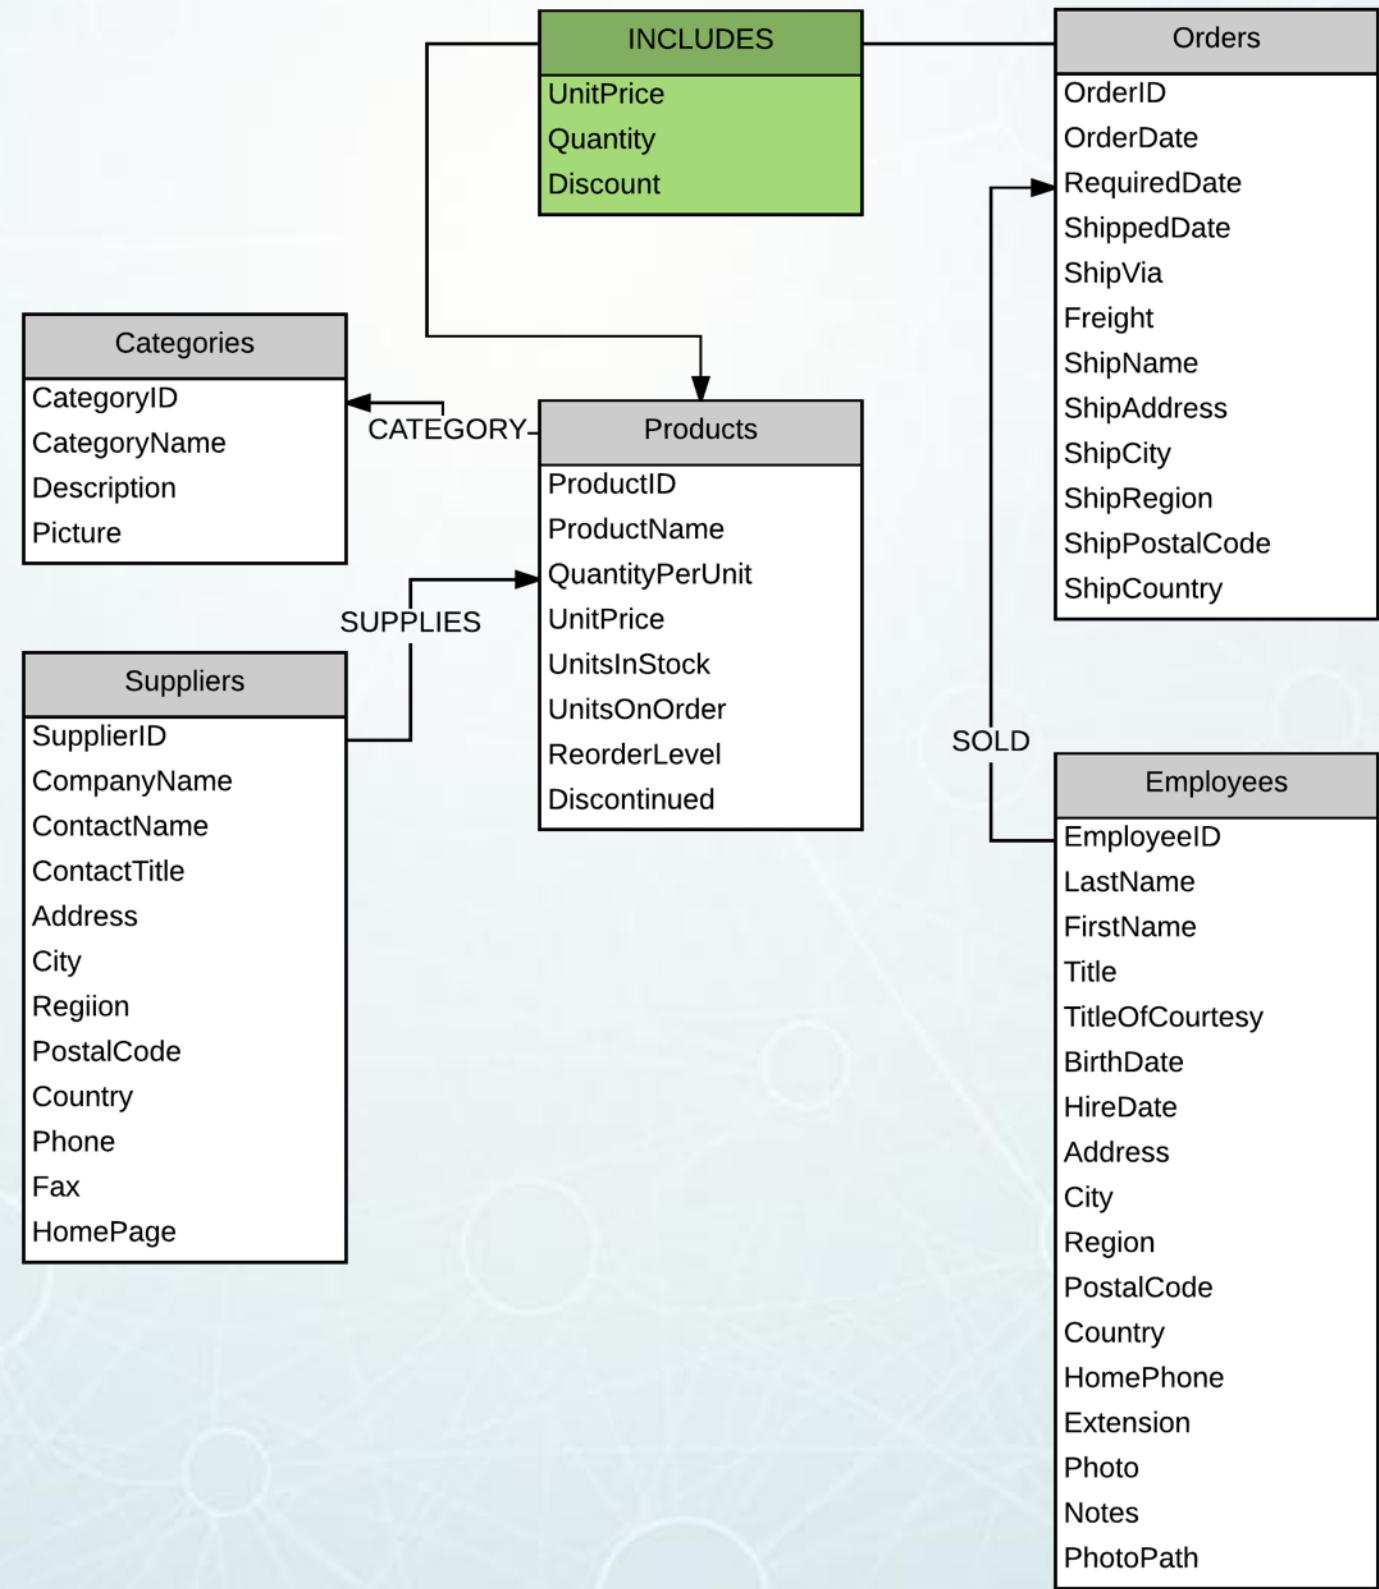
# (Simple) JOIN Tables Become Relationships



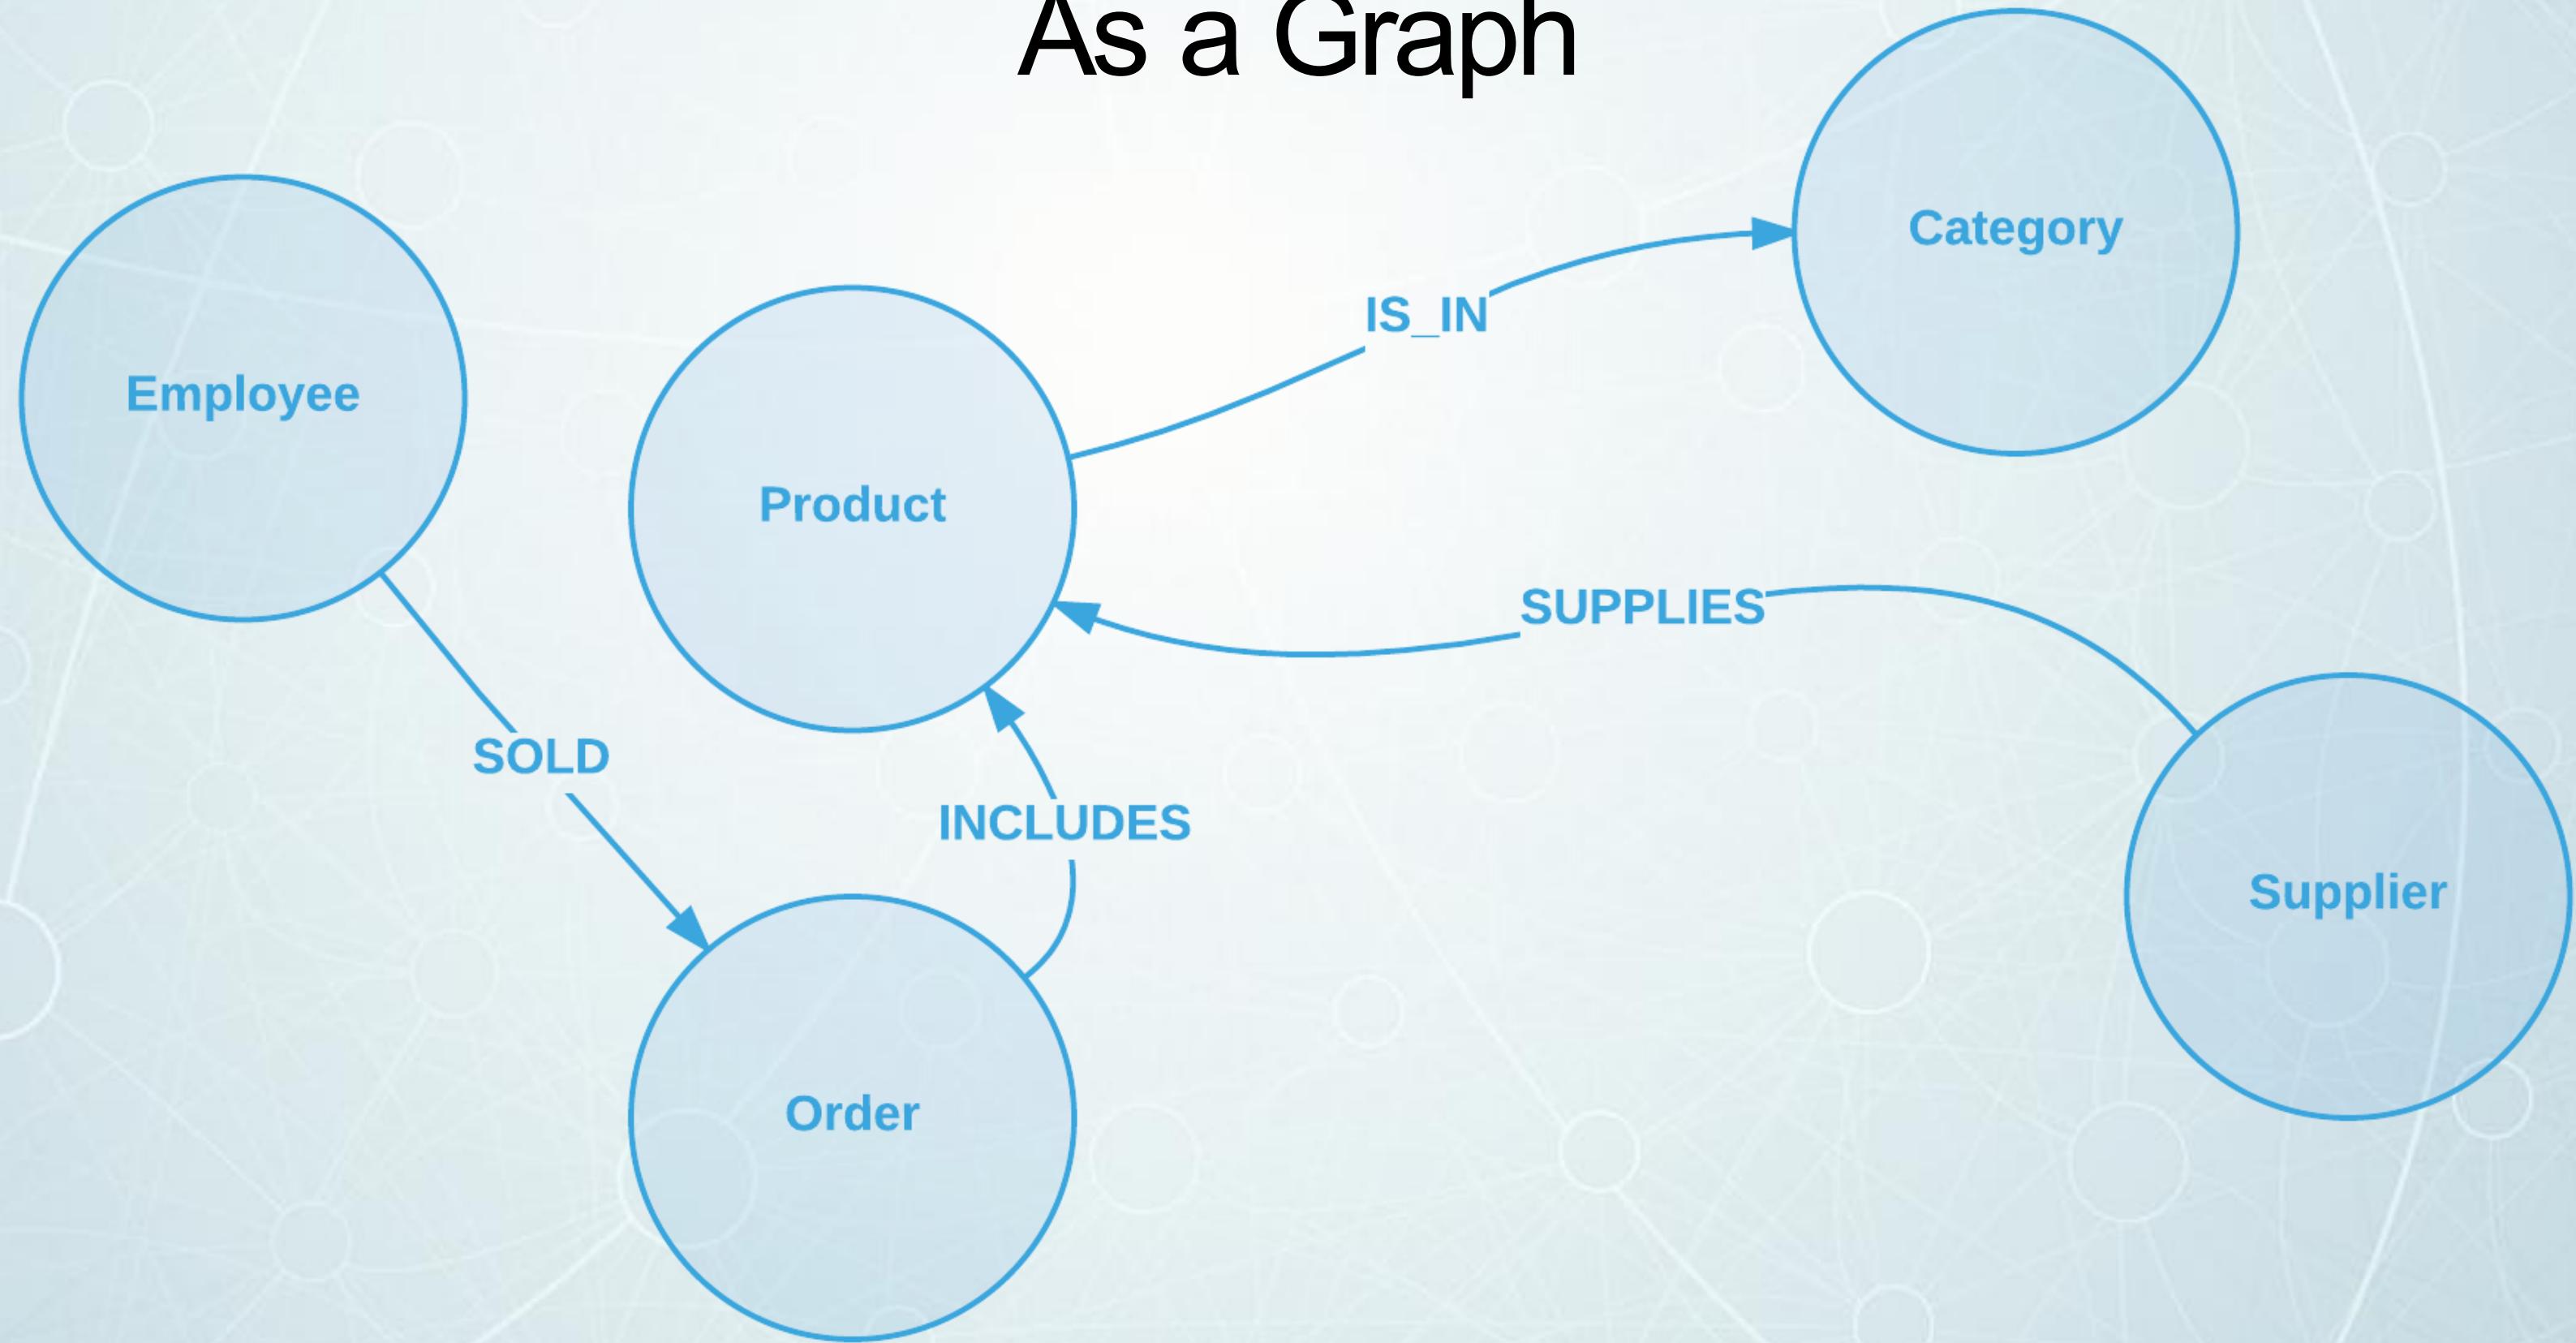
# Attributed JOIN Tables -> Relationships with Properties



# Querying a Subset Today



# As a Graph



# **QUERYING THE GRAPH**



# using openCypher

# Property Graph Model



```
CREATE (:Employee{ firstName:"Steven" }) -[:REPORTS_TO]-> (:Employee{ firstName:"Andrew" })
```

LABEL

PROPERTY

LABEL

PROPERTY

# Who do people report to?

```
MATCH  
  (e:Employee)<- [ :REPORTS_TO] - (sub:Employee)  
RETURN  
  *
```

# Who do people report to?



# Who do people report to?

```
MATCH  
  (e:Employee)<- [ :REPORTS_TO] - (sub:Employee)  
RETURN  
  e.employeeID AS managerID,  
  e.firstName AS managerName,  
  sub.employeeID AS employeeID,  
  sub.firstName AS employeeName;
```

# Who do people report to?

\$ MATCH (e:Employee)<-[**:REPORTS\_TO**]-(sub:Employee) RETURN e.employeeID AS...

The screenshot shows a query results table from a Neo4j browser. The table has five columns: managerID, managerName, employeeID, and employeeName. The rows represent employees and their managers. The data shows that Andrew manages multiple employees, while Steven, Michael, and Robert manage only one each. Nancy, Janet, Margaret, and Anne are managed by others.

	managerID	managerName	employeeID	employeeName
Rows	2	Andrew	1	Nancy
</>	2	Andrew	3	Janet
Code	2	Andrew	4	Margaret
	2	Andrew	5	Steven
	2	Andrew	8	Laura
	5	Steven	9	Anne
	5	Steven	6	Michael
	5	Steven	7	Robert

Returned 8 rows in 111 ms.

# Who does Robert report to?

```
MATCH
  p=(e:Employee)-[:REPORTS_TO]-(sub:Employee)
WHERE
  sub.firstName = 'Robert'
RETURN
  p
```

# Who does Robert report to?



# What is Robert's reporting chain?

```
MATCH
  p=(e:Employee) <- [ :REPORTS_TO* ] - (sub:Employee)
WHERE
  sub.firstName = 'Robert'
RETURN
  p
```

# What is Robert's reporting chain?

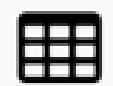


# Who's the Big Boss?

```
MATCH
  (e:Employee)
WHERE
  NOT (e) - [:REPORTS_TO] -> ()
RETURN
  e.firstName as bigBoss
```

# Who's the Big Boss?

```
$ MATCH (e:Employee) WHERE NOT (e)-[:REPORTS_TO]->() RETURN e.firstName a...
```



**bigBoss**

Rows

Andrew



Code

Returned 1 row in 84 ms.

# Product Cross-Selling

```
MATCH
(choc:Product {productName: 'Chocolade' })
<- [ :INCLUDES] - (:Order) <- [ :SOLD] - (employee),
(employee) - [ :SOLD] -> (o2) - [ :INCLUDES] -> (other:Product)
RETURN
employee.firstName,
other.productName,
COUNT(DISTINCT o2) as count
ORDER BY
count DESC
LIMIT 5;
```

# Product Cross-Selling

```
$ MATCH (choc:Product {productName: 'Chocolade'}) <-[ :INCLUDES ] - (:Order) <...  
Rows
```

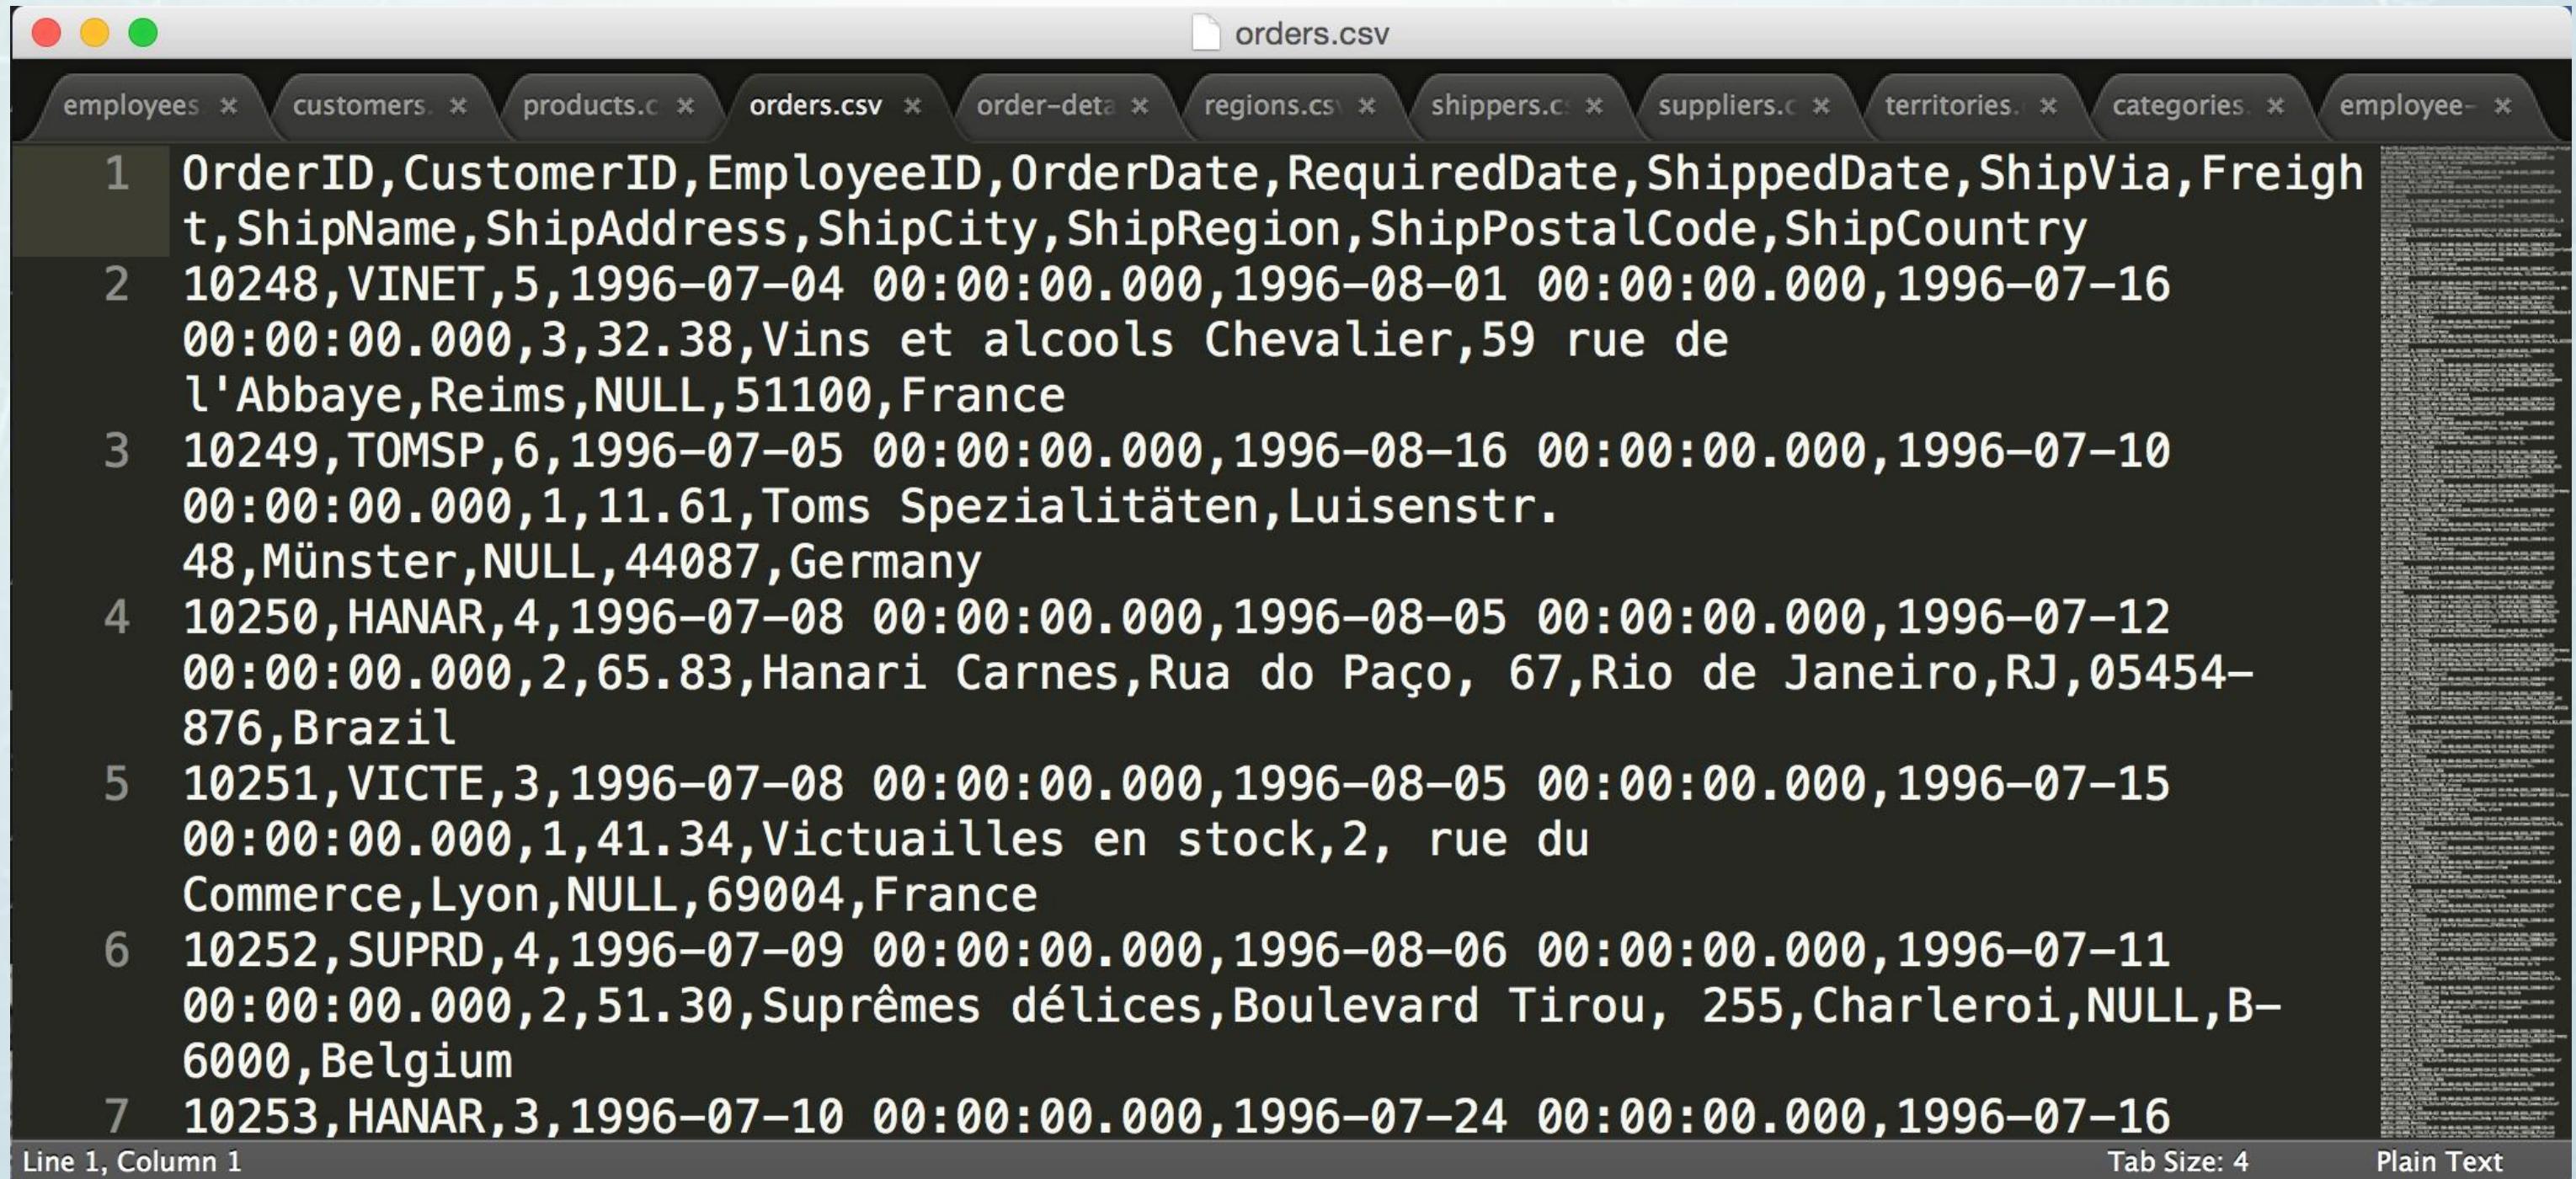
	employee.firstName	other.productName	count
Rows	Margaret	Gnocchi di nonna Alice	14
</>	Janet	Gumbär Gummibärchen	12
Code	Nancy	Flotemysost	12
	Margaret	Pâté chinois	12
	Nancy	Camembert Pierrot	11

Returned 5 rows in 319 ms.

# **LOADING OUR DATA**

# CSV

# CSV files for Northwind

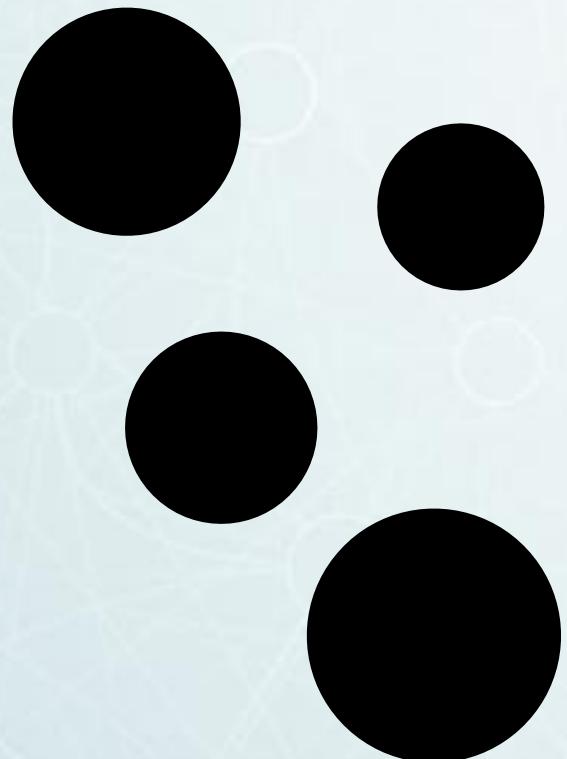


A screenshot of a Mac OS X terminal window titled "orders.csv". The window shows a list of seven orders from the Northwind database. The first column contains numbers 1 through 7, followed by the order details. The columns correspond to the fields in the "orders" table: OrderID, CustomerID, EmployeeID, OrderDate, RequiredDate, ShippedDate, ShipVia, Freight, ShipName, ShipAddress, ShipCity, ShipRegion, ShipPostalCode, and ShipCountry.

	OrderID	CustomerID	EmployeeID	OrderDate	RequiredDate	ShippedDate	ShipVia	Freight	ShipName	ShipAddress	ShipCity	ShipRegion	ShipPostalCode	ShipCountry
1	10248	VINET	5	1996-07-04 00:00:00.000	1996-08-01 00:00:00.000	1996-07-16 00:00:00.000	3	32.38	Vins et alcools Chevalier	59 rue de l'Abbaye	Reims	NULL	51100	France
2	10249	TOMSP	6	1996-07-05 00:00:00.000	1996-08-16 00:00:00.000	1996-07-10 00:00:00.000	1	11.61	Toms Spezialitäten	Luisenstr. 48	Münster	NULL	44087	Germany
3	10250	HANAR	4	1996-07-08 00:00:00.000	1996-08-05 00:00:00.000	1996-07-12 00:00:00.000	2	65.83	Hanari Carnes	Rua do Paço, 67	Rio de Janeiro	RJ	05454-876	Brazil
4	10251	VICTE	3	1996-07-08 00:00:00.000	1996-08-05 00:00:00.000	1996-07-15 00:00:00.000	1	41.34	Victuailles en stock	2, rue du Commerce	Lyon	NULL	69004	France
5	10252	SUPRD	4	1996-07-09 00:00:00.000	1996-08-06 00:00:00.000	1996-07-11 00:00:00.000	2	51.30	Suprêmes délices	Boulevard Tirou, 255	Charleroi	NULL	B-6000	Belgium
6	10253	HANAR	3	1996-07-10 00:00:00.000	1996-07-24 00:00:00.000	1996-07-16 00:00:00.000								

Line 1, Column 1      Tab Size: 4      Plain Text

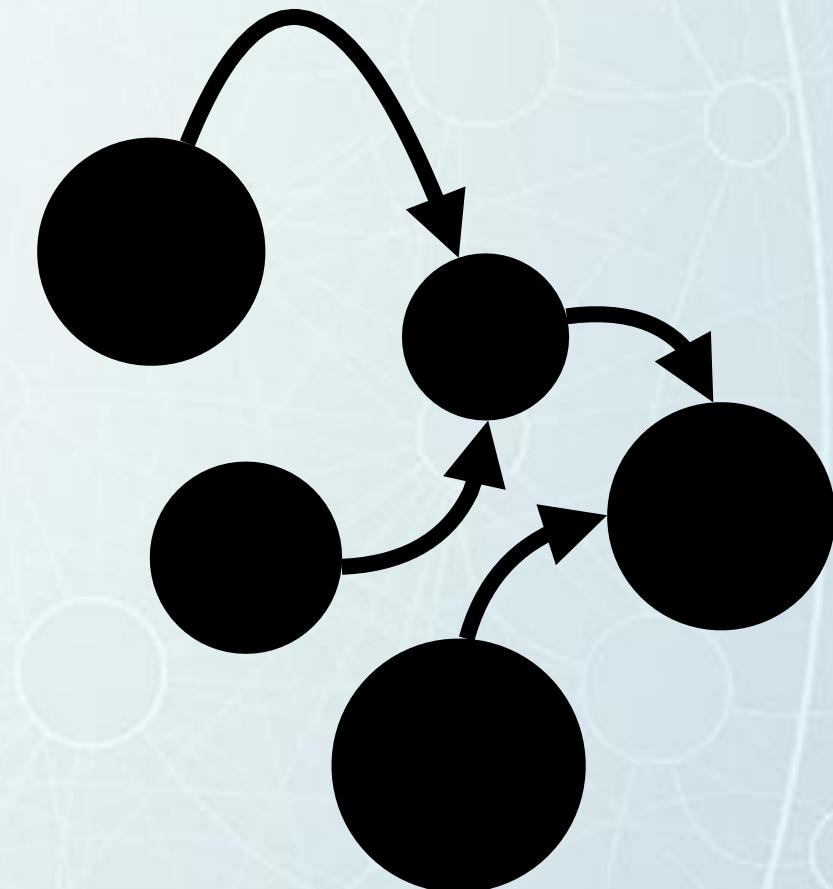
# 3 Steps to Creating the Graph



IMPORT NODES



CREATE INDEXES



IMPORT RELATIONSHIPS

# Importing Nodes

```
// Create customers
USING PERIODIC COMMIT
LOAD CSV WITH HEADERS FROM "https://
raw.githubusercontent.com/neo4j-contrib/developer-resources/
gh-pages/data/northwind/customers.csv" AS row
CREATE (:Customer {companyName: row.CompanyName, customerID:
row.CustomerID, fax: row.Fax, phone: row.Phone});
```

```
// Create products
USING PERIODIC COMMIT
LOAD CSV WITH HEADERS FROM "https://
raw.githubusercontent.com/neo4j-contrib/developer-resources/
gh-pages/data/northwind/products.csv" AS row
CREATE (:Product {productName: row.ProductName, productID:
row.ProductID, unitPrice: toFloat(row.UnitPrice)});
```

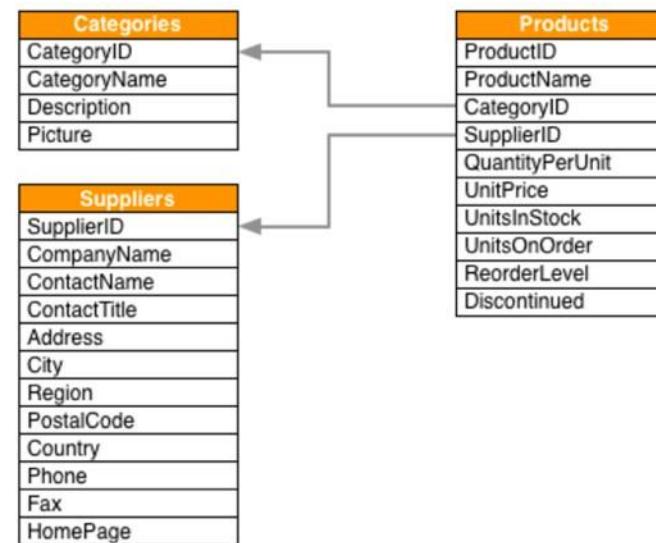
# :play northward graph

```
:play northwind graph
```

## Product Catalog

Northwind sells food products in a few categories, provided by suppliers. Let's start by loading the product catalog tables.

The load statements to the right require public internet access. **LOAD CSV** will retrieve a CSV file from a valid URL, applying a Cypher statement to each row using a named map (here we're using the name `row`).



:help [cypher](#) [LOAD CSV](#)

## Load records

```
LOAD CSV WITH HEADERS FROM "http://data.neo4j.com/northwind/products.csv" AS row
CREATE (n:Product)
SET n = row,
    n.unitPrice =toFloat(row.unitPrice),
    n.unitsInStock =toInt(row.unitsInStock), n.unitsOnOrder =toInt(row.unitsOnOrder),
    n.reorderLevel =toInt(row.reorderLevel), n.discontinued = (row.discontinued <> "0")
```

```
LOAD CSV WITH HEADERS FROM "http://data.neo4j.com/northwind/categories.csv" AS row
CREATE (n:Category)
SET n = row
```

```
LOAD CSV WITH HEADERS FROM "http://data.neo4j.com/northwind/suppliers.csv" AS row
CREATE (n:Supplier)
SET n = row
```

## Create indexes

```
CREATE INDEX ON :Product(productID)
```

```
CREATE INDEX ON :Category(categoryID)
```

```
CREATE INDEX ON :Supplier(supplierID)
```

# High Performance LOADING

**neo4j-import**

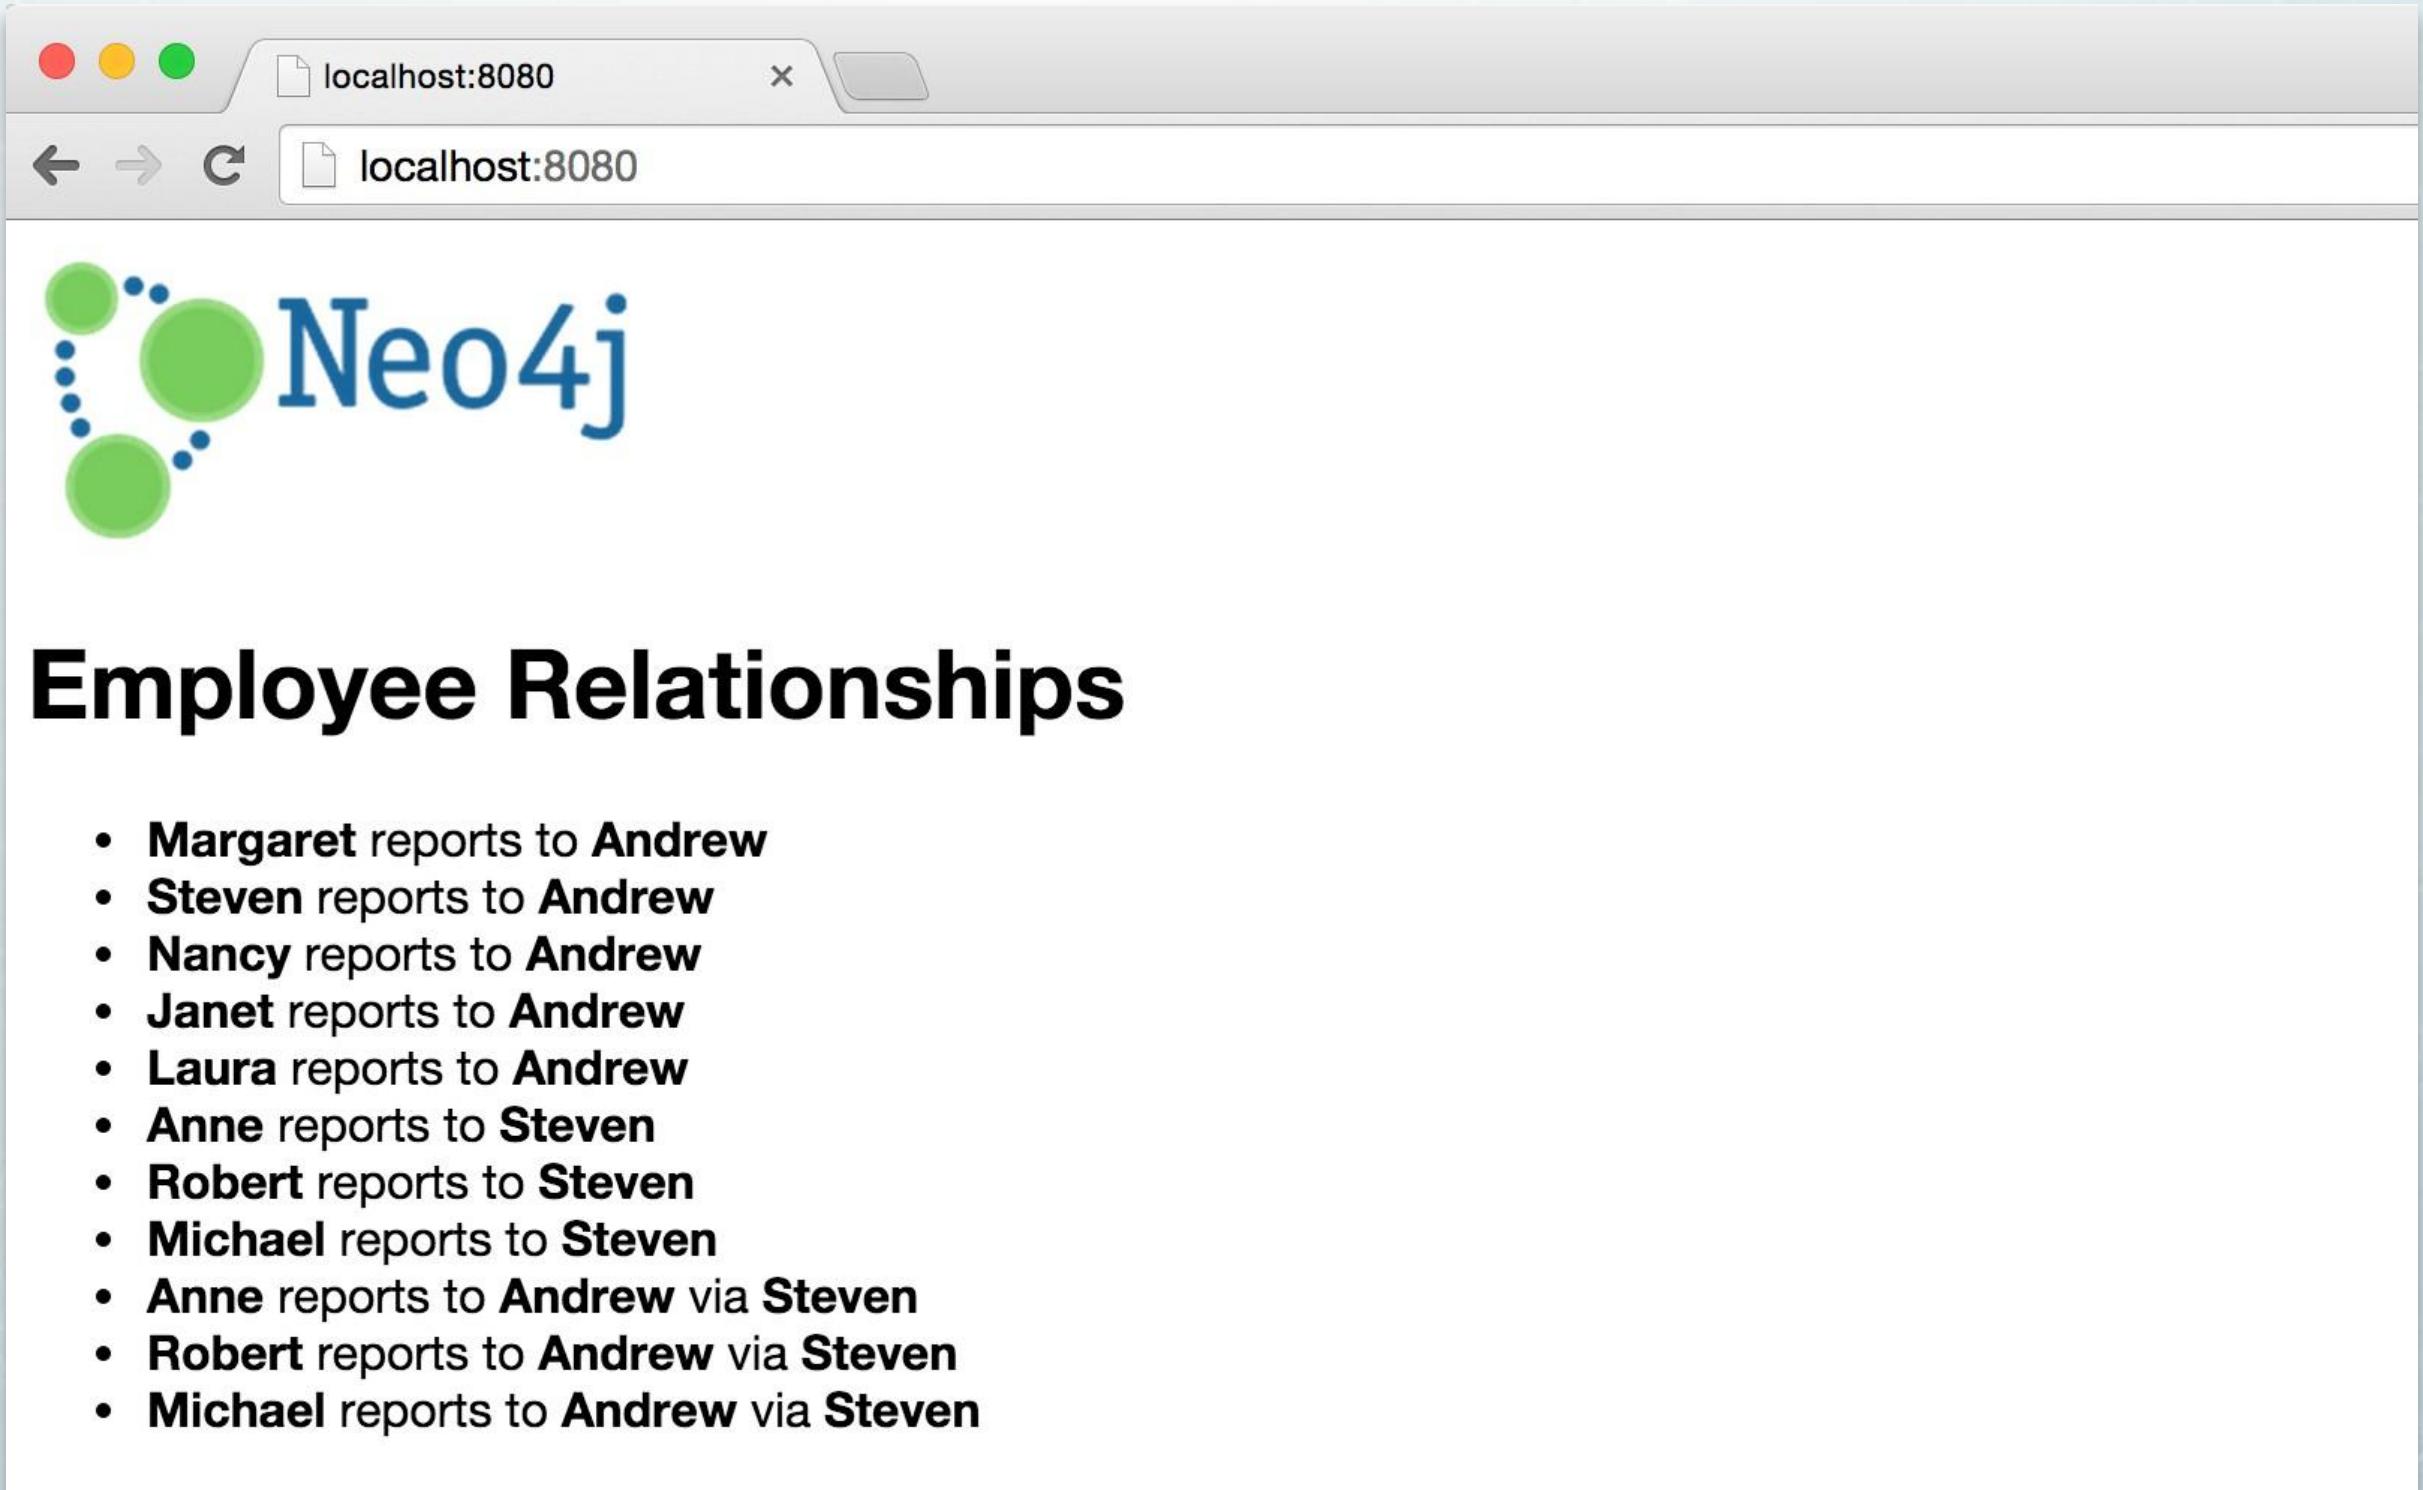


4.58 million things  
and their relationships...

**Loads in 100 seconds!**

# **POWERING AN APP**

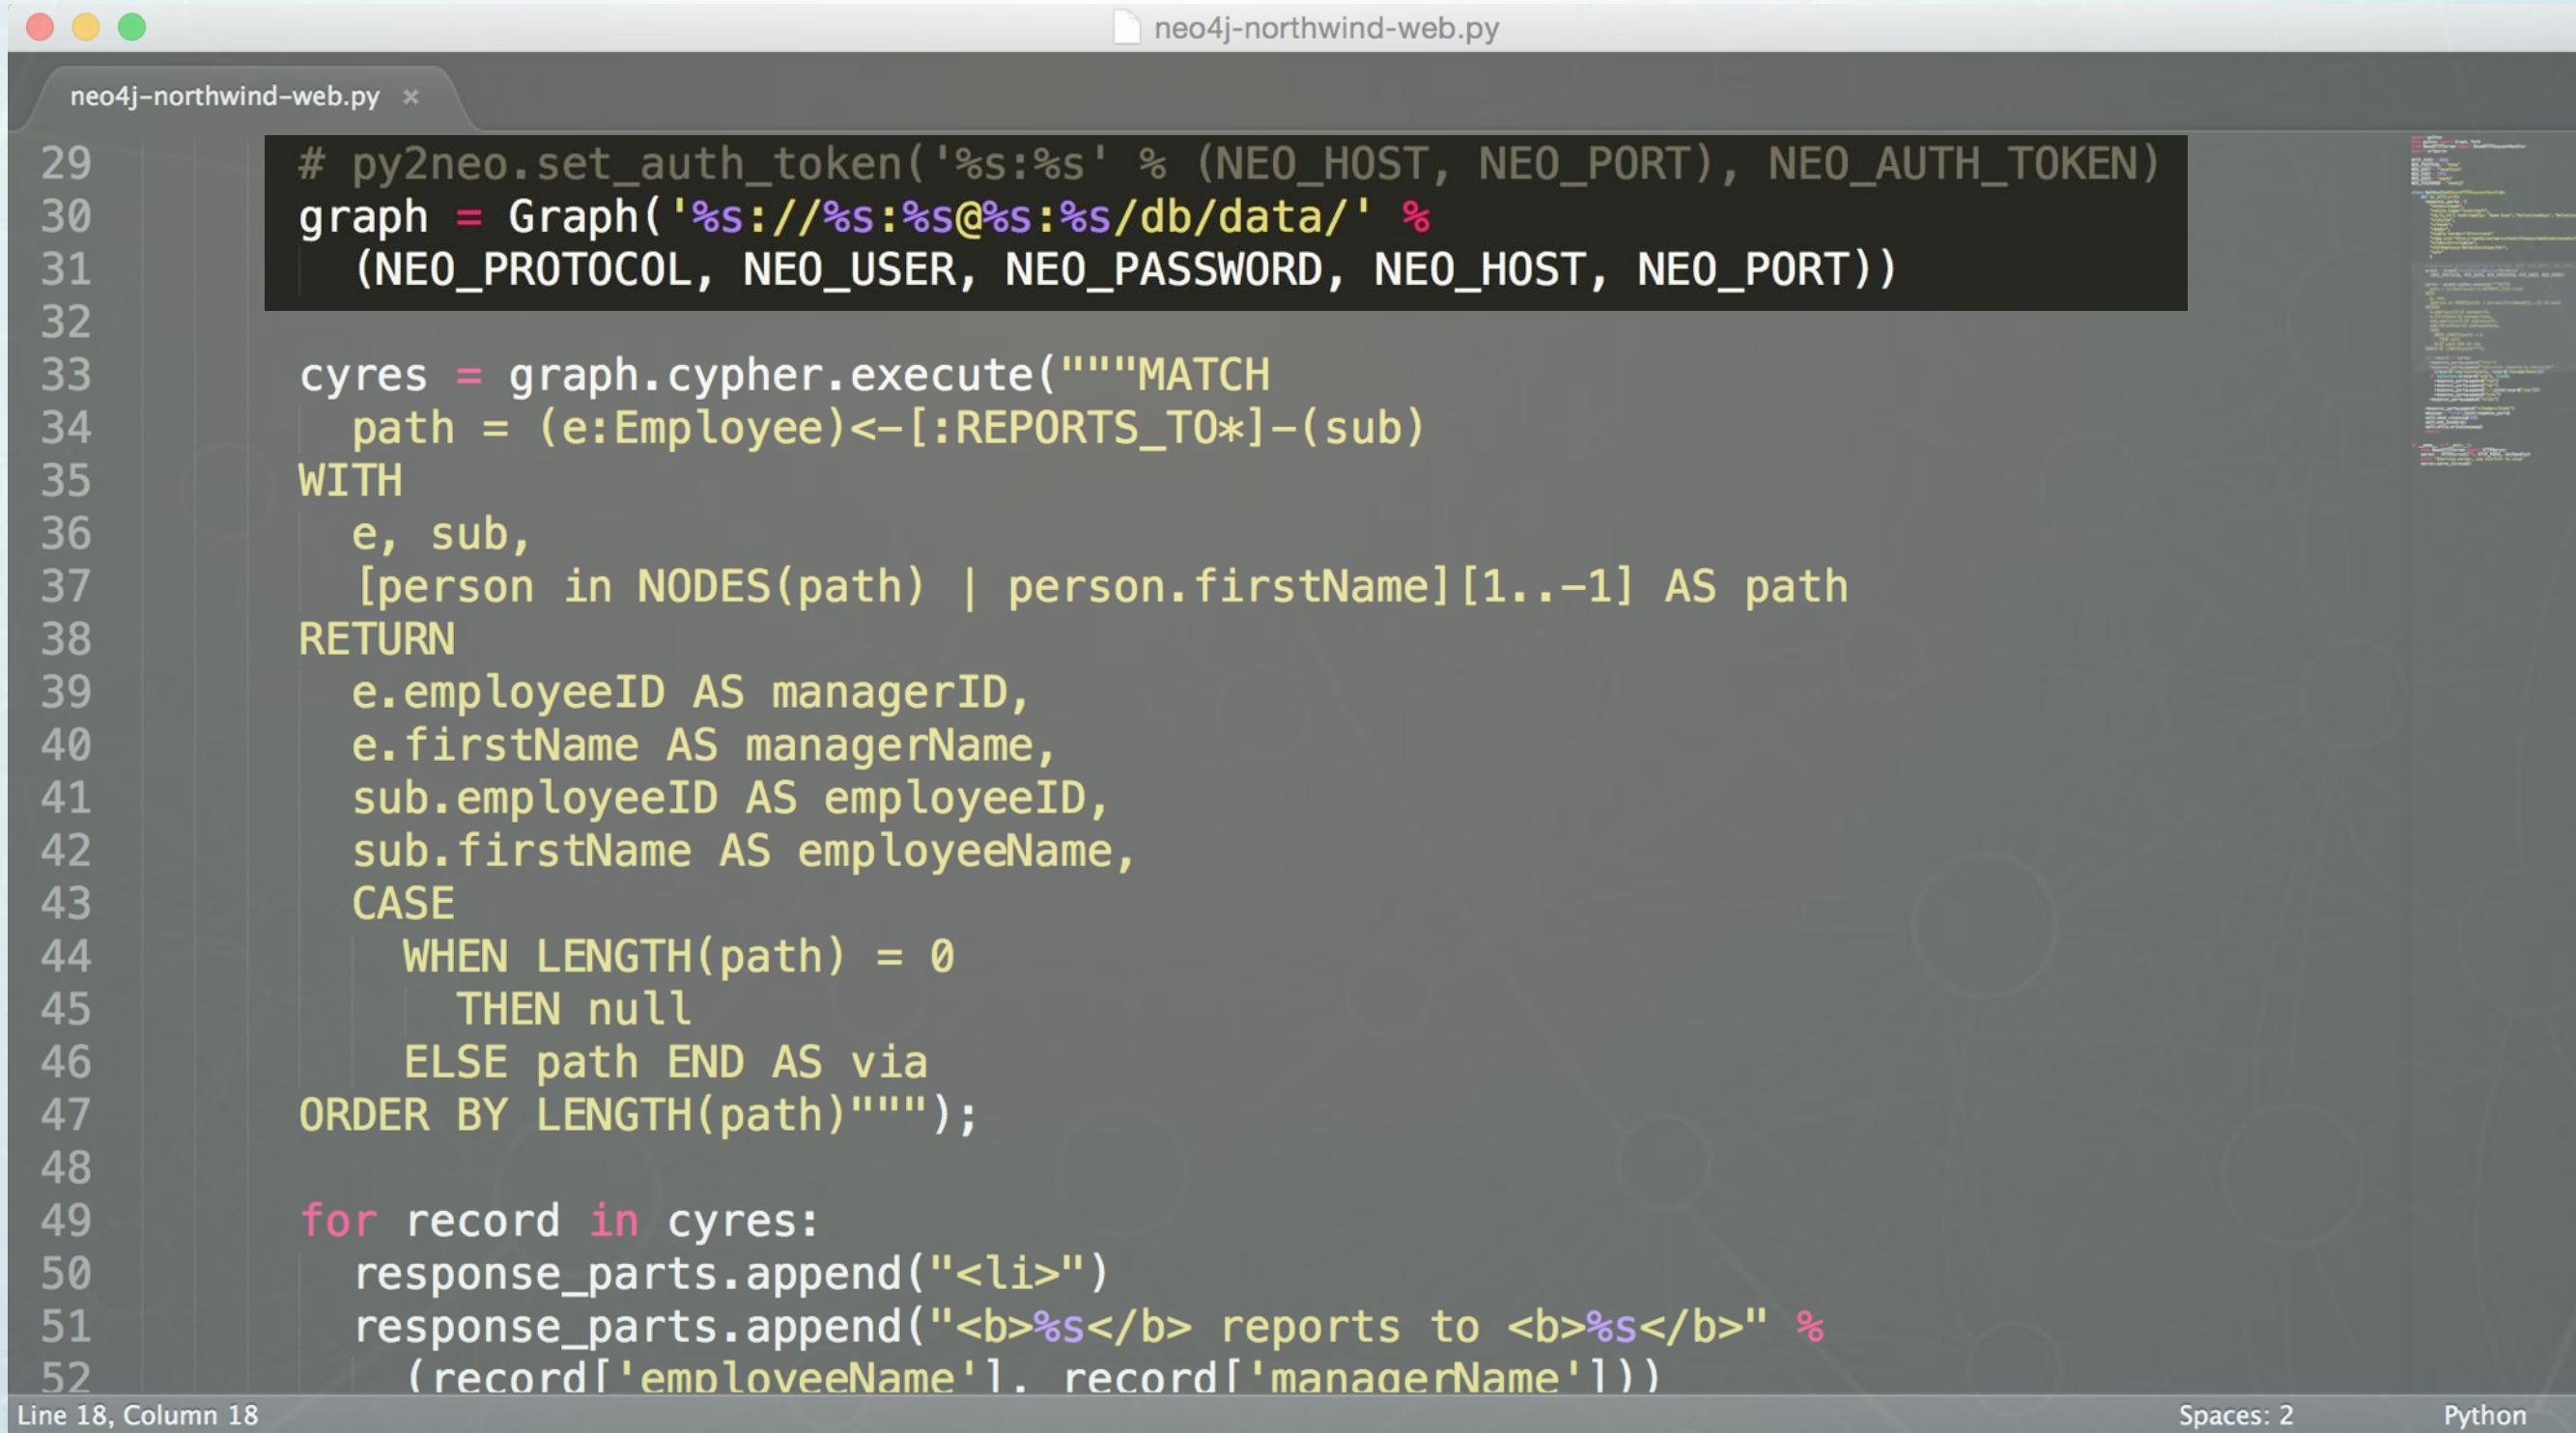
# Simple App



The screenshot shows a web browser window with the URL `localhost:8080`. The page features the Neo4j logo (three green circles) and the title "Employee Relationships". Below the title is a bulleted list of employee relationships:

- Margaret reports to Andrew
- Steven reports to Andrew
- Nancy reports to Andrew
- Janet reports to Andrew
- Laura reports to Andrew
- Anne reports to Steven
- Robert reports to Steven
- Michael reports to Steven
- Anne reports to Andrew via Steven
- Robert reports to Andrew via Steven
- Michael reports to Andrew via Steven

# Simple Python Code

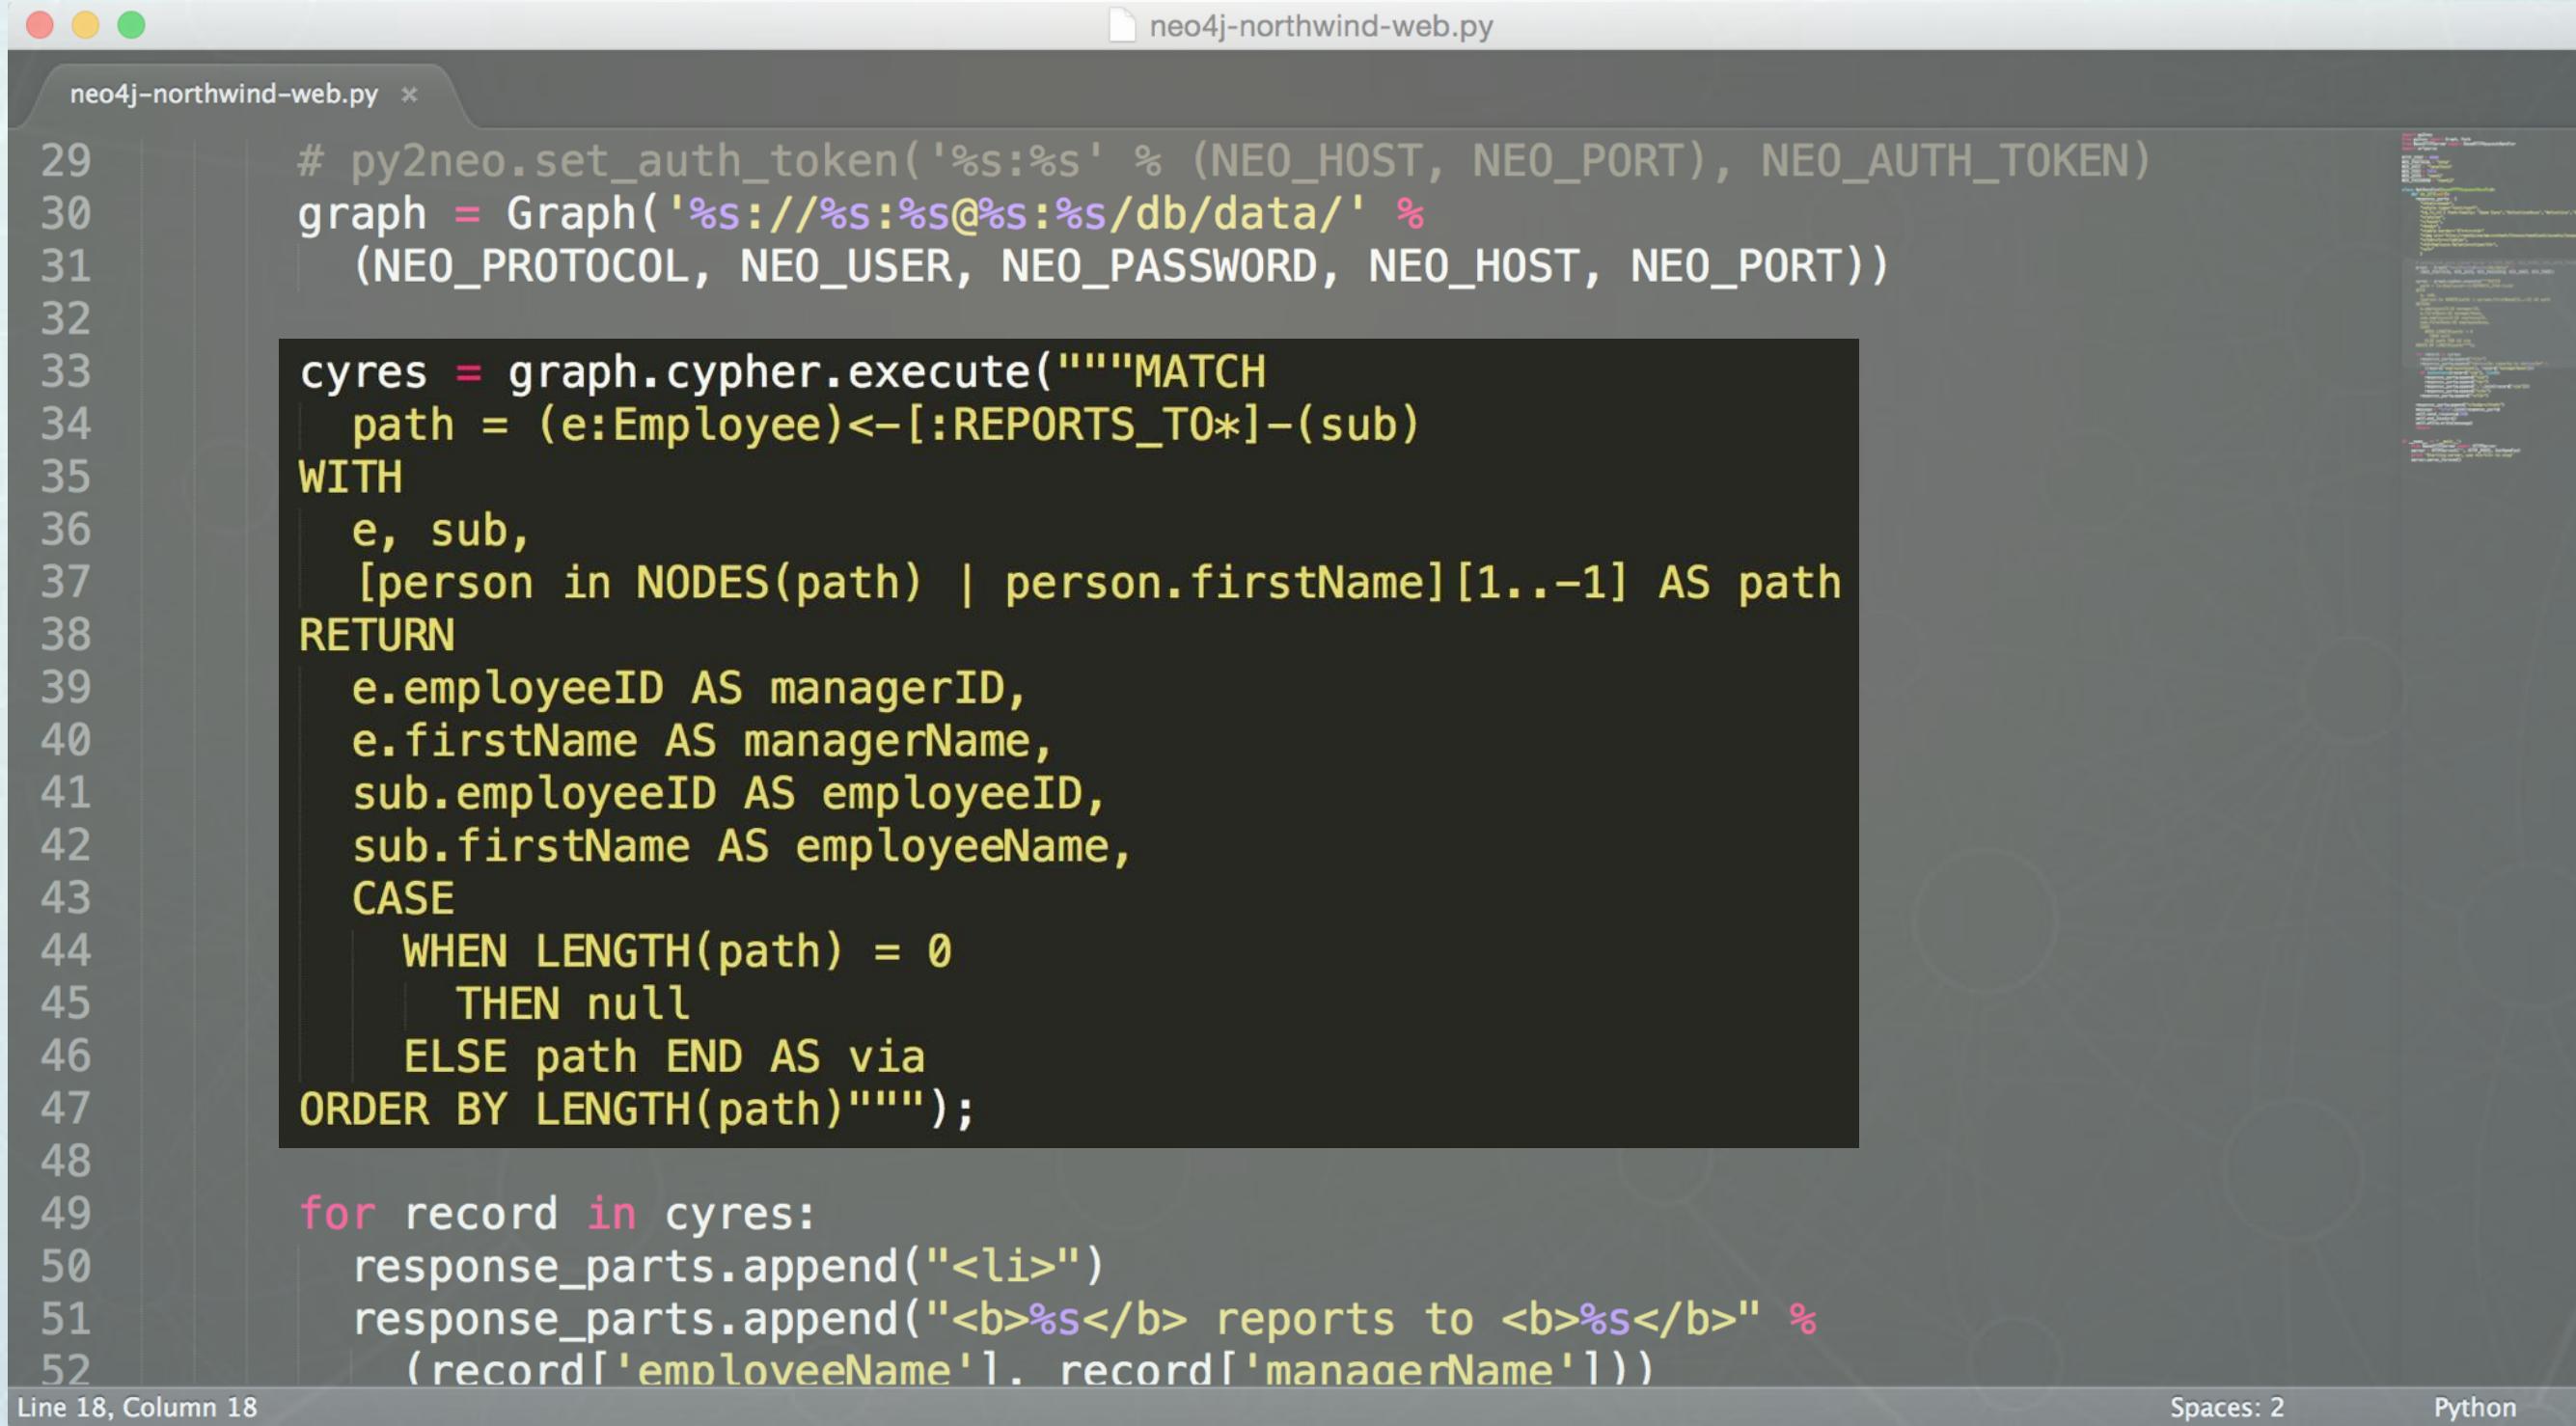


A screenshot of a code editor window titled "neo4j-northwind-web.py". The code uses the py2neo library to connect to a Neo4j database and execute Cypher queries. The code is annotated with line numbers from 29 to 52.

```
29 # py2neo.set_auth_token('%s:%s' % (NEO_HOST, NEO_PORT), NEO_AUTH_TOKEN)
30 graph = Graph('"%s://:%s@%s:%s/db/data/" % '
31               (NEO_PROTOCOL, NEO_USER, NEO_PASSWORD, NEO_HOST, NEO_PORT))
32
33 cyres = graph.cypher.execute("""MATCH
34   path = (e:Employee)-[:REPORTS_TO*]-(sub)
35 WITH
36   e, sub,
37   [person in NODES(path) | person.firstName[1..-1]] AS path
38 RETURN
39   e.employeeID AS managerID,
40   e.firstName AS managerName,
41   sub.employeeID AS employeeID,
42   sub.firstName AS employeeName,
43 CASE
44   WHEN LENGTH(path) = 0
45   THEN null
46   ELSE path END AS via
47 ORDER BY LENGTH(path)""");
48
49 for record in cyres:
50   response_parts.append("<li>")
51   response_parts.append("<b>%s</b> reports to <b>%s</b>" %
52                         (record['employeeName'], record['managerName']))
```

Line 18, Column 18      Spaces: 2      Python

# Simple Python Code



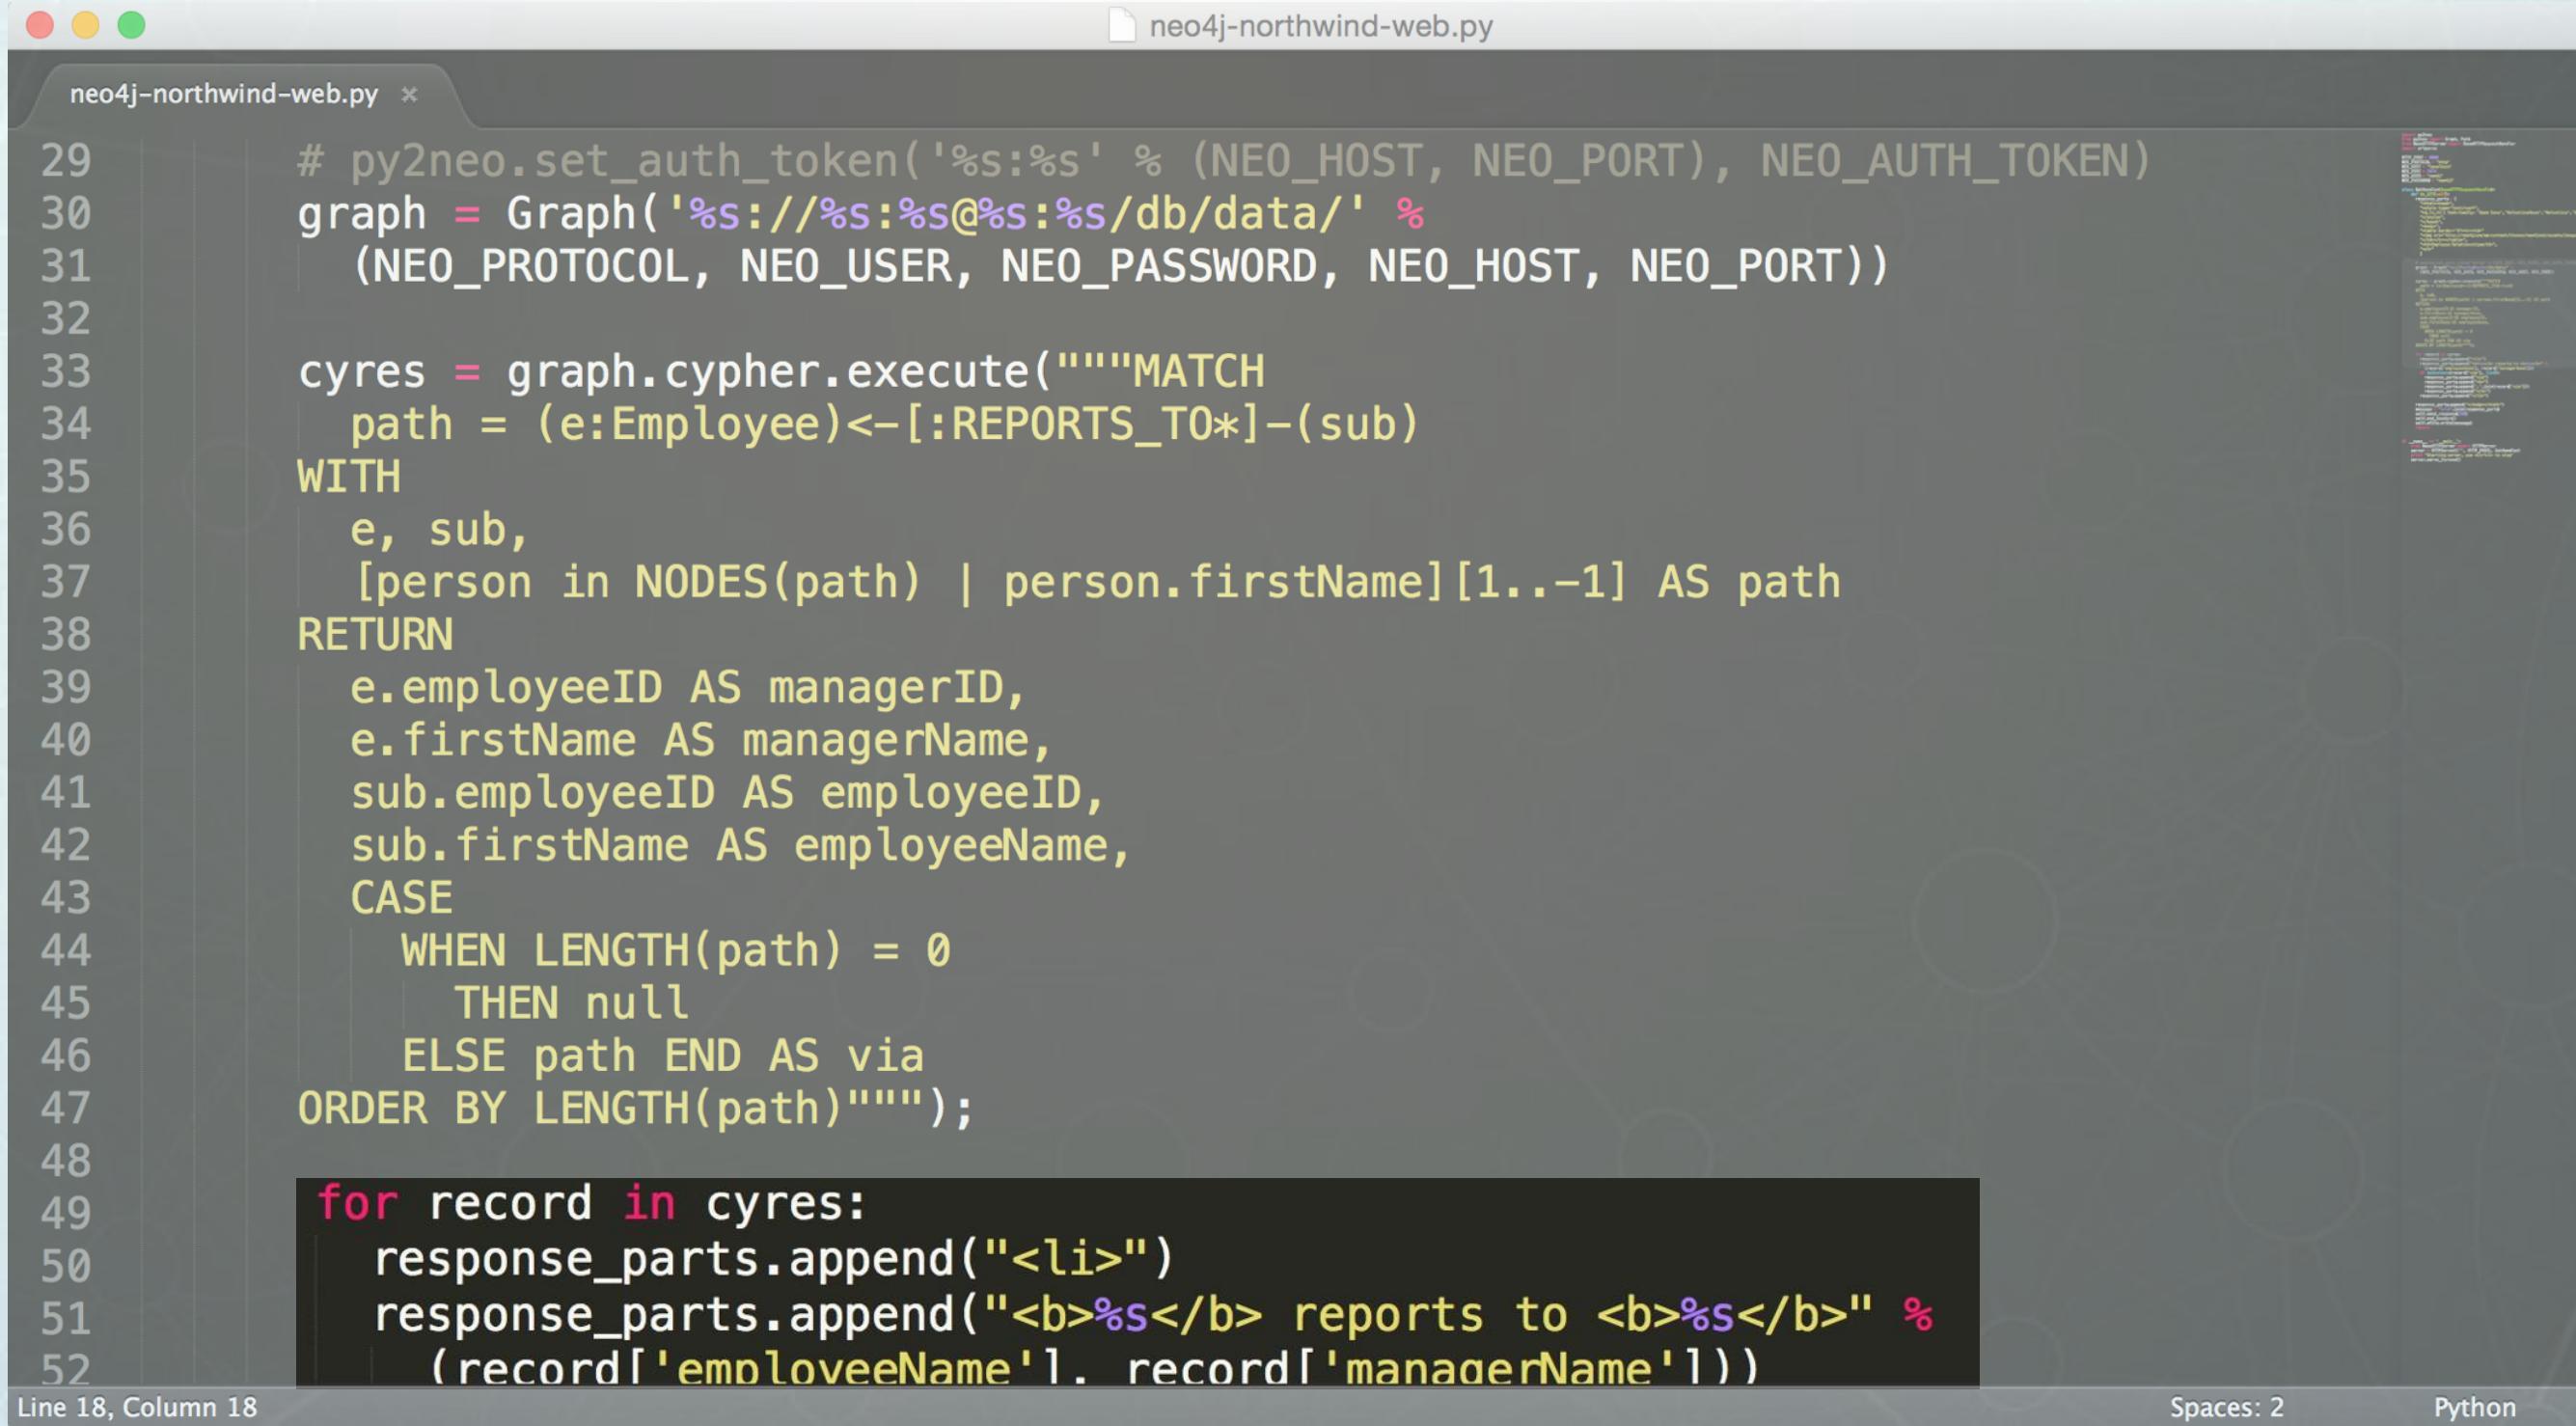
A screenshot of a code editor window titled "neo4j-northwind-web.py". The code uses the py2neo library to interact with a Neo4j database. It defines a Cypher query to find management paths and iterates over the results to build a response.

```
neo4j-northwind-web.py * neo4j-northwind-web.py

29 # py2neo.set_auth_token('%s:%s' % (NEO_HOST, NEO_PORT), NEO_AUTH_TOKEN)
30 graph = Graph('%s://'%s:%s@%s:%s/db/data/' %
31     (NEO_PROTOCOL, NEO_USER, NEO_PASSWORD, NEO_HOST, NEO_PORT))
32
33 cyres = graph.cypher.execute("""MATCH
34     path = (e:Employee)-[:REPORTS_TO*]-(sub)
35 WITH
36     e, sub,
37     [person in NODES(path) | person.firstName][1..-1] AS path
38 RETURN
39     e.employeeID AS managerID,
40     e.firstName AS managerName,
41     sub.employeeID AS employeeID,
42     sub.firstName AS employeeName,
43     CASE
44         WHEN LENGTH(path) = 0
45             THEN null
46         ELSE path END AS via
47 ORDER BY LENGTH(path)""");
48
49 for record in cyres:
50     response_parts.append("<li>")
51     response_parts.append("<b>%s</b> reports to <b>%s</b>" %
52         (record['employeeName'], record['managerName']))
```

Line 18, Column 18      Spaces: 2      Python

# Simple Python Code



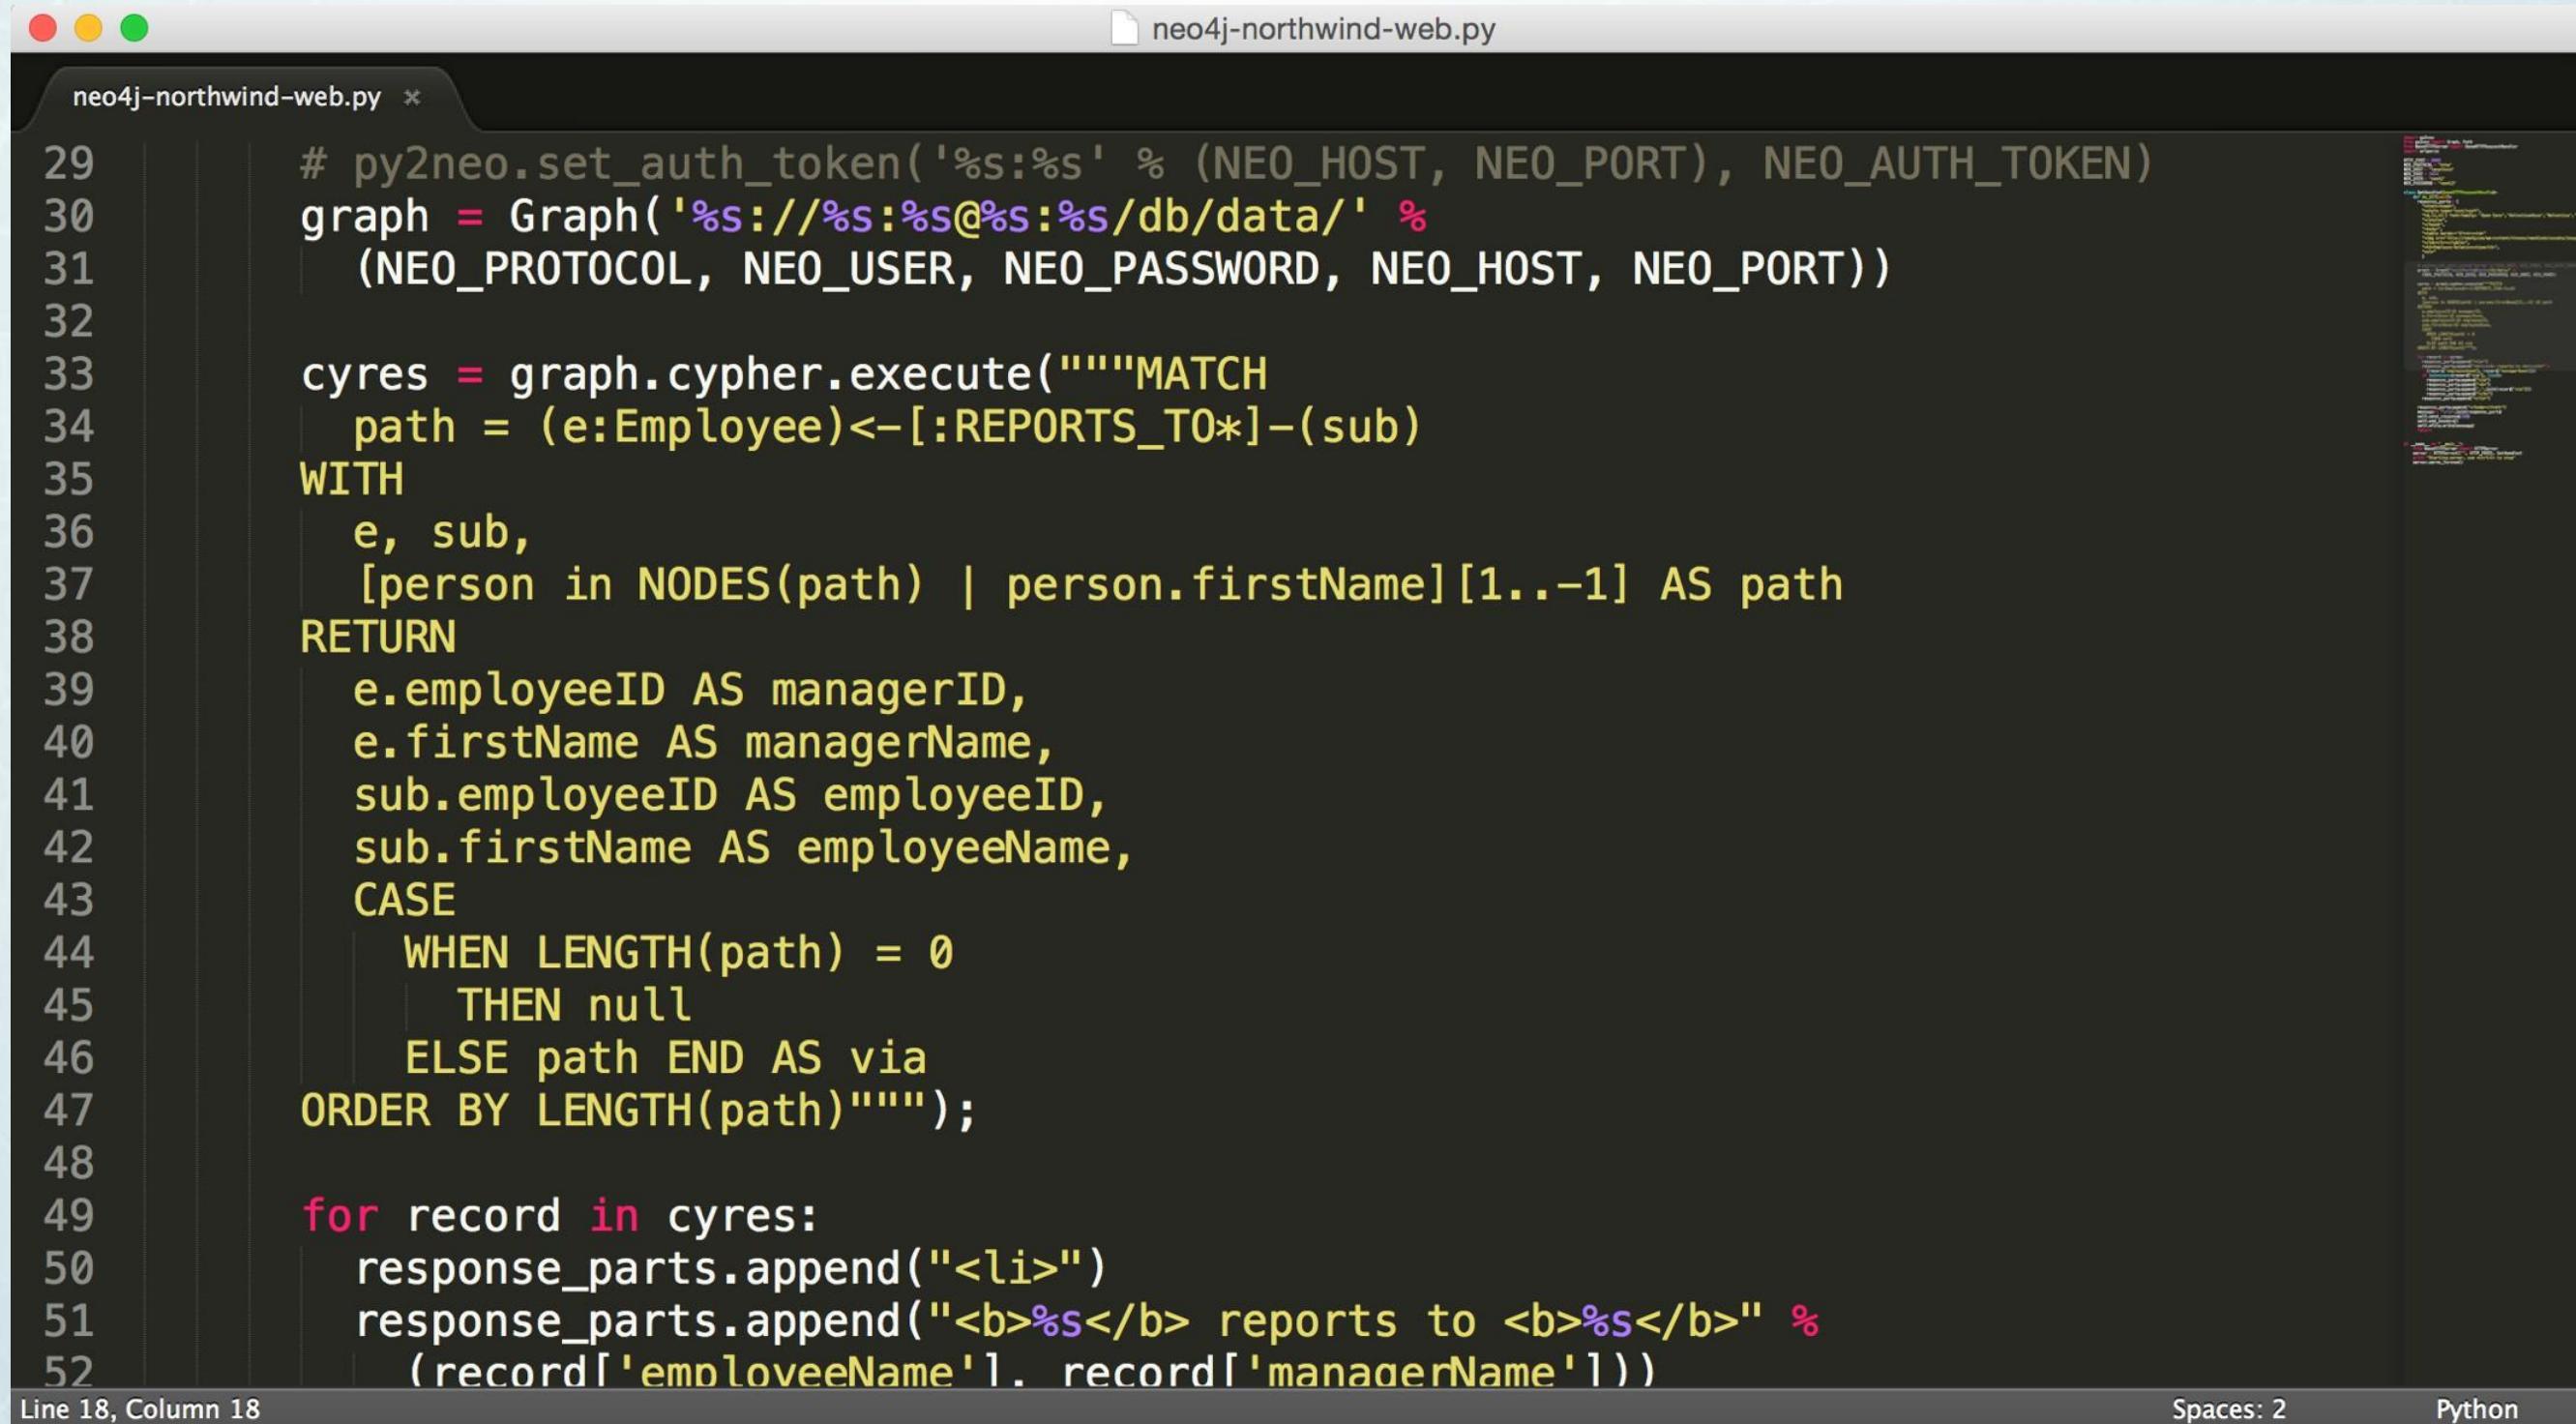
The screenshot shows a code editor window titled "neo4j-northwind-web.py". The code is a combination of Python and Cypher. It sets up a connection to a Neo4j database using py2neo and executes a Cypher query to find management paths. The Python code then processes the results and appends them to a list of HTML fragments.

```
neo4j-northwind-web.py * neo4j-northwind-web.py

29 # py2neo.set_auth_token('%s:%s' % (NEO_HOST, NEO_PORT), NEO_AUTH_TOKEN)
30 graph = Graph('%s://'%s:%s@%s:%s/db/data/' %
31     (NEO_PROTOCOL, NEO_USER, NEO_PASSWORD, NEO_HOST, NEO_PORT))
32
33 cyres = graph.cypher.execute("""MATCH
34     path = (e:Employee)<-[ :REPORTS_TO* ]-(sub)
35 WITH
36     e, sub,
37     [person in NODES(path) | person.firstName][1..-1] AS path
38 RETURN
39     e.employeeID AS managerID,
40     e.firstName AS managerName,
41     sub.employeeID AS employeeID,
42     sub.firstName AS employeeName,
43     CASE
44         WHEN LENGTH(path) = 0
45             THEN null
46         ELSE path END AS via
47 ORDER BY LENGTH(path)""");
48
49 for record in cyres:
50     response_parts.append("<li>")
51     response_parts.append("<b>%s</b> reports to <b>%s</b>" %
52         (record['employeeName'], record['managerName']))
```

Line 18, Column 18      Spaces: 2      Python

# Simple Python Code



The screenshot shows a code editor window titled "neo4j-northwind-web.py". The code is a Python script that performs a Cypher query to find employee management paths and appends the results to a list of HTML list items.

```
neo4j-northwind-web.py *
29     # py2neo.set_auth_token('%s:%s' % (NEO_HOST, NEO_PORT), NEO_AUTH_TOKEN)
30     graph = Graph('"%s://%s:%s@%s:%s/db/data/" %'
31                   (NEO_PROTOCOL, NEO_USER, NEO_PASSWORD, NEO_HOST, NEO_PORT))
32
33     cyres = graph.cypher.execute("""MATCH
34         path = (e:Employee)-[:REPORTS_TO*]-(sub)
35     WITH
36         e, sub,
37         [person in NODES(path) | person.firstName][1..-1] AS path
38     RETURN
39         e.employeeID AS managerID,
40         e.firstName AS managerName,
41         sub.employeeID AS employeeID,
42         sub.firstName AS employeeName,
43         CASE
44             WHEN LENGTH(path) = 0
45                 THEN null
46             ELSE path END AS via
47     ORDER BY LENGTH(path)""");
48
49     for record in cyres:
50         response_parts.append("<li>")
51         response_parts.append("<b>%s</b> reports to <b>%s</b>" %
52                             (record['employeeName'], record['managerName']))
```

Line 18, Column 18      Spaces: 2      Python

# Resources



# Install Neo4j and Take it for a Spin

Experience powerful scalability, blazing speed and unparalleled flexibility – download and try Neo4j today.

## For Business

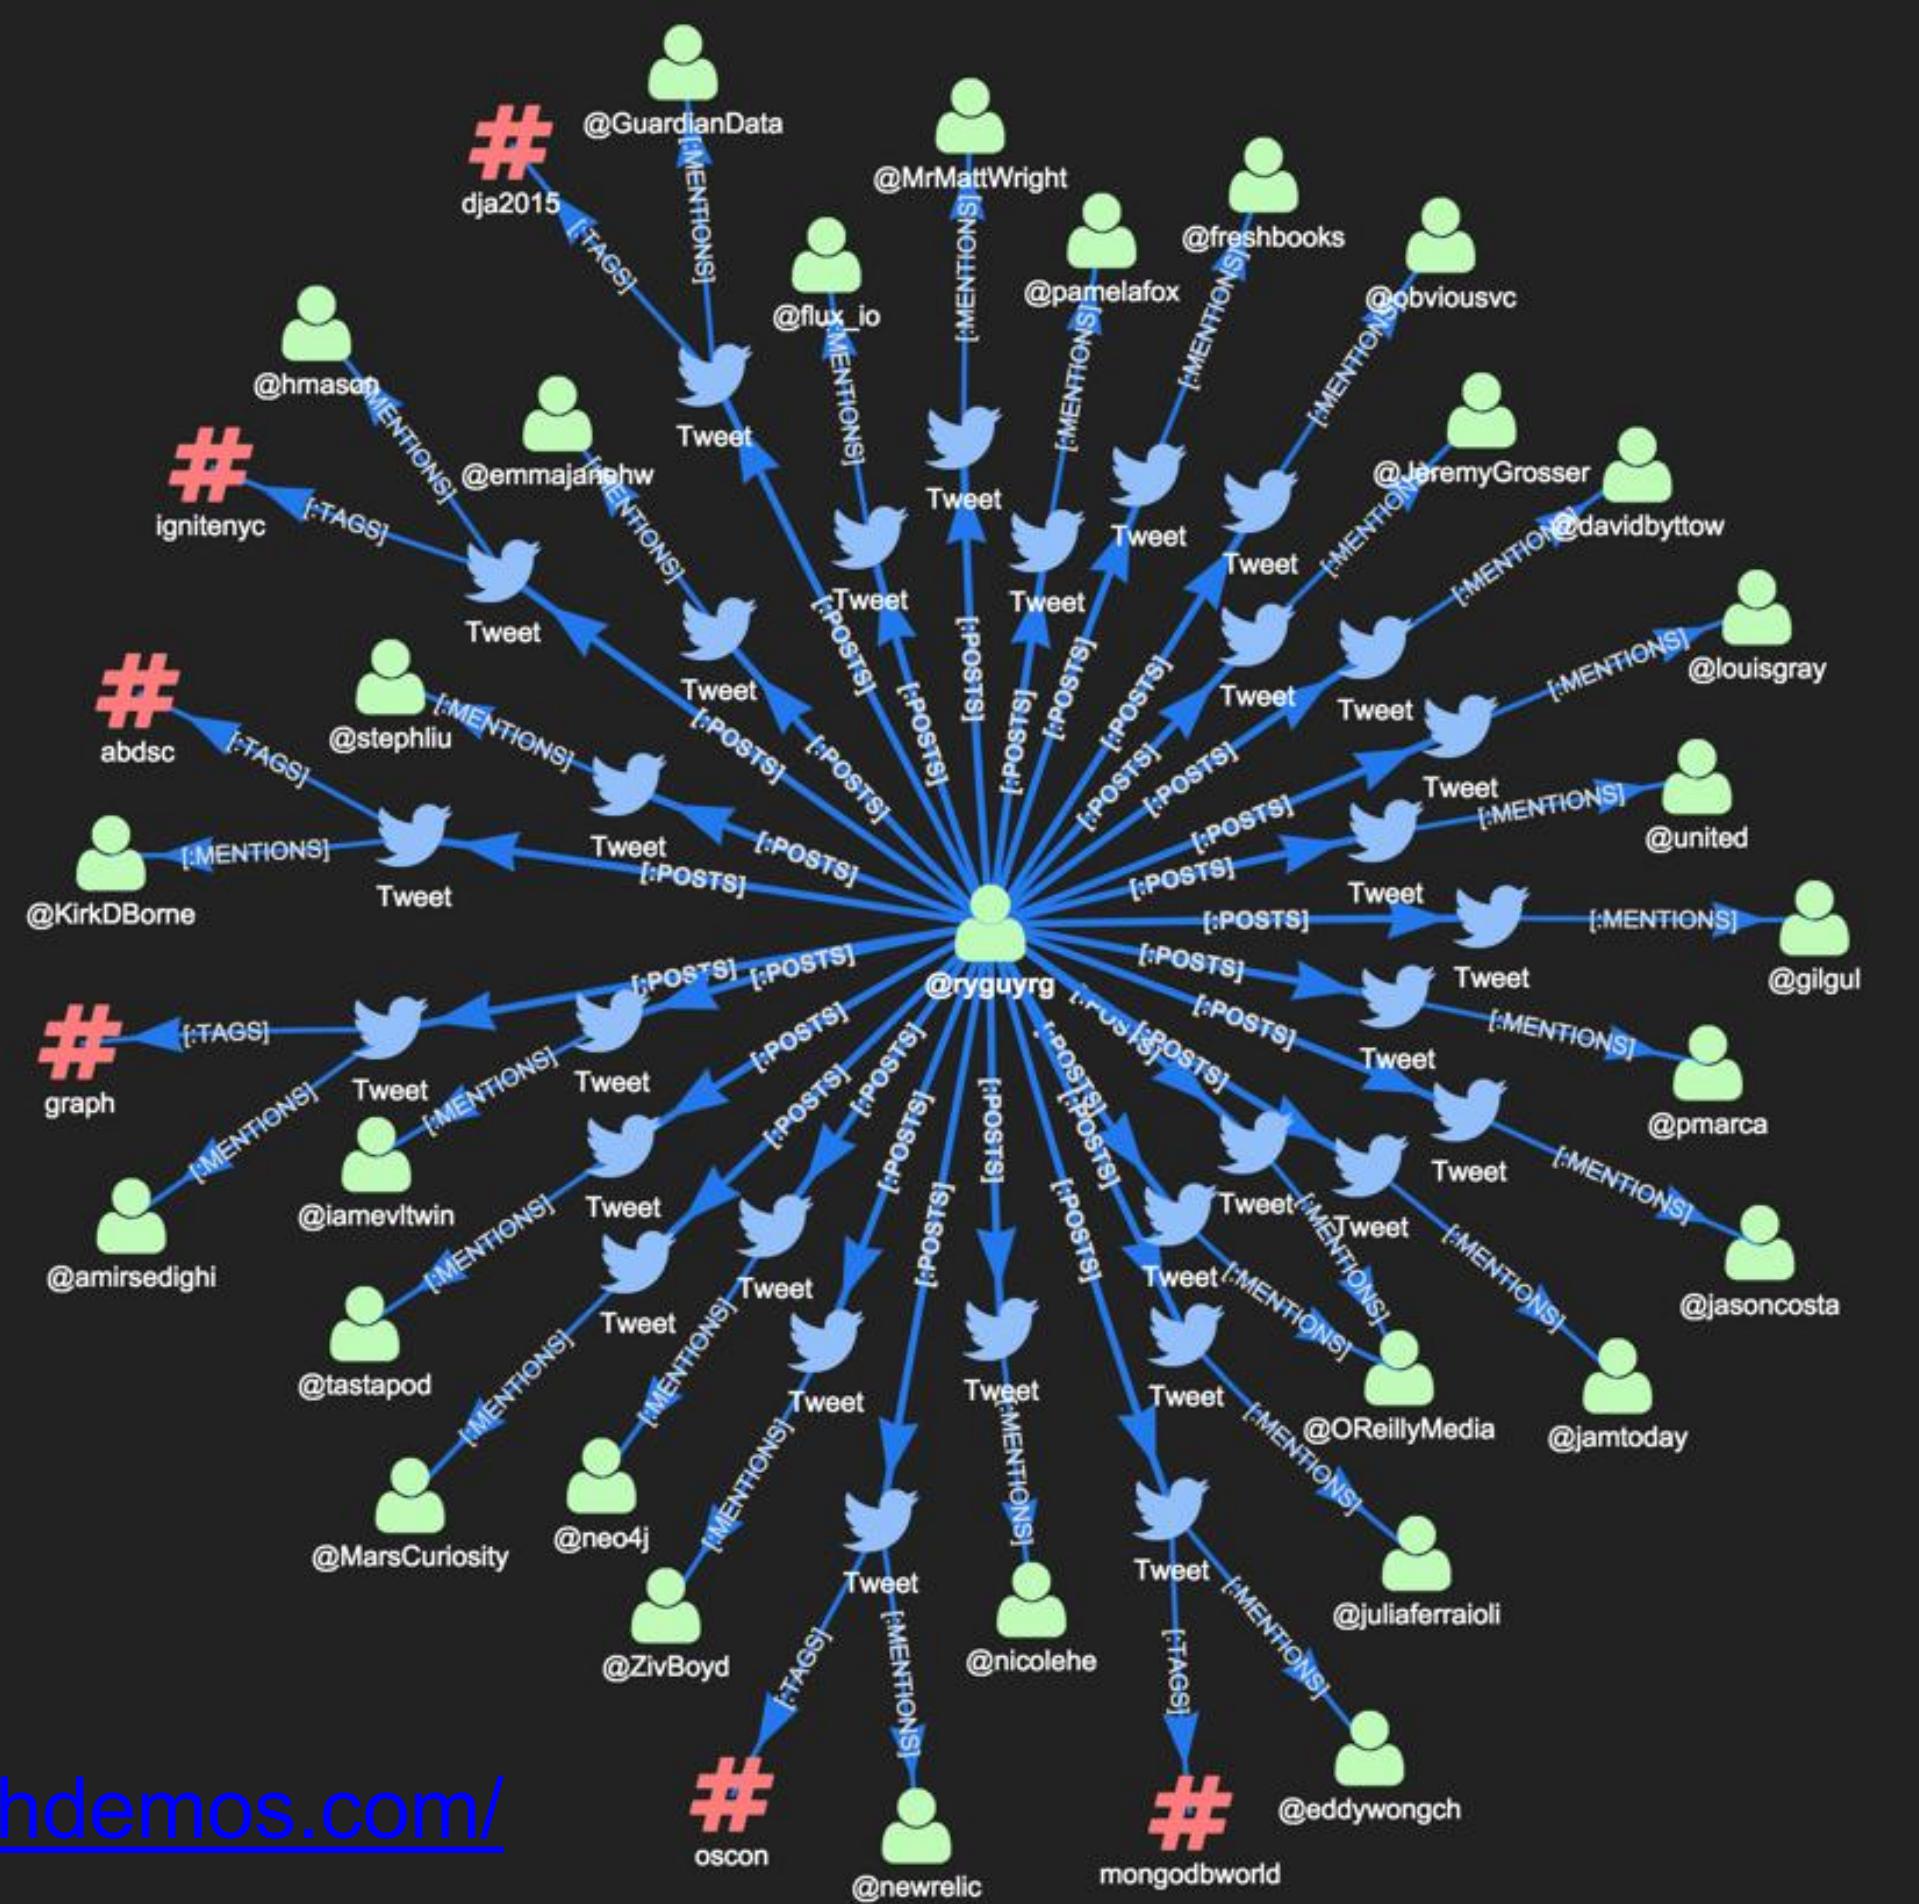
The Neo4j Enterprise Edition offers incredible power and flexibility, with enterprise-grade availability, management and scale-up & scale-out capabilities.

[Download Free Enterprise Trial](#)

## For Individuals

Ideal for learning, and smaller do-it-yourself projects that do not require high levels of scaling. Excludes professional services and support.

[Download Community Edition](#)



<http://network.graphdemos.com/>

# [neo4j.com/developer](http://neo4j.com/developer)



## Test-Drive Neo4j with Cypher

Social

Network Management

Fraud Detection

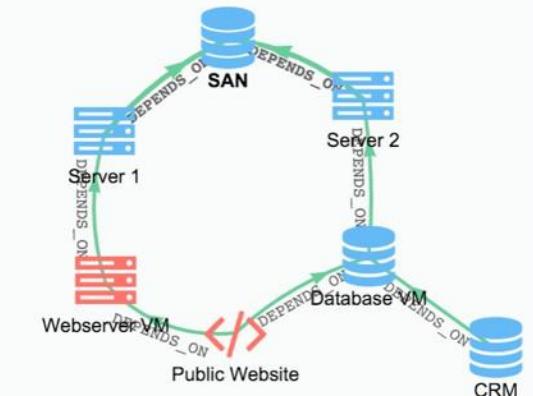
[Impact Analysis](#) | [Dependency Analysis](#) | [Statistics](#)

### Impact Analysis

Find all services that depend on Server 1. These would be impacted by an outage of that server.

```
MATCH
  (n:Service)<-[ :DEPENDS_ON* ]-(dependent:Service)
WHERE
  n.name = "Server 1"
RETURN
  dependent
```

Only Webserver VM depends on Server 1. Because we're looking at variable length paths of DEPENDS\_ON relationships, we're also able to determine that Public Website would be impacted by an outage of Server 1.



See Code In: [JAVA](#) [PYTHON](#) [RUBY](#) [PHP](#) [C#](#) [NODE.JS](#)

**Downloading and Installing Java**

- [Download Neo4j JDBC](#)
- Copy and paste code at left into [Social.java](#)
- Run [javac Social.java](#)
- Run [java -cp /path/to/neo4j-jdbc-2.3-SNAPSHOT-jar-with-dependencies.jar:. Social](#)

```
// javac Network.java
// java -cp /path/to/neo4j-jdbc-2.3-SNAPSHOT-jar-with-dependencies.jar:. Network

import java.sql.*;
import static java.util.Arrays.asList;
import java.util.List;

public class Network {

    public static void query( Connection con,
                           String query, String[] columns, Object
...params)
        throws SQLException {
        try (PreparedStatement pst = con.prepareStatement(query)) {
            for (int i=0;i<params.length;i++)
                pst.setObject(i + 1, params[i]);
            ResultSet rs = pst.executeQuery();
            int count = 0;
            while (rs.next()) {
                for (int i=0;i<columns.length;i++)
                    System.out.print(rs.getString(columns[i])+"\t"
);
        }
    }
}
```

# There Are Lots of Ways to Easily Learn Neo4j



community



Documentation



## Graph Academy

Learn. Graph. Deploy.



BOOKS

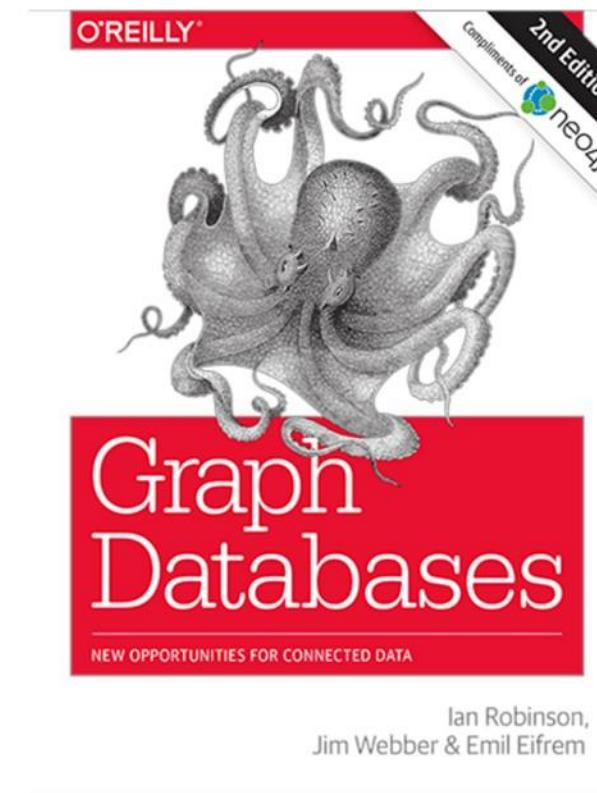
Built-in  
Guides

On-site  
Training

GraphGist

Online Training





**UPDATE! 2nd Edition**

## O'Reilly's *Graph Databases*

The Definitive Book on  
**Graph Databases.**

Official Released Version of O'Reilly's Graph  
Databases.

Copyright (c) 2015, Neo Technology, Inc. All rights reserved. The reproduction  
or distribution of this copyrighted work is strictly prohibited.

---

***Graph Databases***  
The Definitive Book on Graph Databases and Introduction to Neo4j

[graphdatabases.com](http://graphdatabases.com)



## What is a GraphGist?

With Neo4j GraphGists you can describe and model your domain in a simple text file ([AsciiDoc](#)) and render it as a rich, interactive, database-backed page in any browser. It is perfect to document a specific domain, use-case, question or graph problem.

## Examples

GraphGists work like any AsciiDoc document, but they allow you to insert special comments to define how data from Neo4j can be displayed and interacted with.

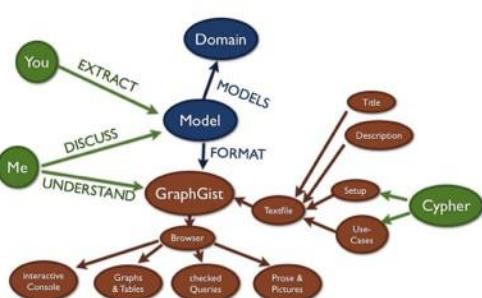
First define a block of [Cypher](#) code to setup the database:

```
//setup
//hide
[source,cypher]
-----
CREATE (:Database {name:'Neo4j'})-[:SUPPORTS]->(:Language {name:'Cypher'})
-----
```

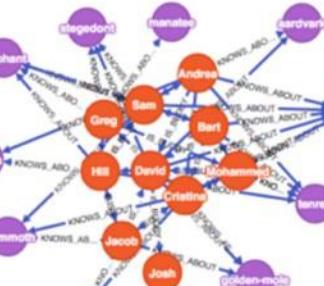
You can then make queries to the data and output them as a table and/or a graph:

```
[source,cypher]
-----
MATCH (db:Database)-[:SUPPORTS]->(language:Language)
RETURN db.name, collect(language.name)
-----
//table
```

```
[source,cypher]
-----
MATCH (db:Database)-[rel:SUPPORTS]->(:Language)
RETURN rel
-----
//graph
```



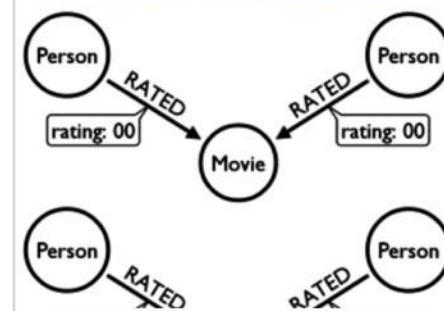
## Aardvark: The Anatomy of a ...



Aardvark (2008 - 2011) was a social search service that connected users live with friends or fr...

**Author:** whatSocks

## Movie Recommendations wi...



In this Graph Gist, I'm using k-NN with cosine similarity as the similarity metric to calcula...

**Author:** Nicole White

## Insight #4: Similar Repositories

Grouping by #contributors will exclude repositories with commits using a single account - [/limitation](#)

```
Query 5
MATCH (a)-[r1:IS_ACTOR]->(MATCH)-<-[r2:IS_ACTOR]-(b) WHERE a.id > b.id
WITH a,b, collect(DISTINCT MATCH.id) AS connections, collect(DISTINCT type(r1)) AS rel1
WHERE length(connections) >= 1 //set minimum # of connections
RETURN a.id, b.id,length(connections) AS count ORDER BY length(connections) DESC
```

Test run OK

Show

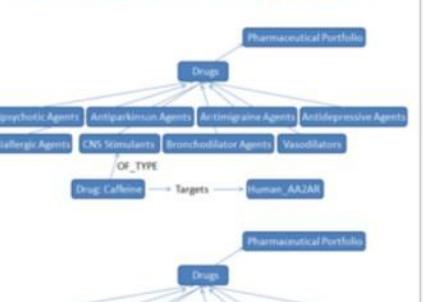
10

entries

Search:

a.id	b.id	count
hakujin/homebrew	BrewTestBot/homebrew	3
x1022as/docker	oceannwu/docker	3
cocos2d/cocos2d-x	andyque/cocos2d-x	3
dingpingli/cocos2d-x	andyque/cocos2d-x	3
		2
		2
		2
		2
		2

## Pharmaceutical Drugs and t...



This Graph Gist explores how to represent a pharmaceutical portfolio in a property graph. A phar...

**Author:** Josh Kunken

<http://neo4j.com/graphgists/>

# THANK YOU!



[will@neo4j.com](mailto:will@neo4j.com)

@lyonwj